

# **Neural Networks**

## **via FluCoMa In Max**

### **Music Hackspace**

**Ted Moore**

[tedmooremusic.com](http://tedmooremusic.com)

[ted@tedmooremusic.com](mailto:ted@tedmooremusic.com)



# FluCoMa: Fluid Corpus Manipulation



- Toolbox for Programmatic Data Mining of Sound Banks
- Integrating Machine Listening and Machine Learning in...
- Max, SuperCollider & Pure Data
- Learning Resources ([learn.flucoma.org](https://learn.flucoma.org))
- Discourse Community ([discourse.flucoma.org](https://discourse.flucoma.org))



Ted Moore, (James Bradbury, Owen Green, Jacob Hart, Gerard Roma, & Pierre Alexandre Tremblay)

University of Huddersfield | UK

# Plan for Today

- **Training a Neural Network (refresher)**
  - classification
- **Validating a Trained Model**
- **MLP Parameters Overview**
- **Adjusting parameters during training**
  - regression with control data
  - regression with real-time analysis
- ***tapin* and *tapout***
  - wavetable autoencoder (advanced)

# Classification



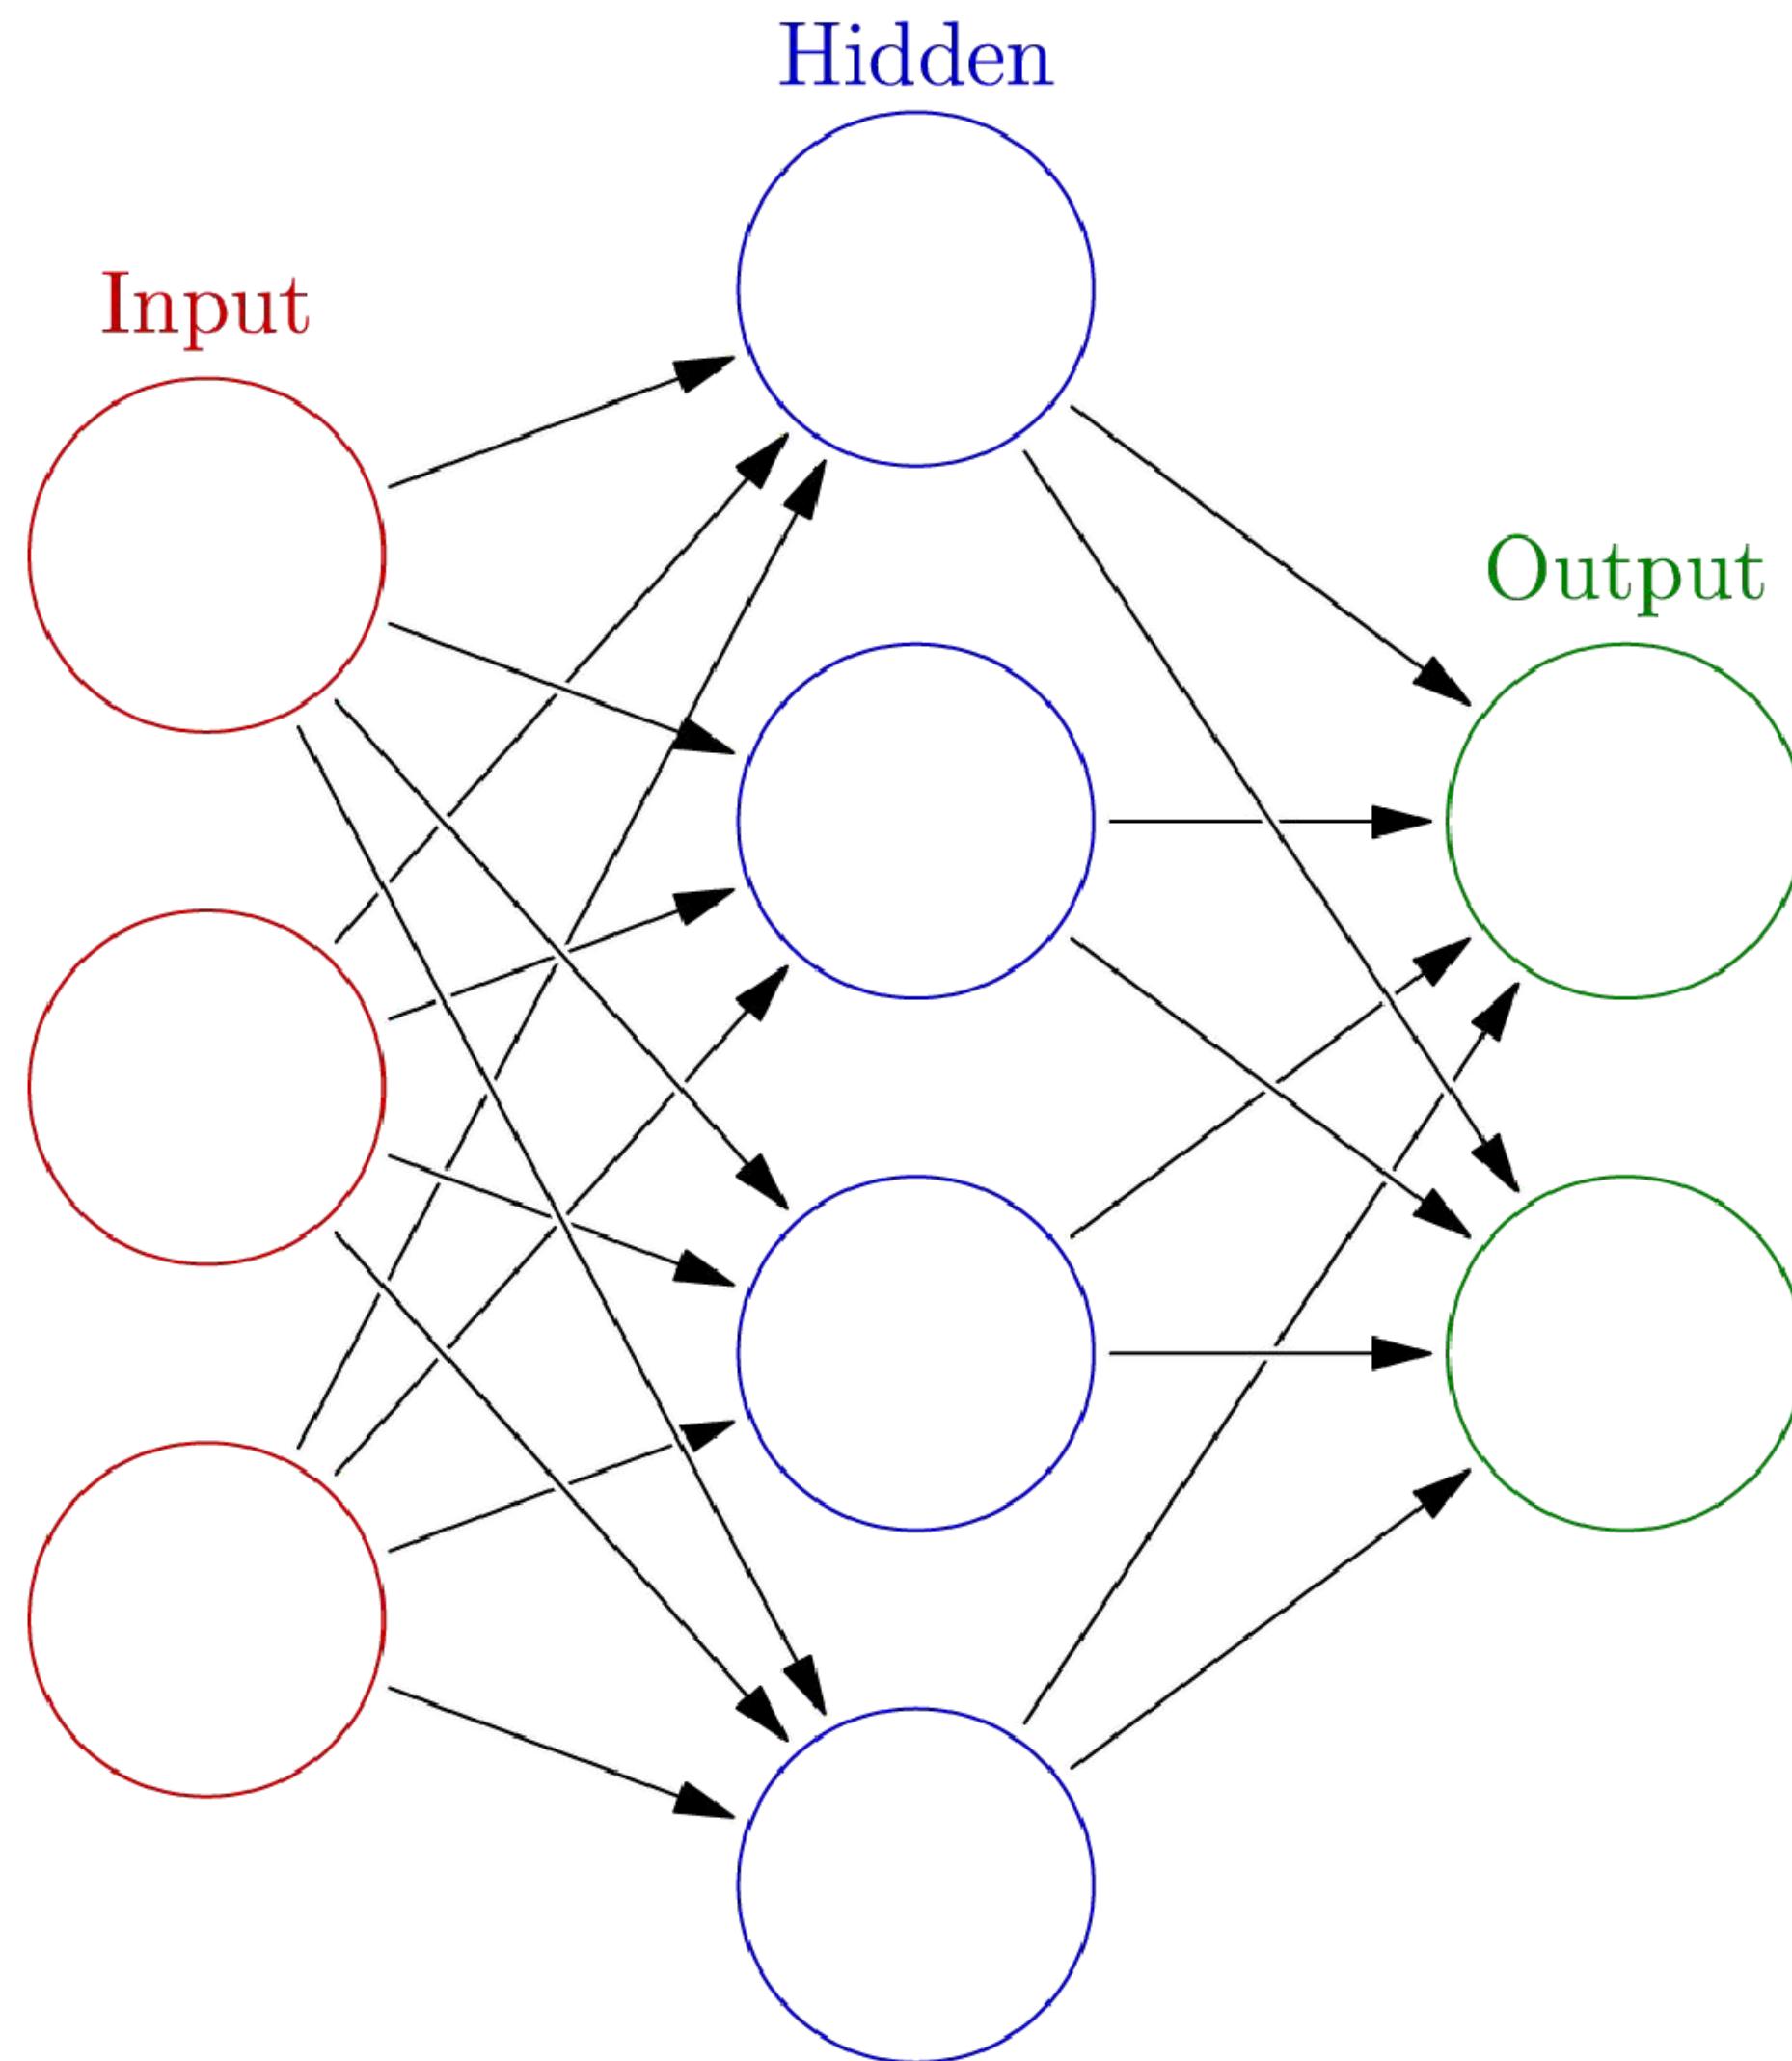
# **feed**

using a neural network for real-time  
audio classification

training a machine to hear the way I hear

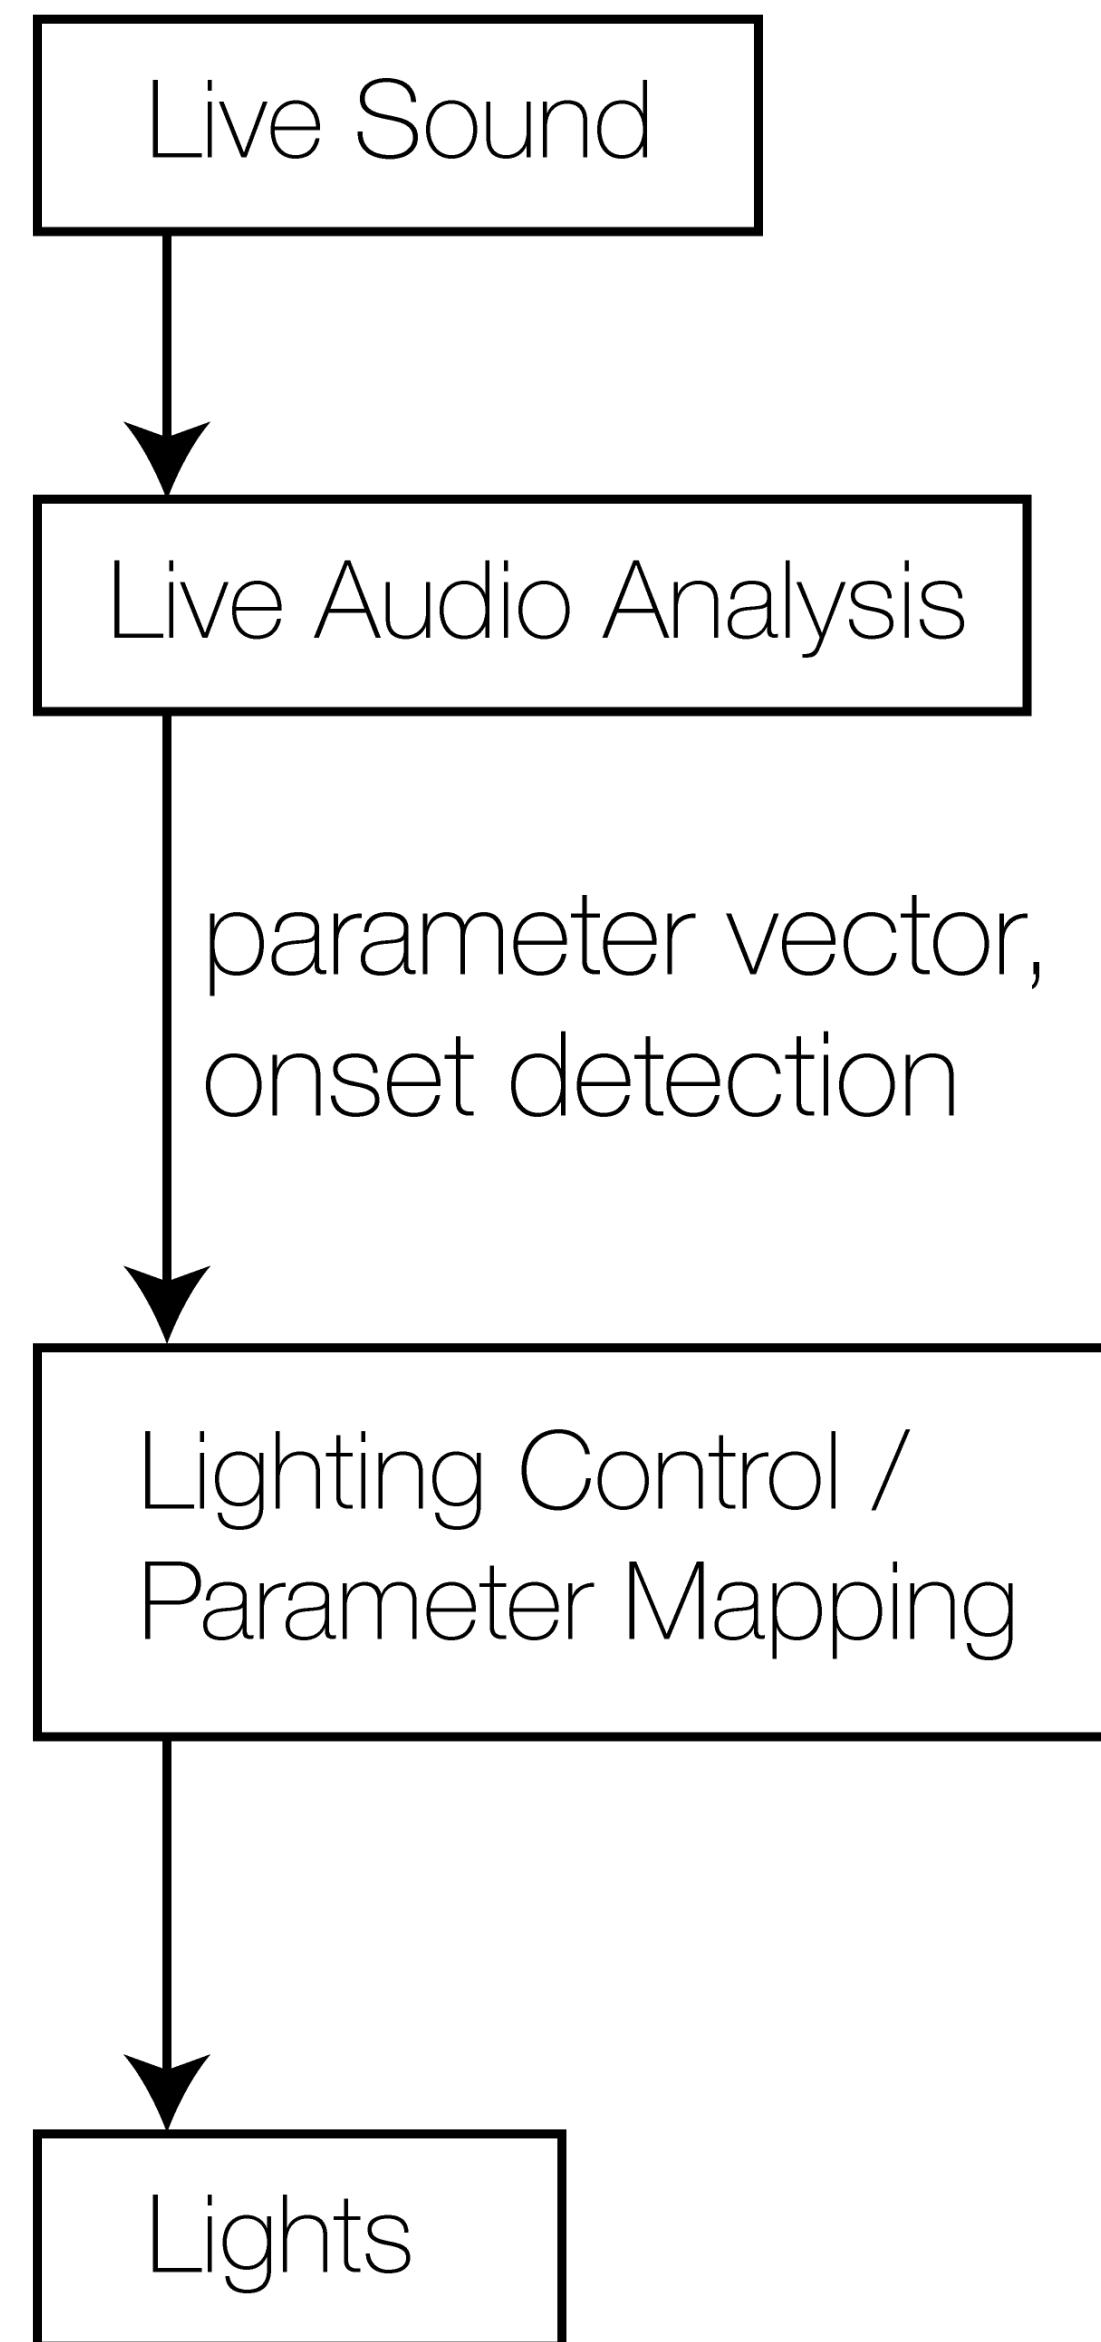


# *Neural Network (Multi-Layer Perceptron)*

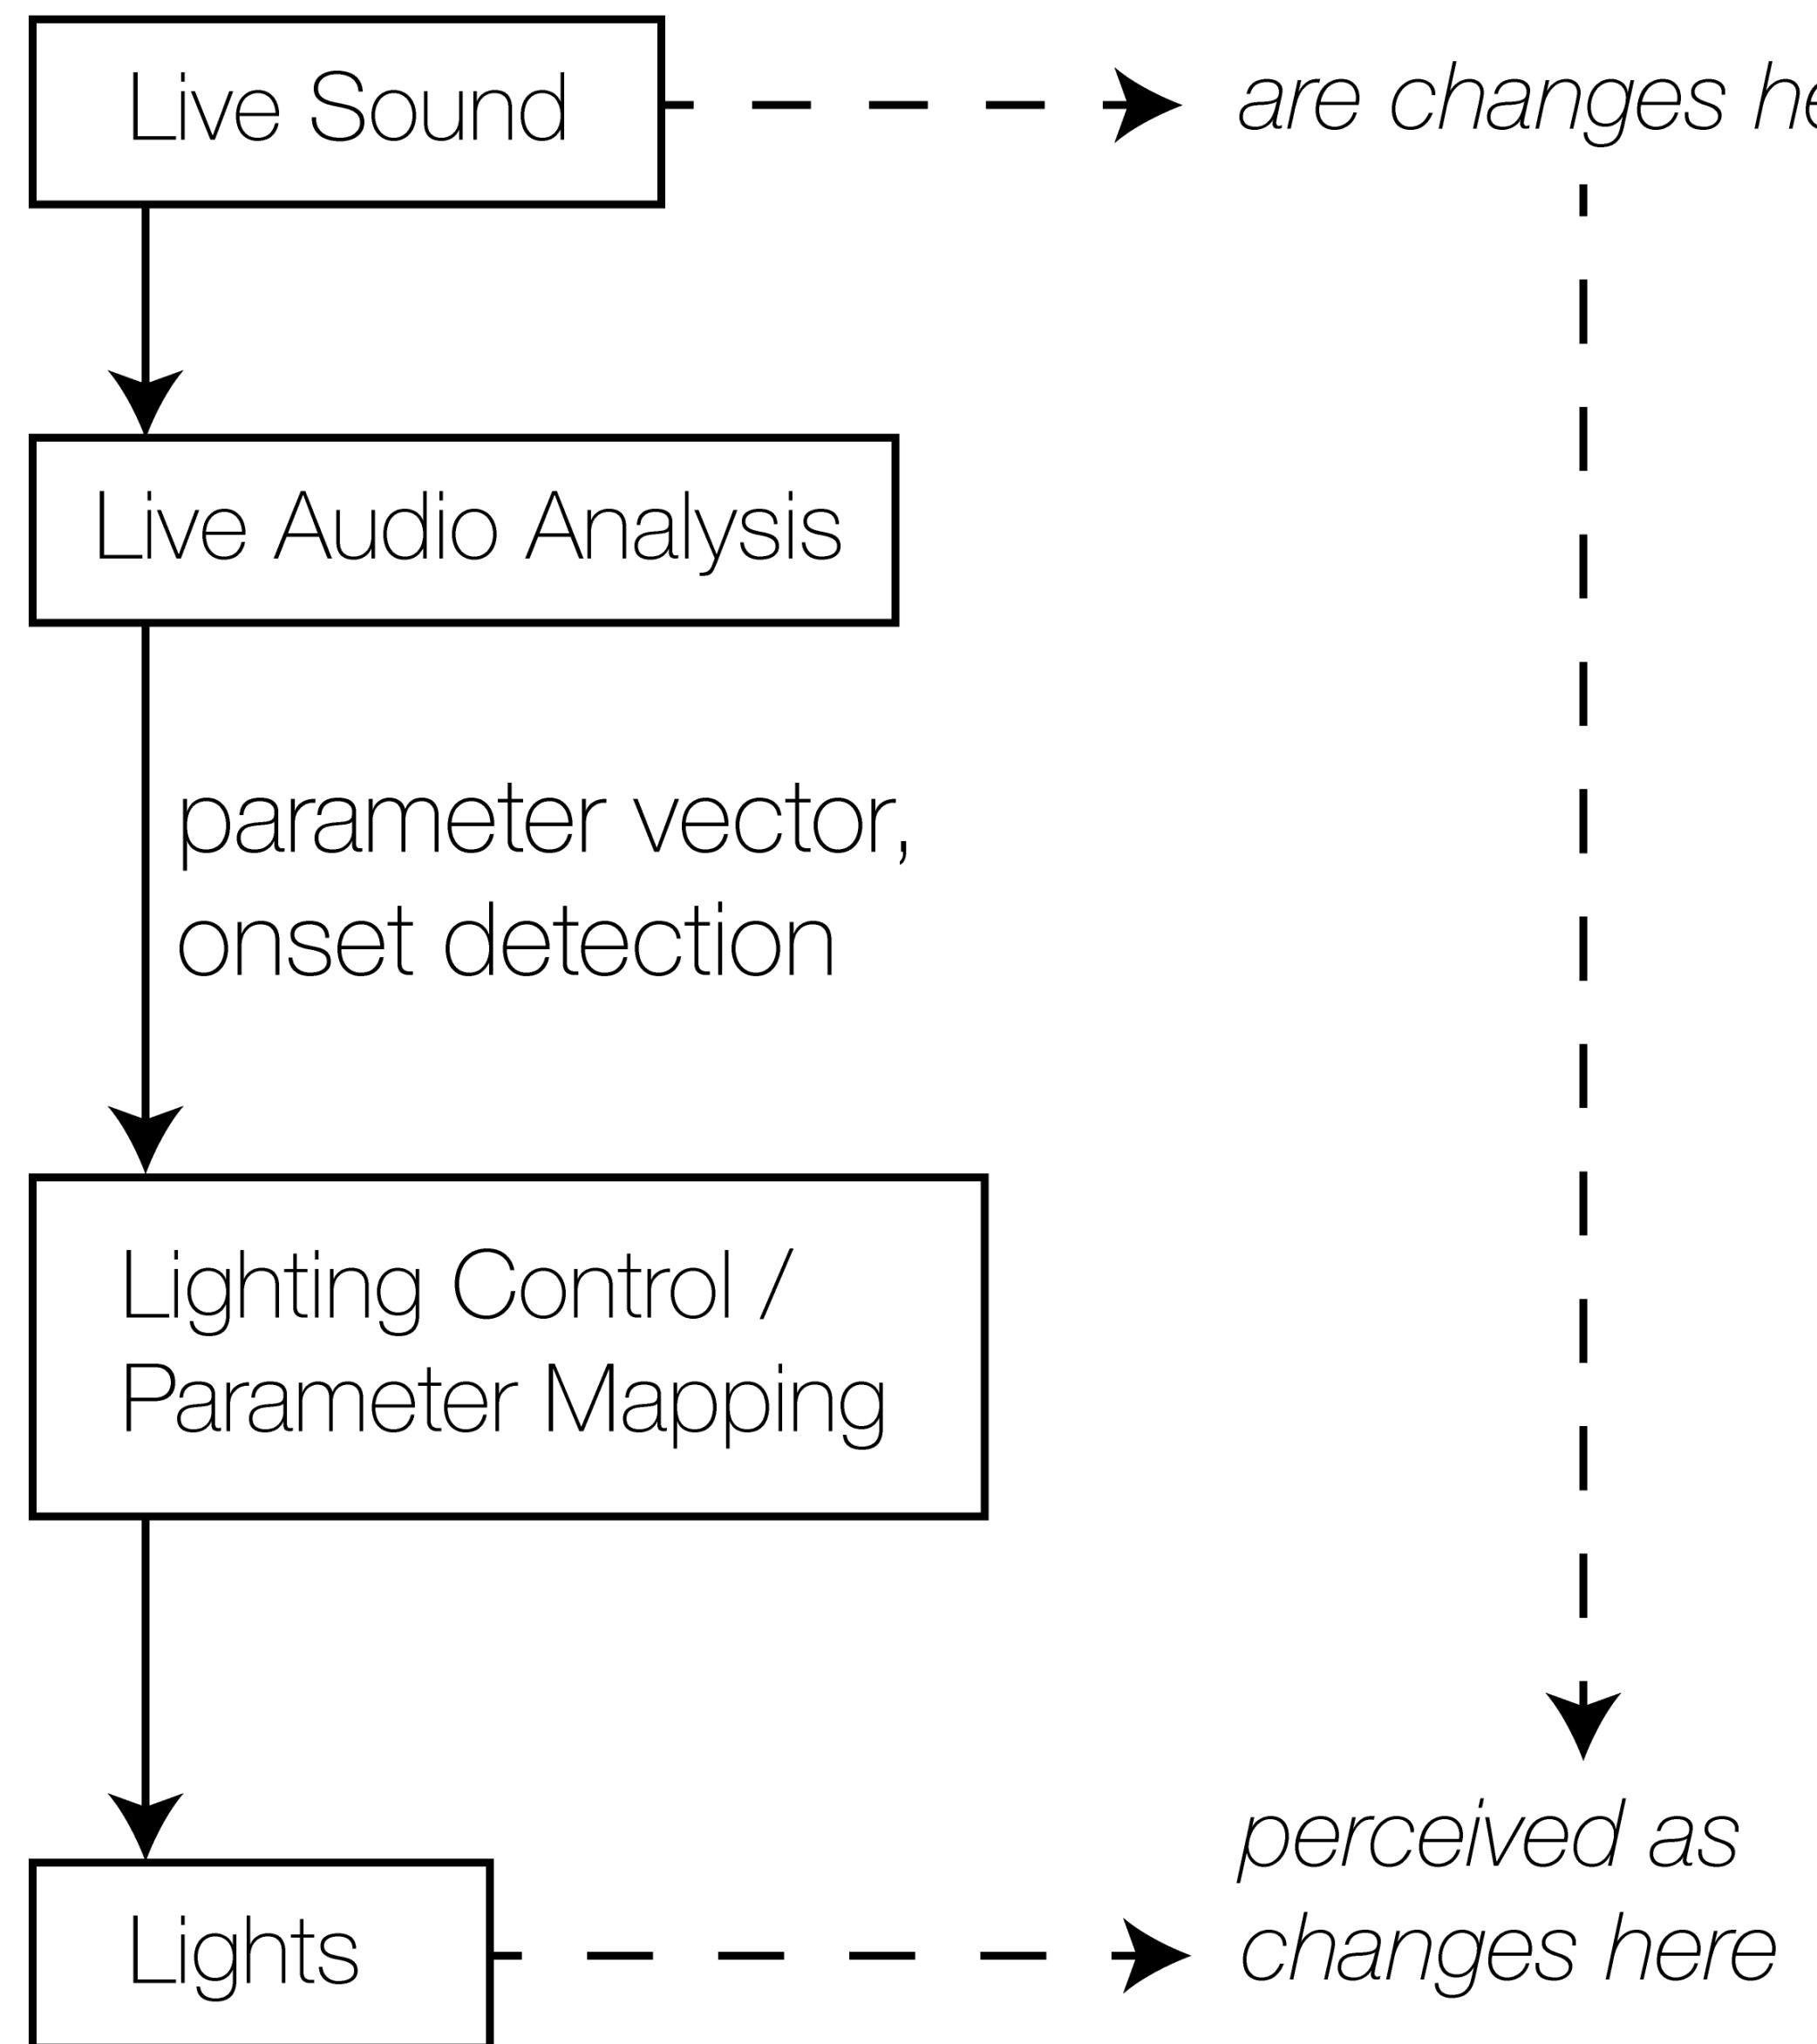


**Classification:**  
a neural network predicts  
which category (or “class”)  
an input belongs to

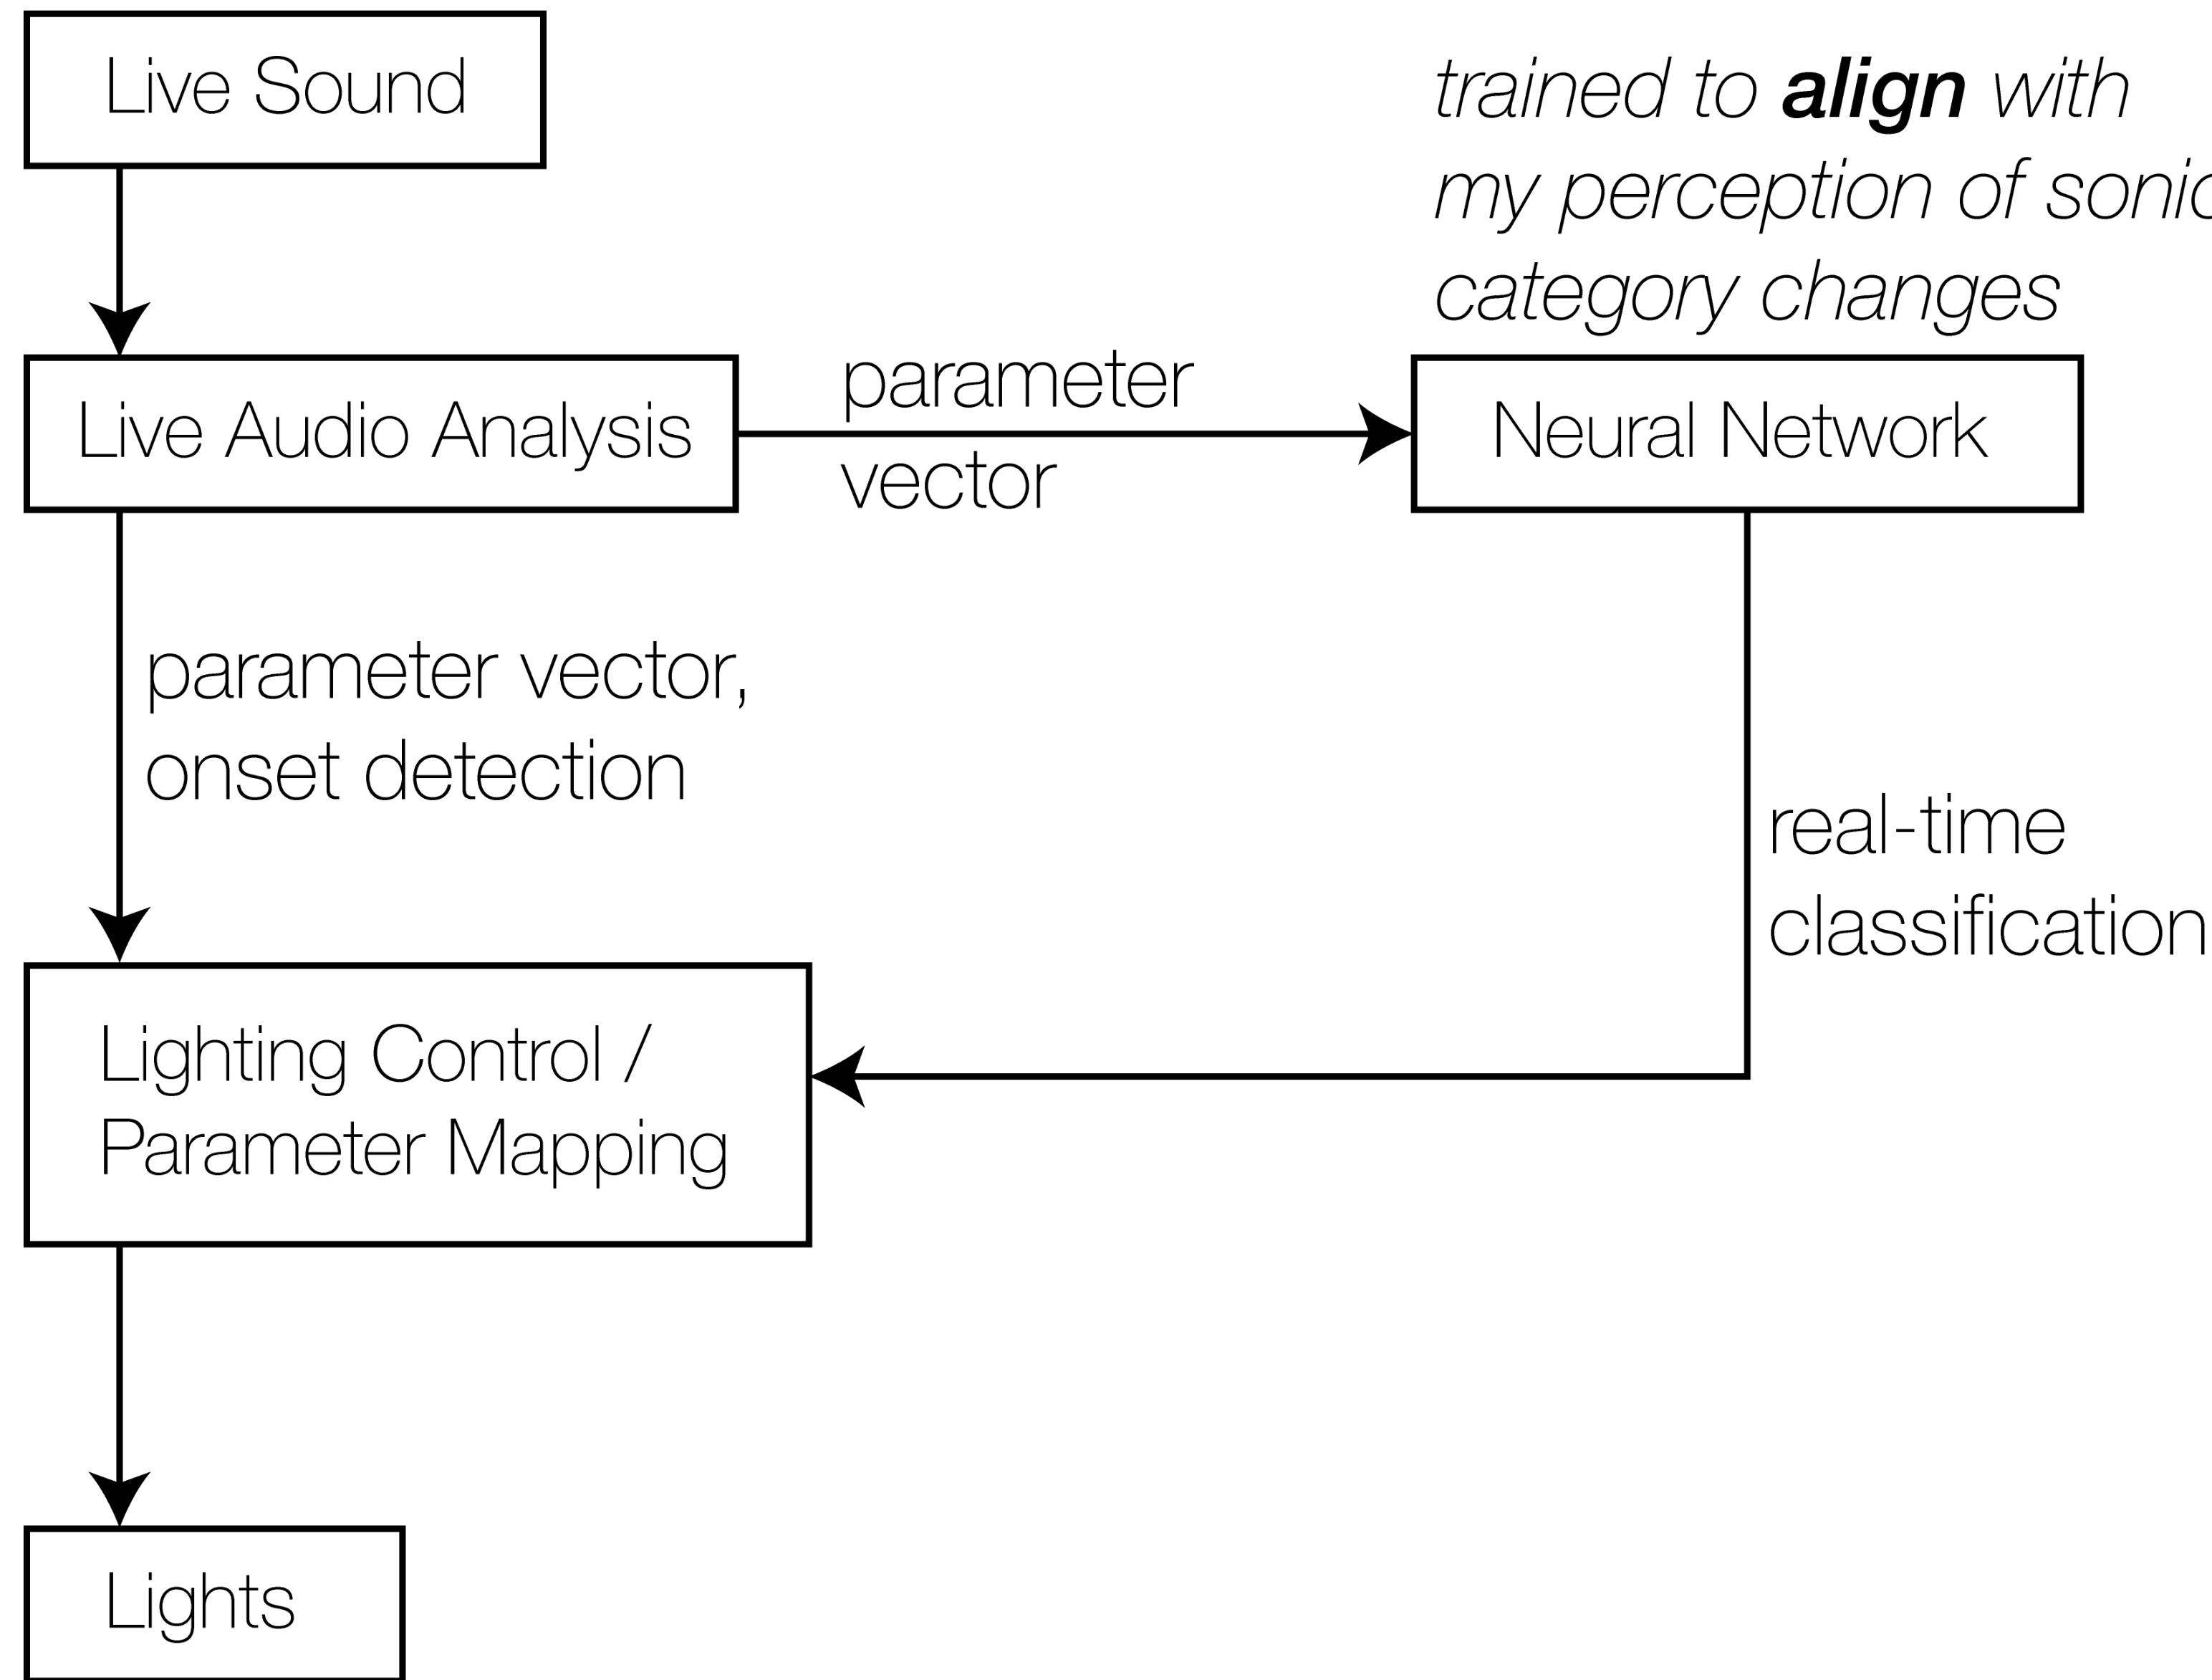
# Machine Listening System



# Machine Listening System



# Machine Learning System





distorted\_noise



high\_squeal



low\_impulses



sus\_noise\_quiet

**flashing lights warning**





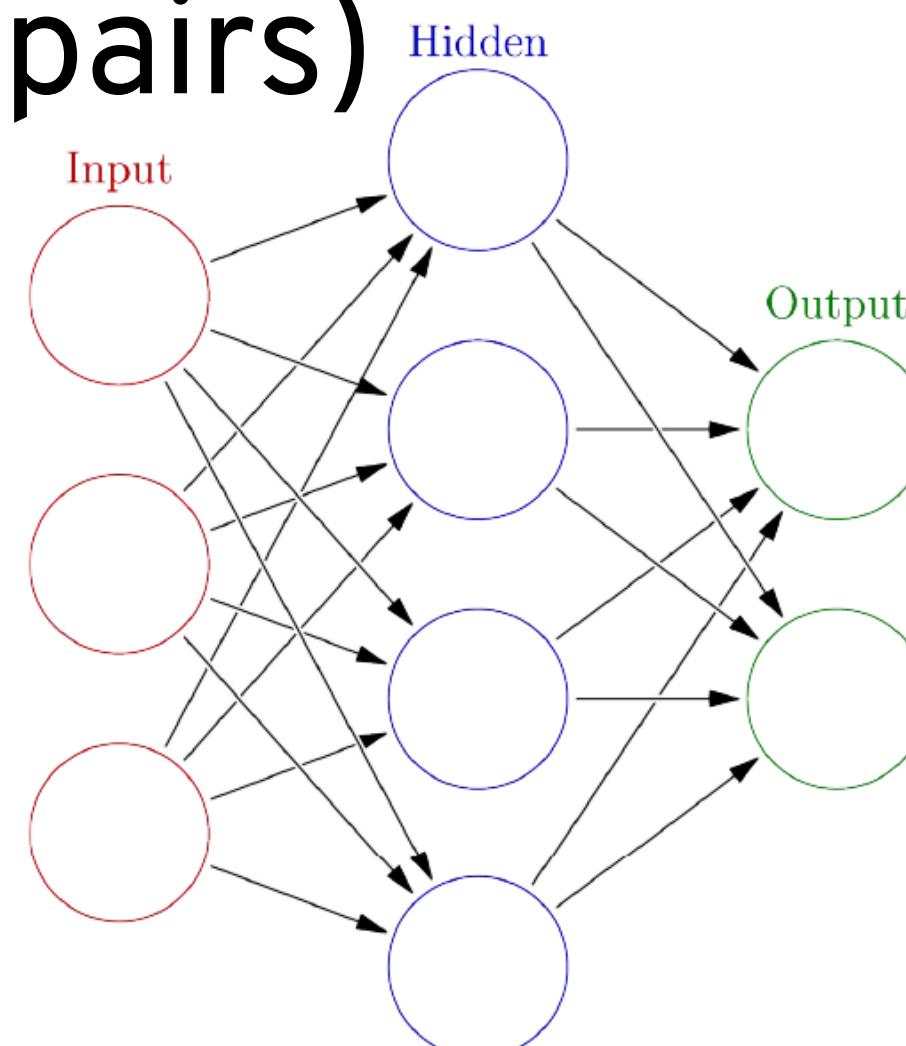
# *Neural Network* **Training a Classifier**

trombone  
examples

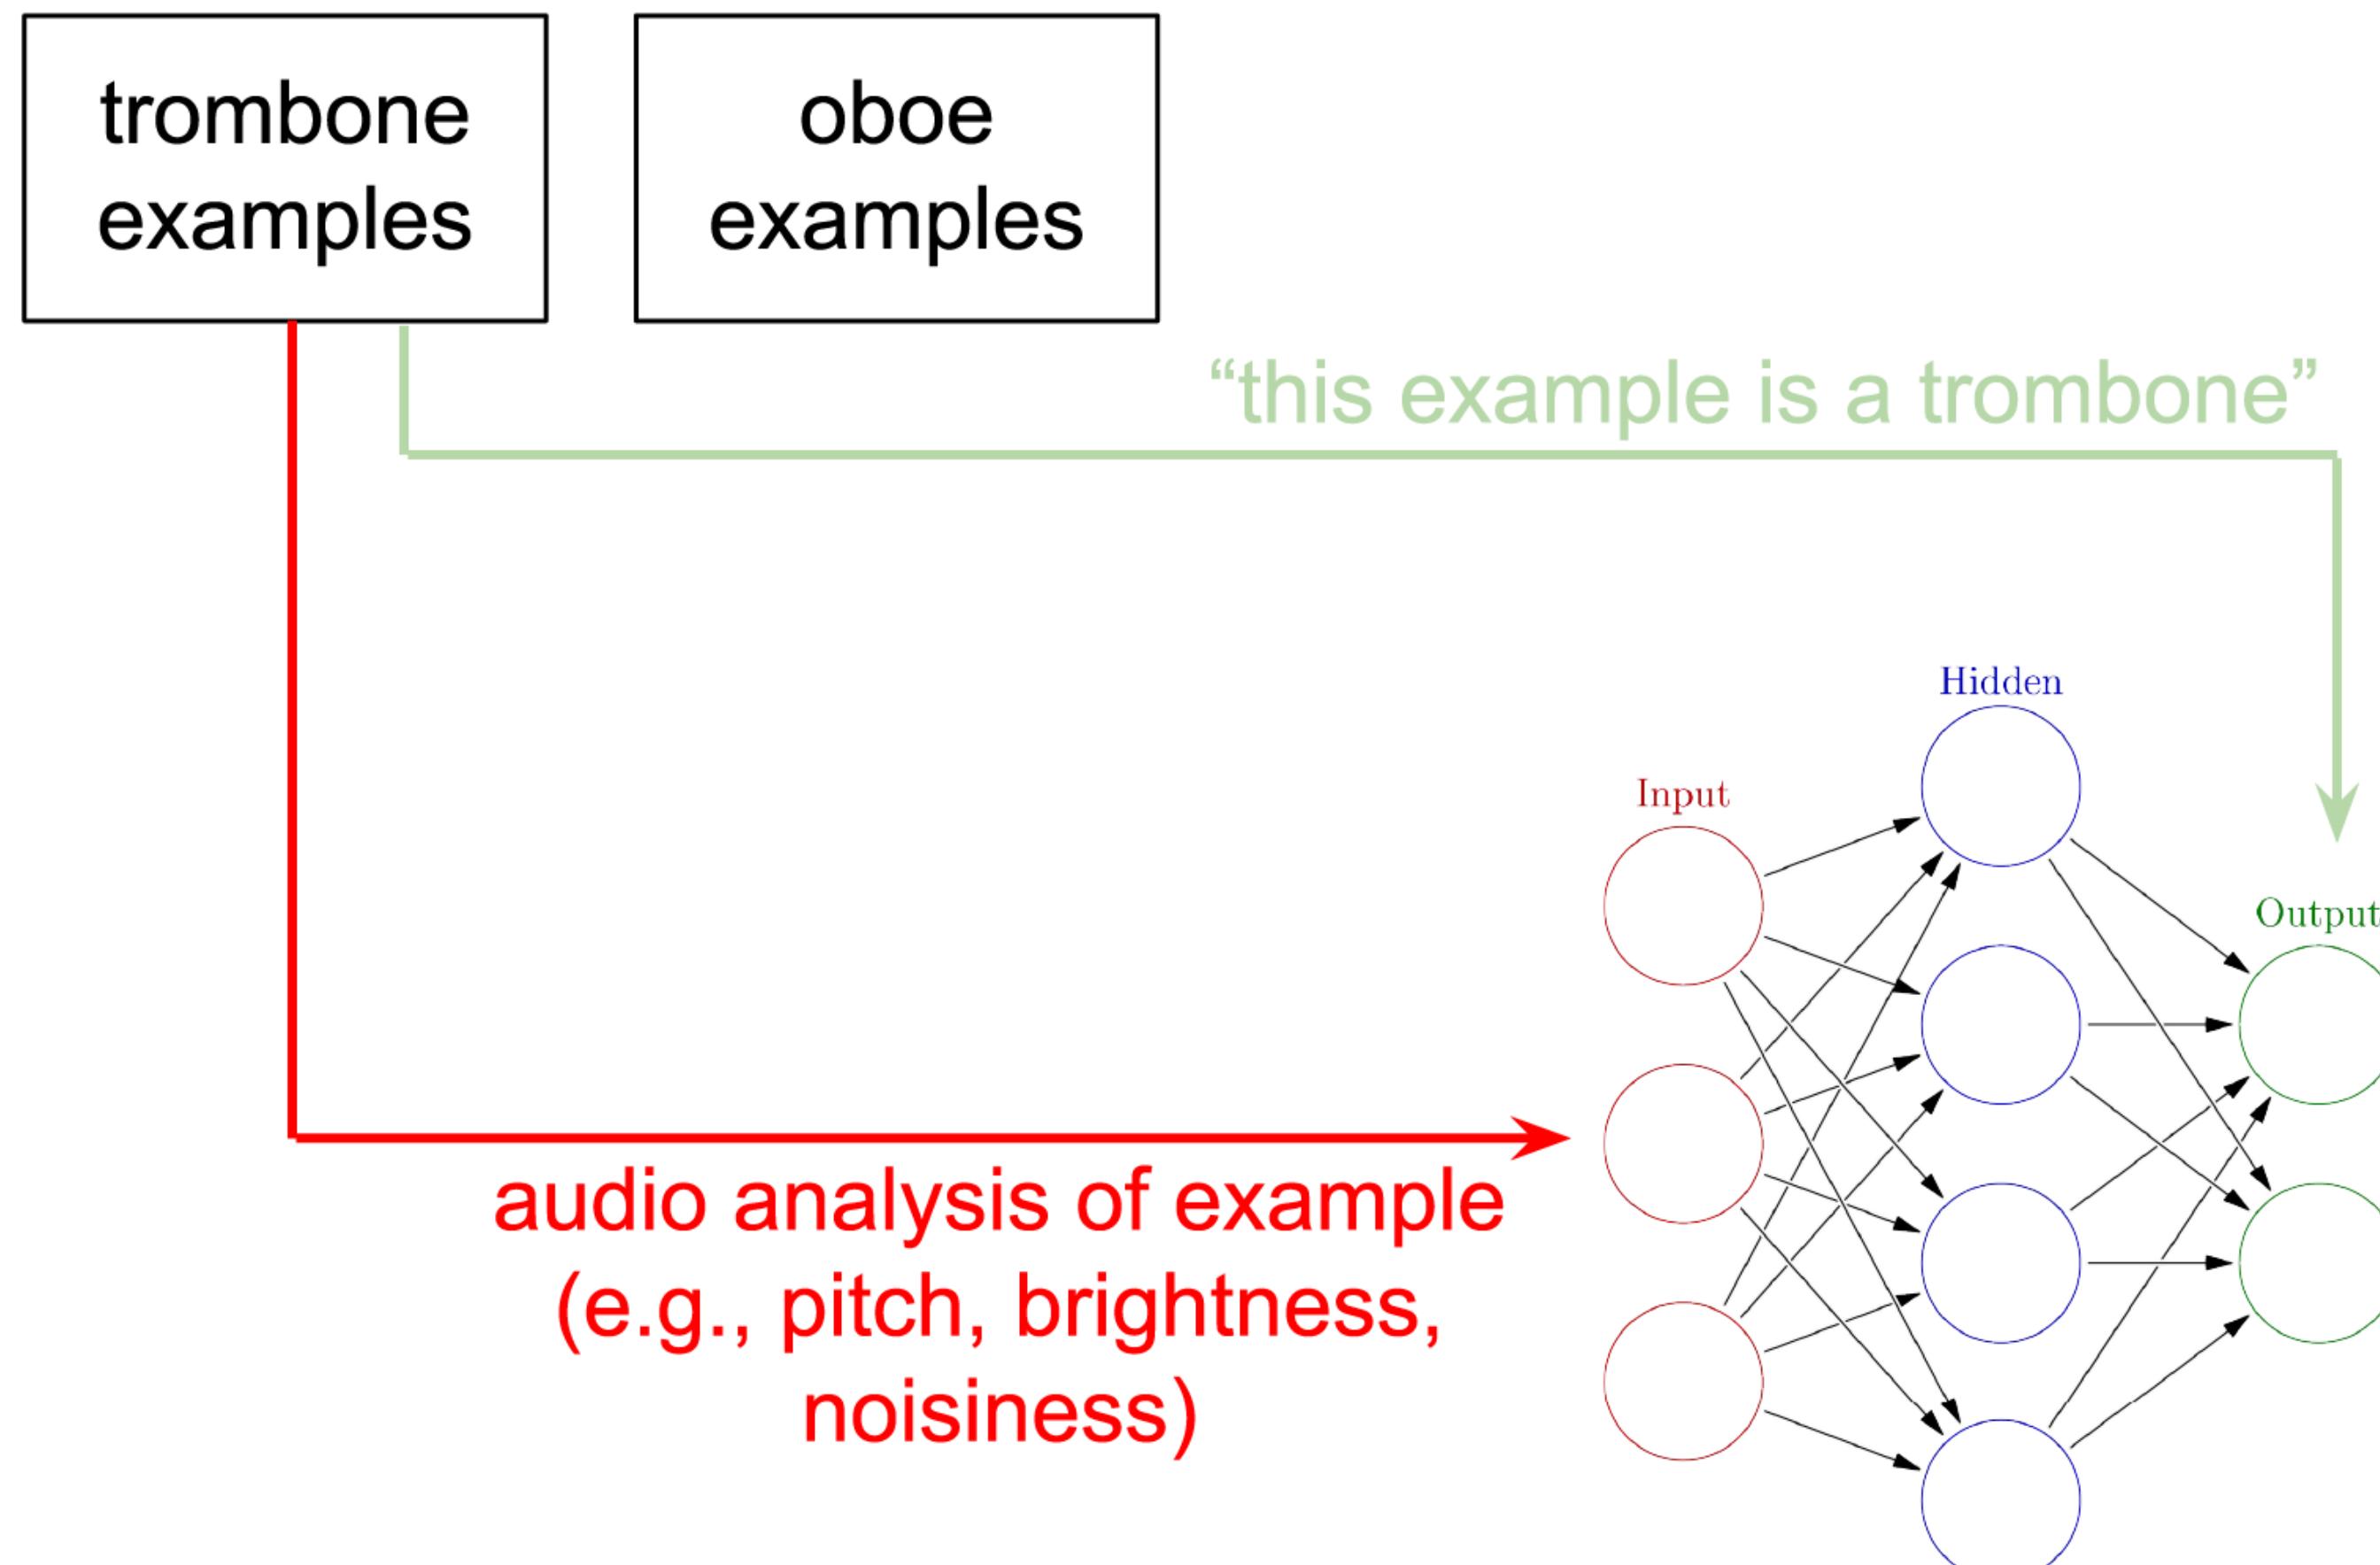
oboe  
examples

## **Supervised Learning**

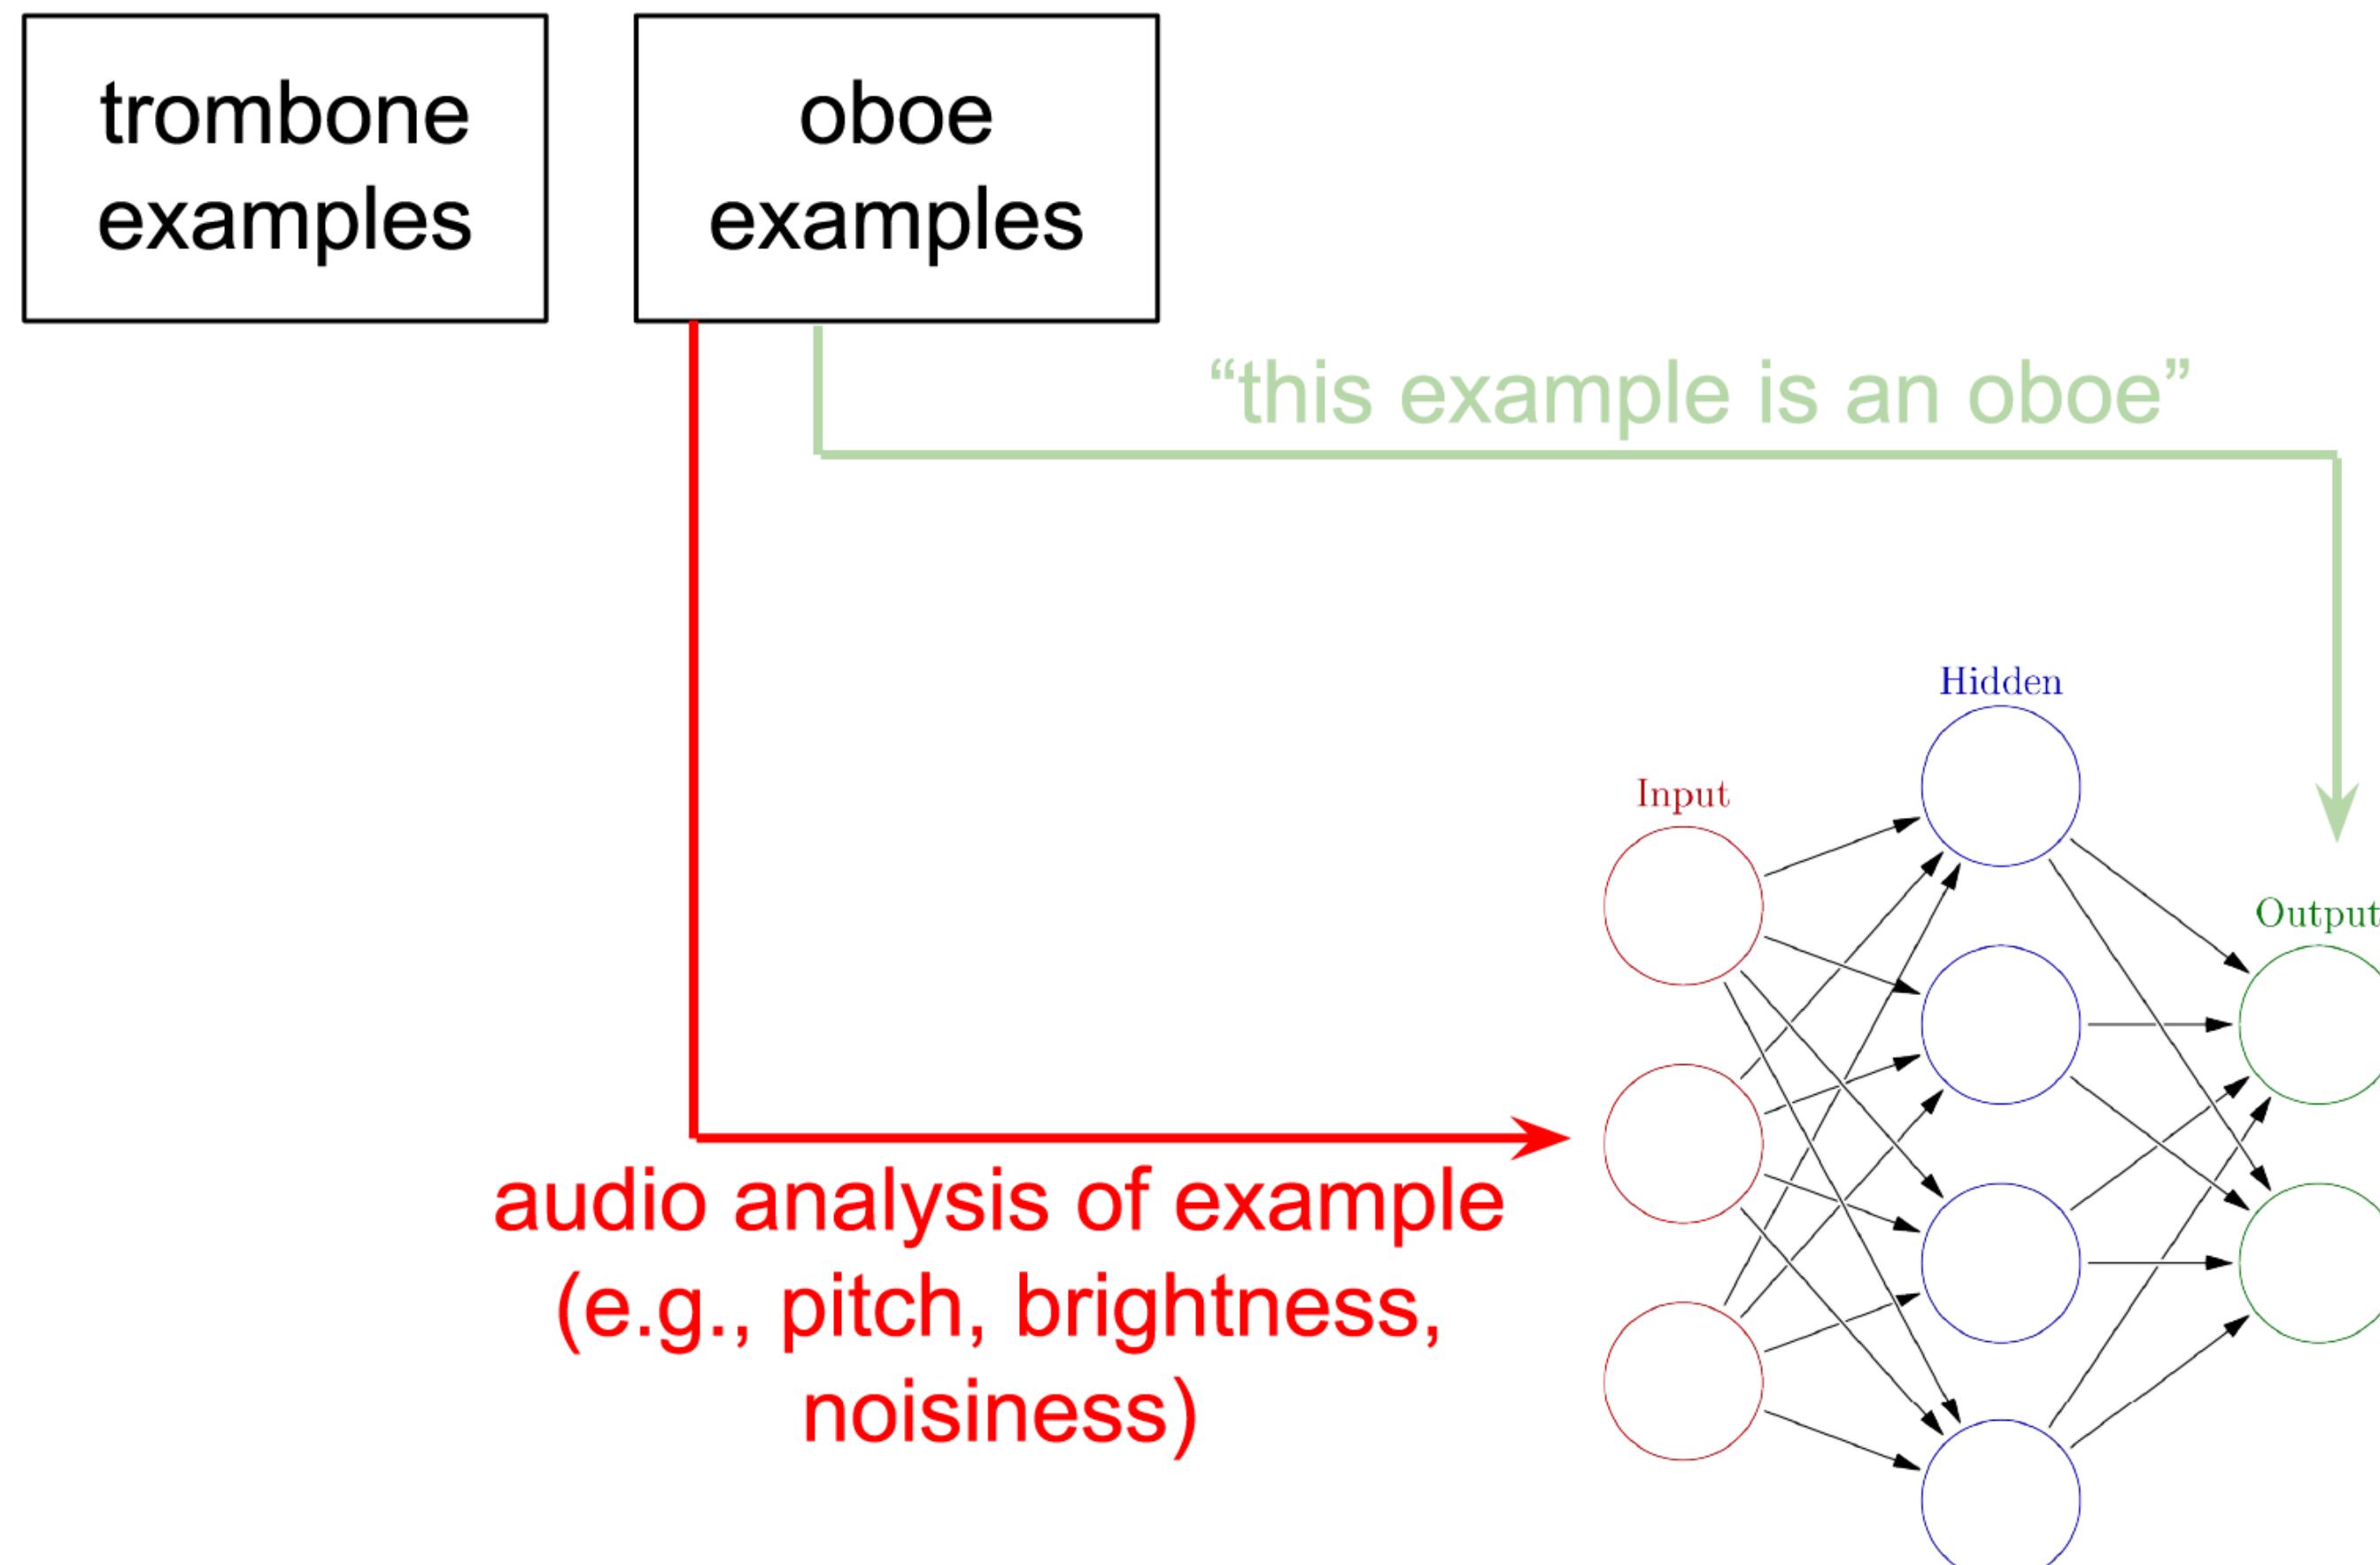
(uses input-output example pairs)



# *Neural Network* **Training a Classifier**



# *Neural Network Training a Classifier*



*feed-forward*  
vs.  
*back-propagation*

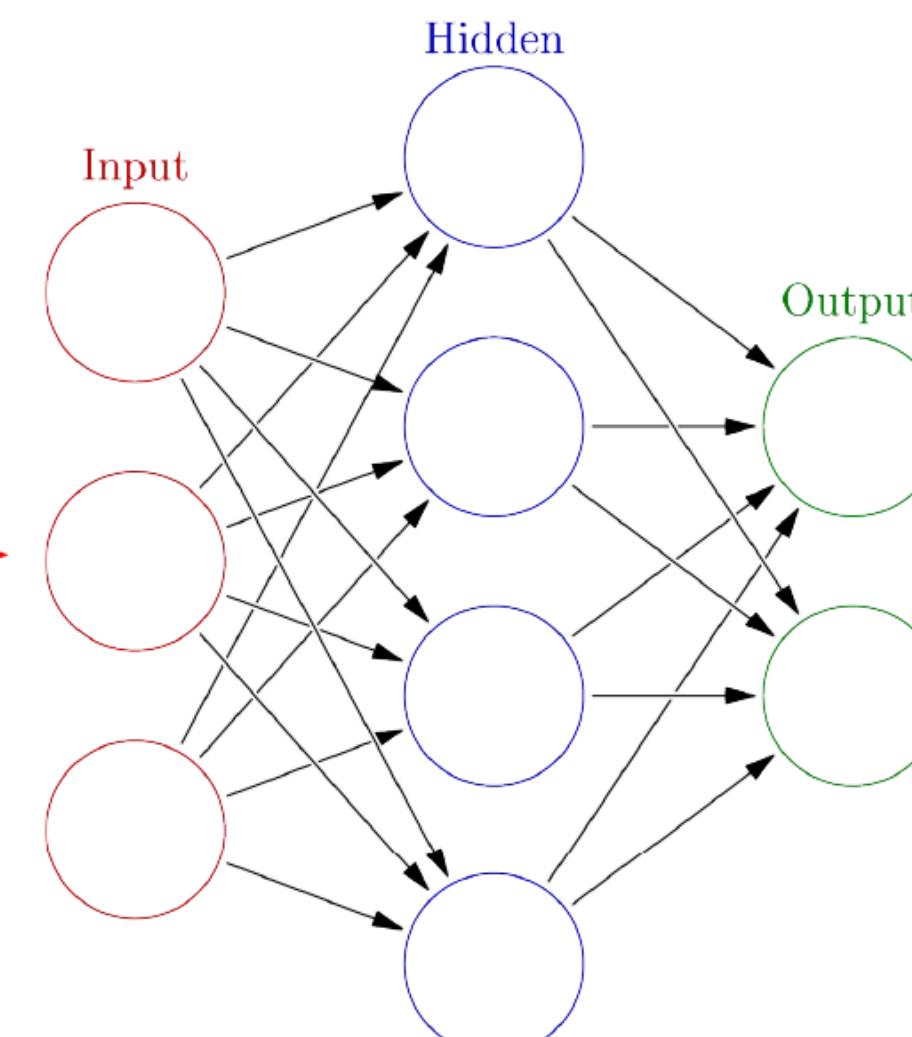
# Neural Network Predicting a Classification

trombone  
examples

oboe  
examples

new example it has  
never seen before

audio analysis of example  
(e.g., pitch, brightness,  
noisiness)



“this new example is most  
like the oboe examples you  
showed me before”  
(or trombone...)

*fluid.mlpclassifier~*

oboe vs. trombone



***fluid.mlpclassifier~***

change color mapping  
parameters



***fluid.mlpclassifier~***

validation



- **overfitting**
  - excellent on training data, poor on anything else
- **underfitting**
  - poor on everything...keep training!

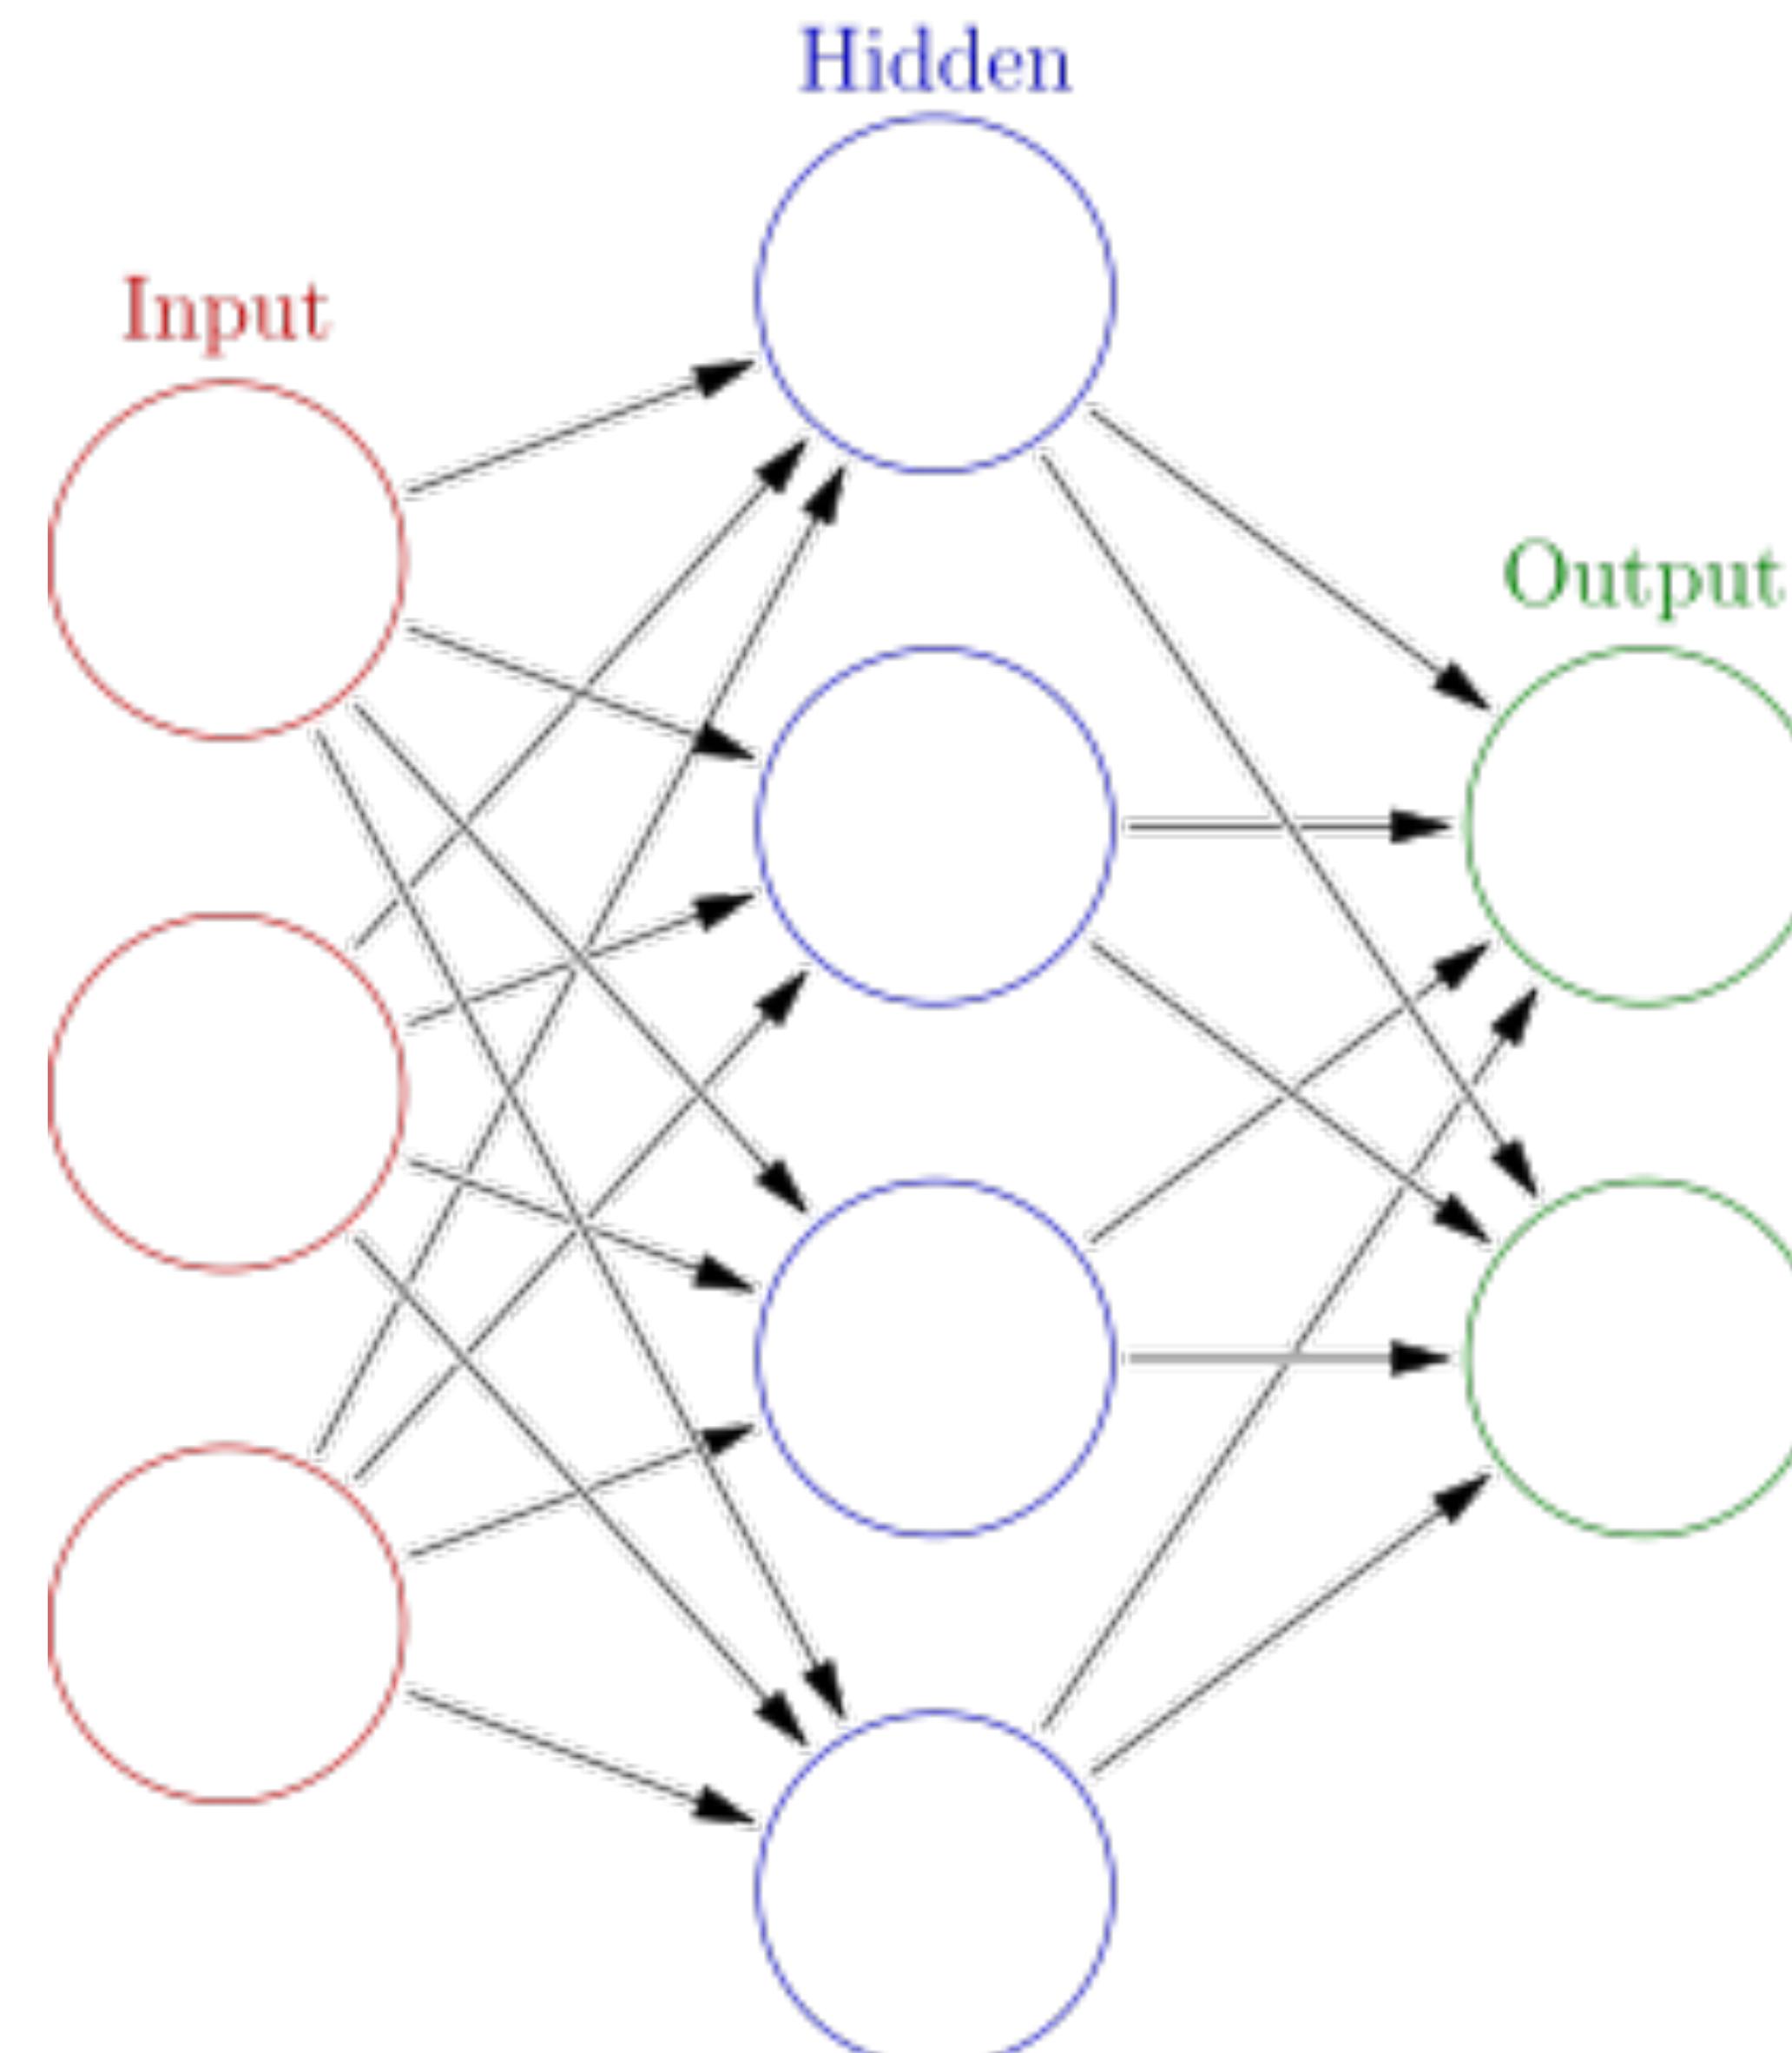
# *MLP Parameters*



@hiddenlayers (list)

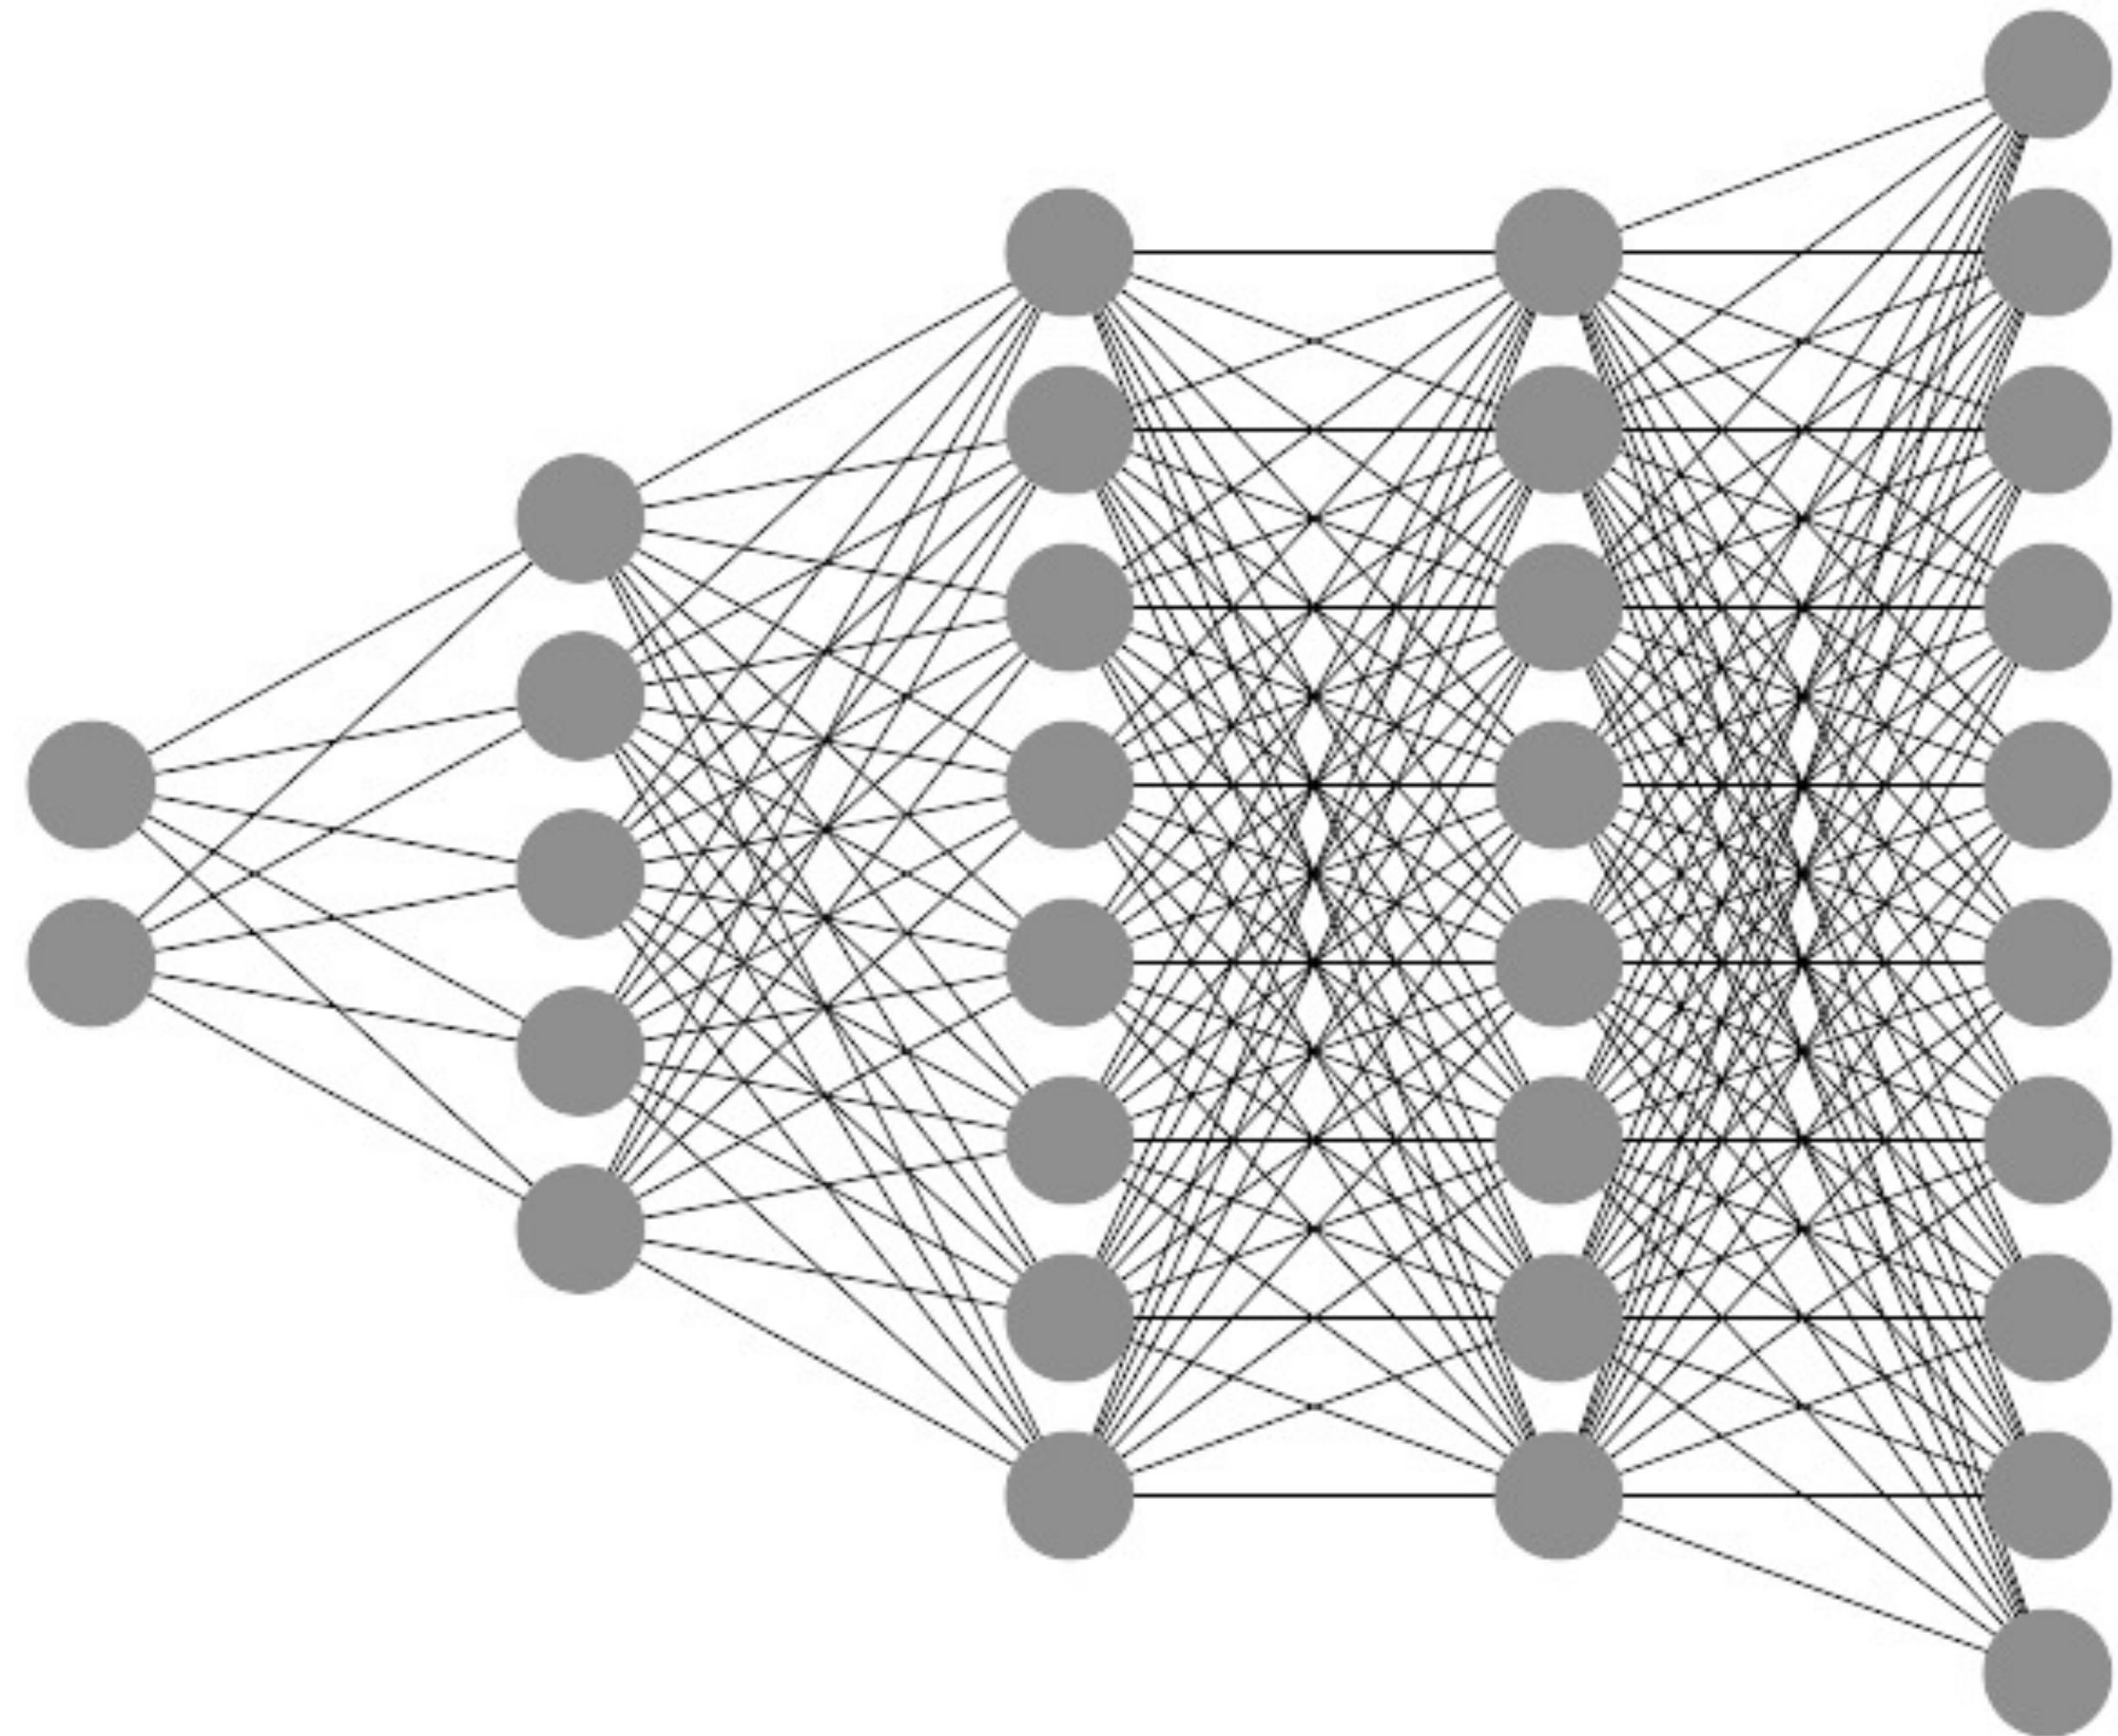
n elements = n layers

Integers = n neurons

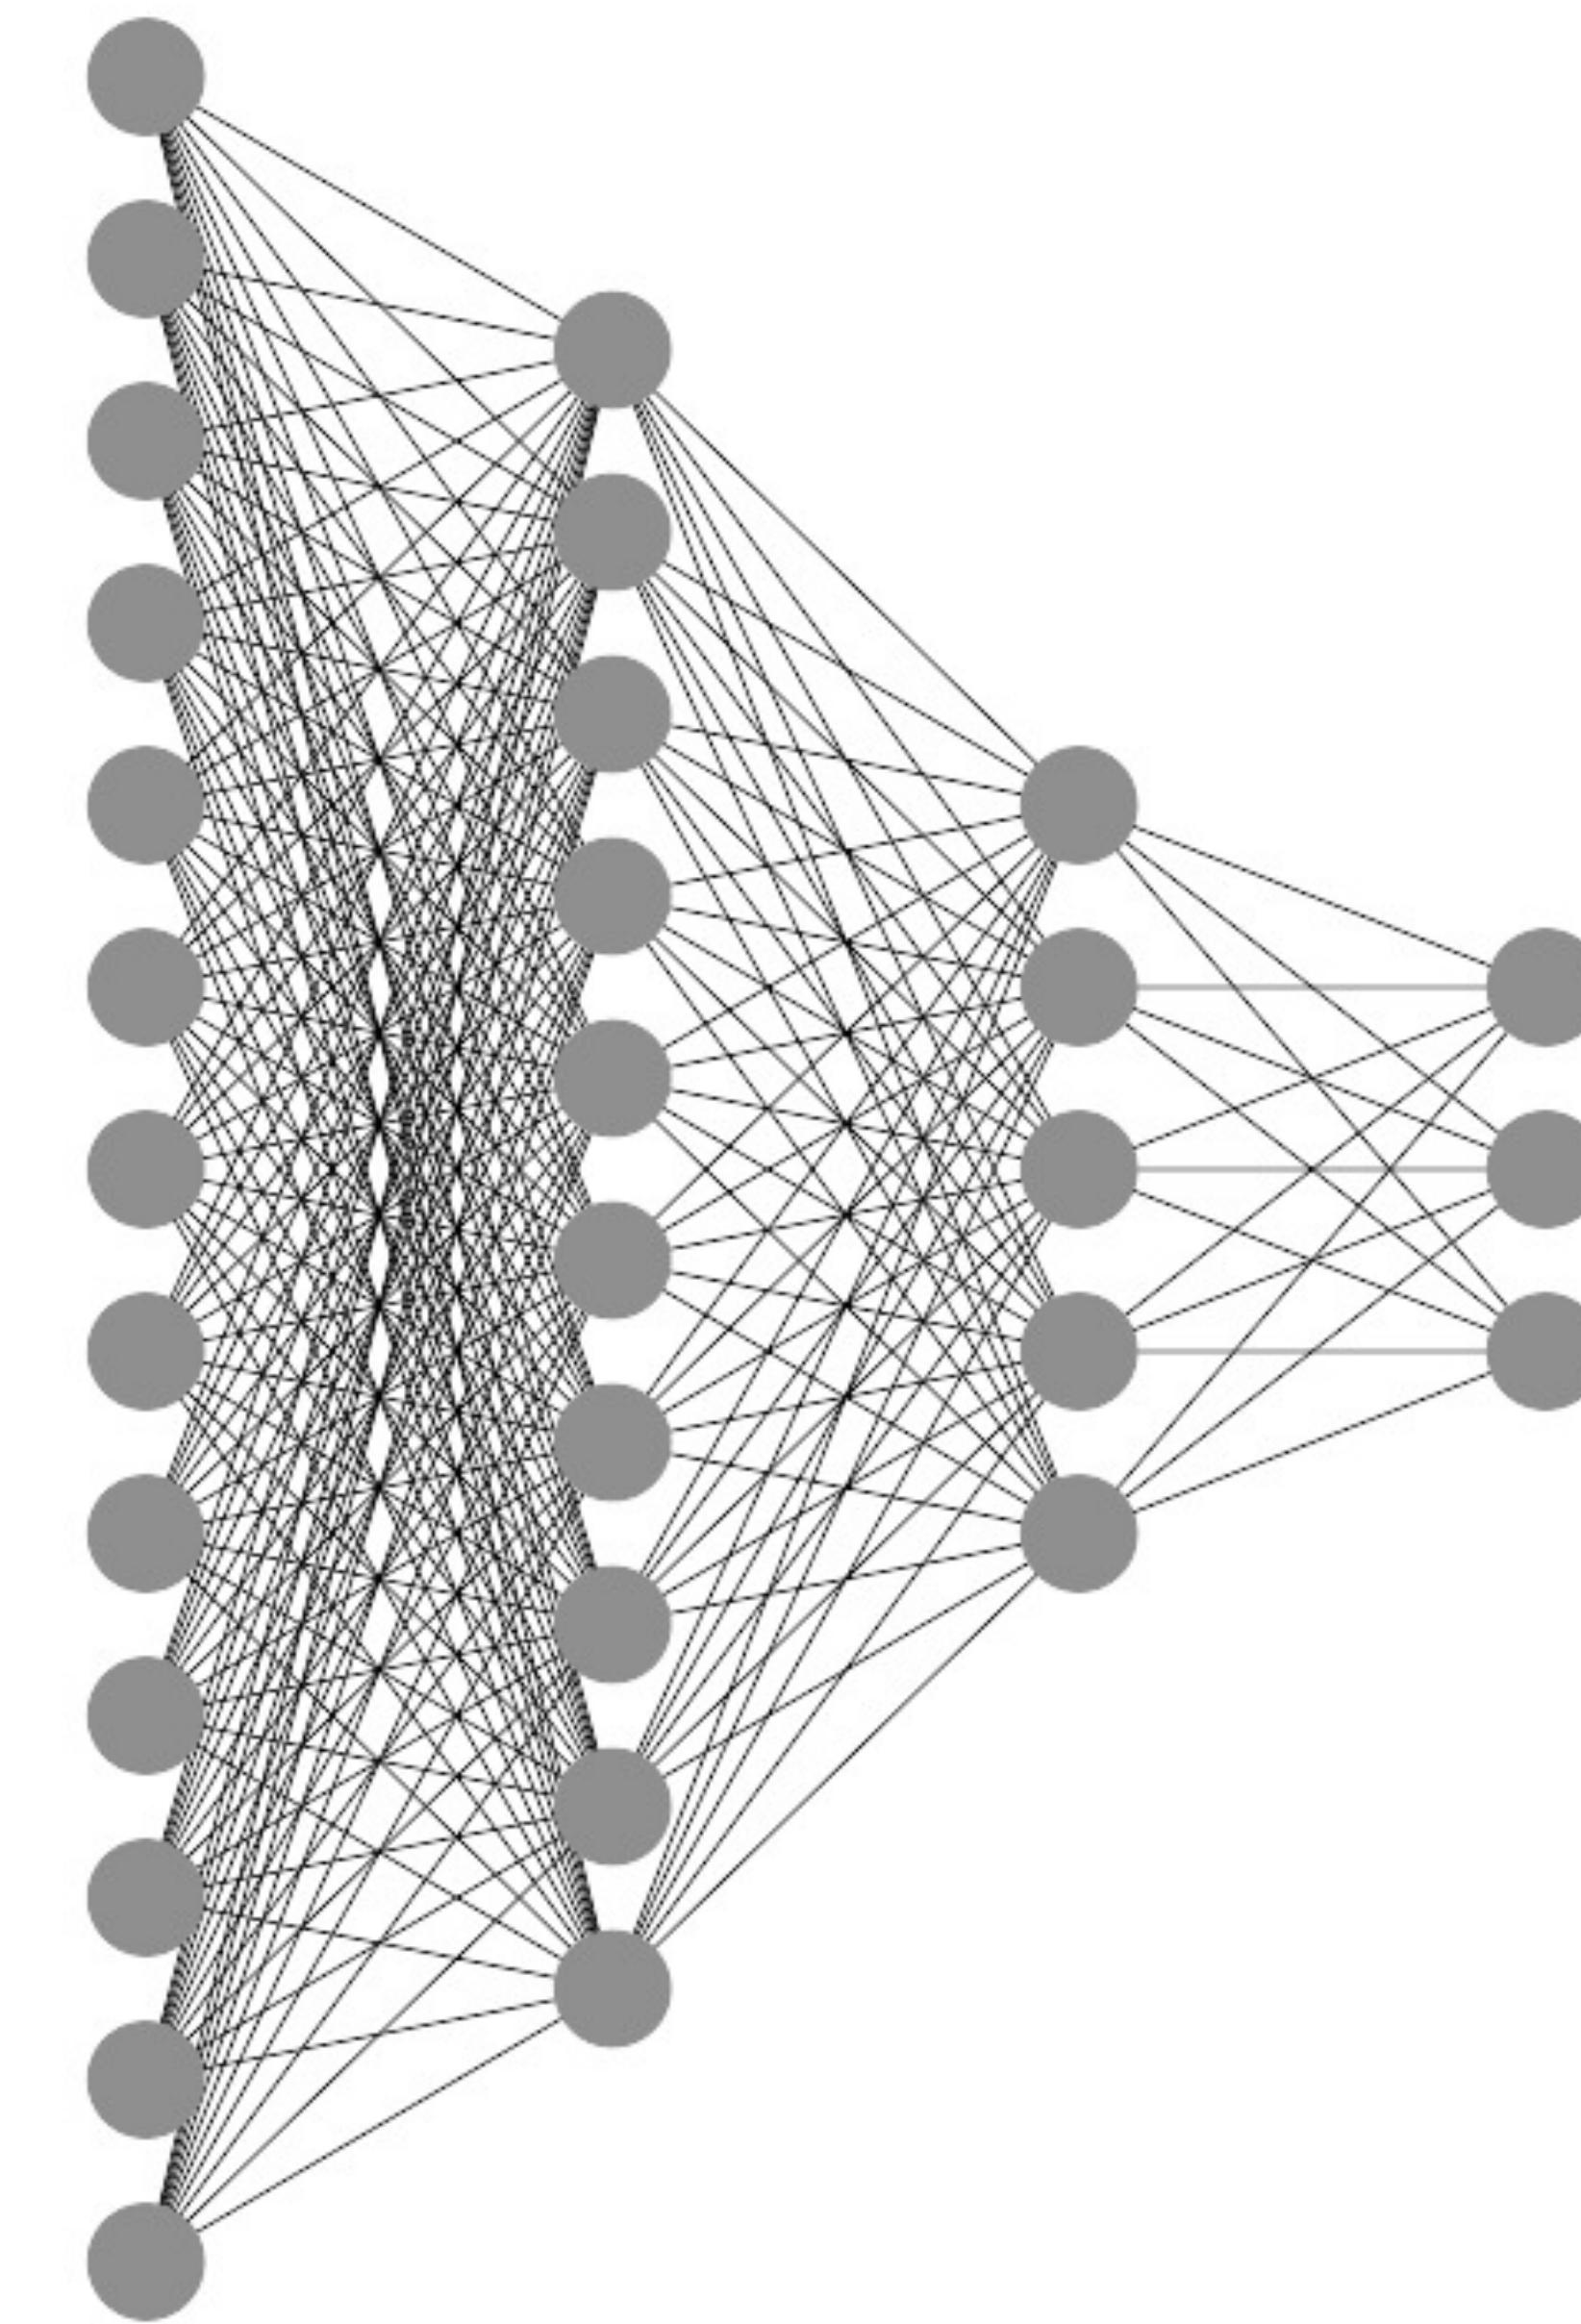


(@hiddenlayers 4)

*will reset network*



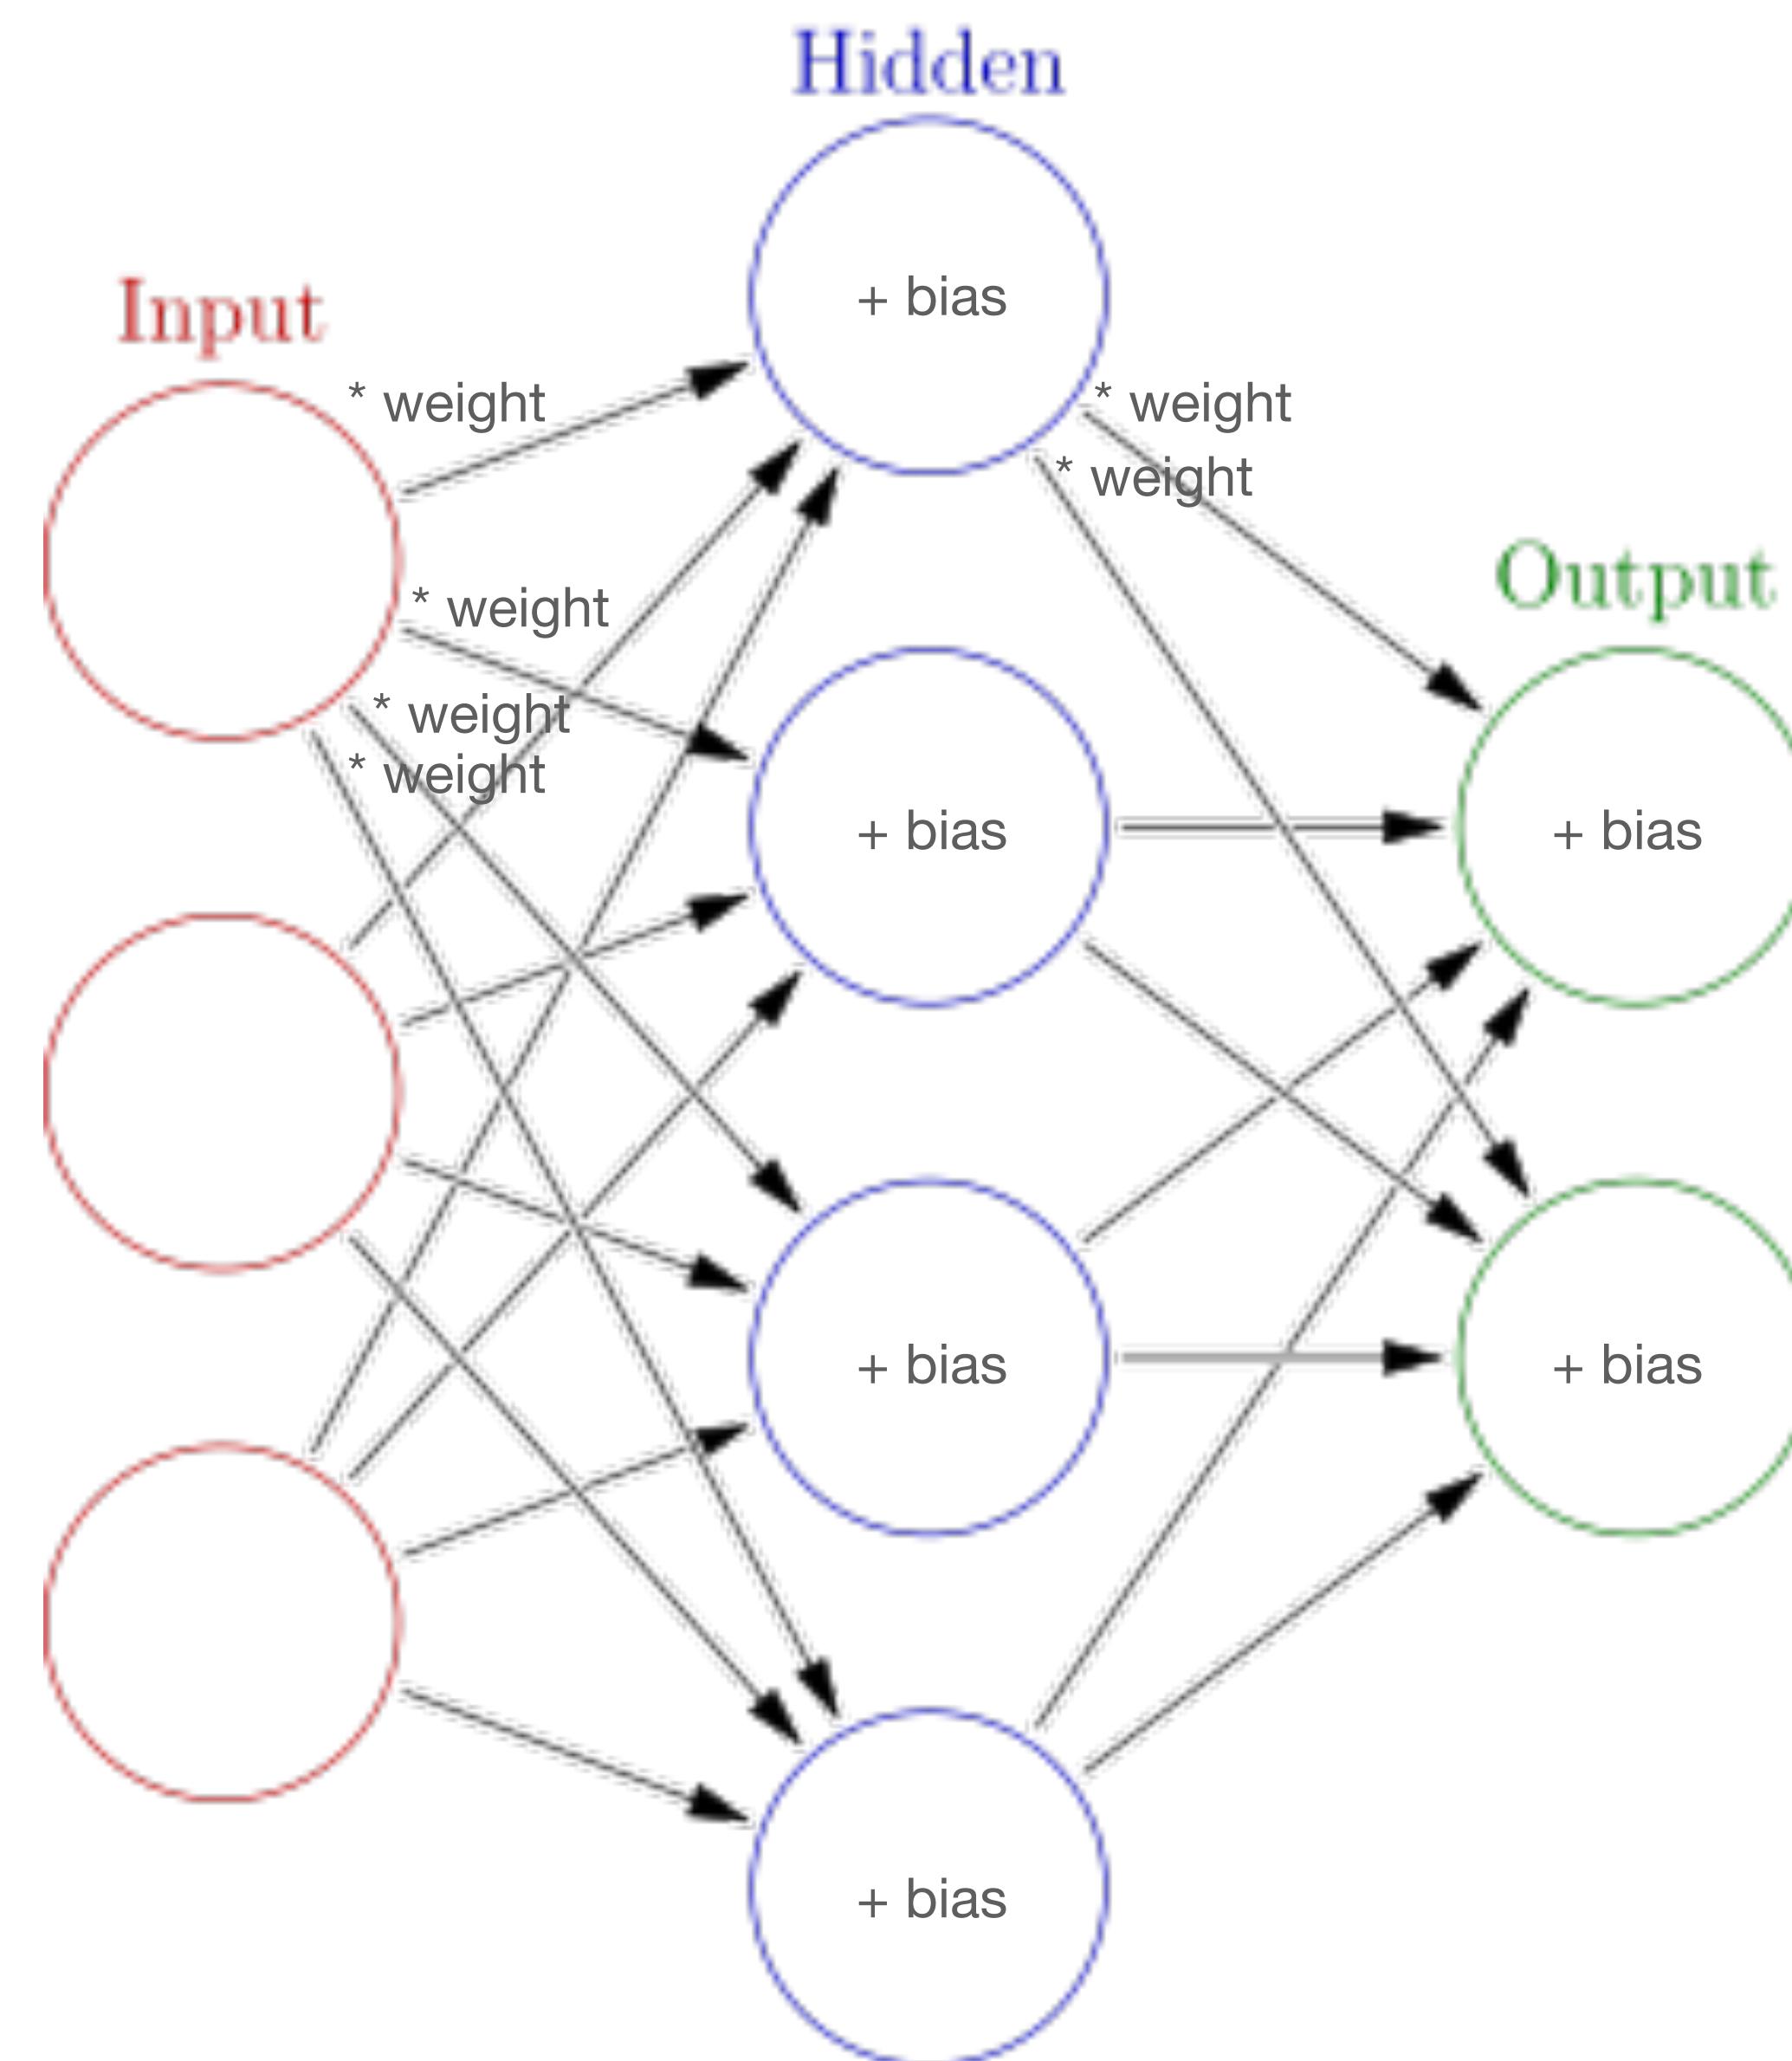
*(@hiddenlayers 5 8 8)*



(@hiddenlayers 10 5)

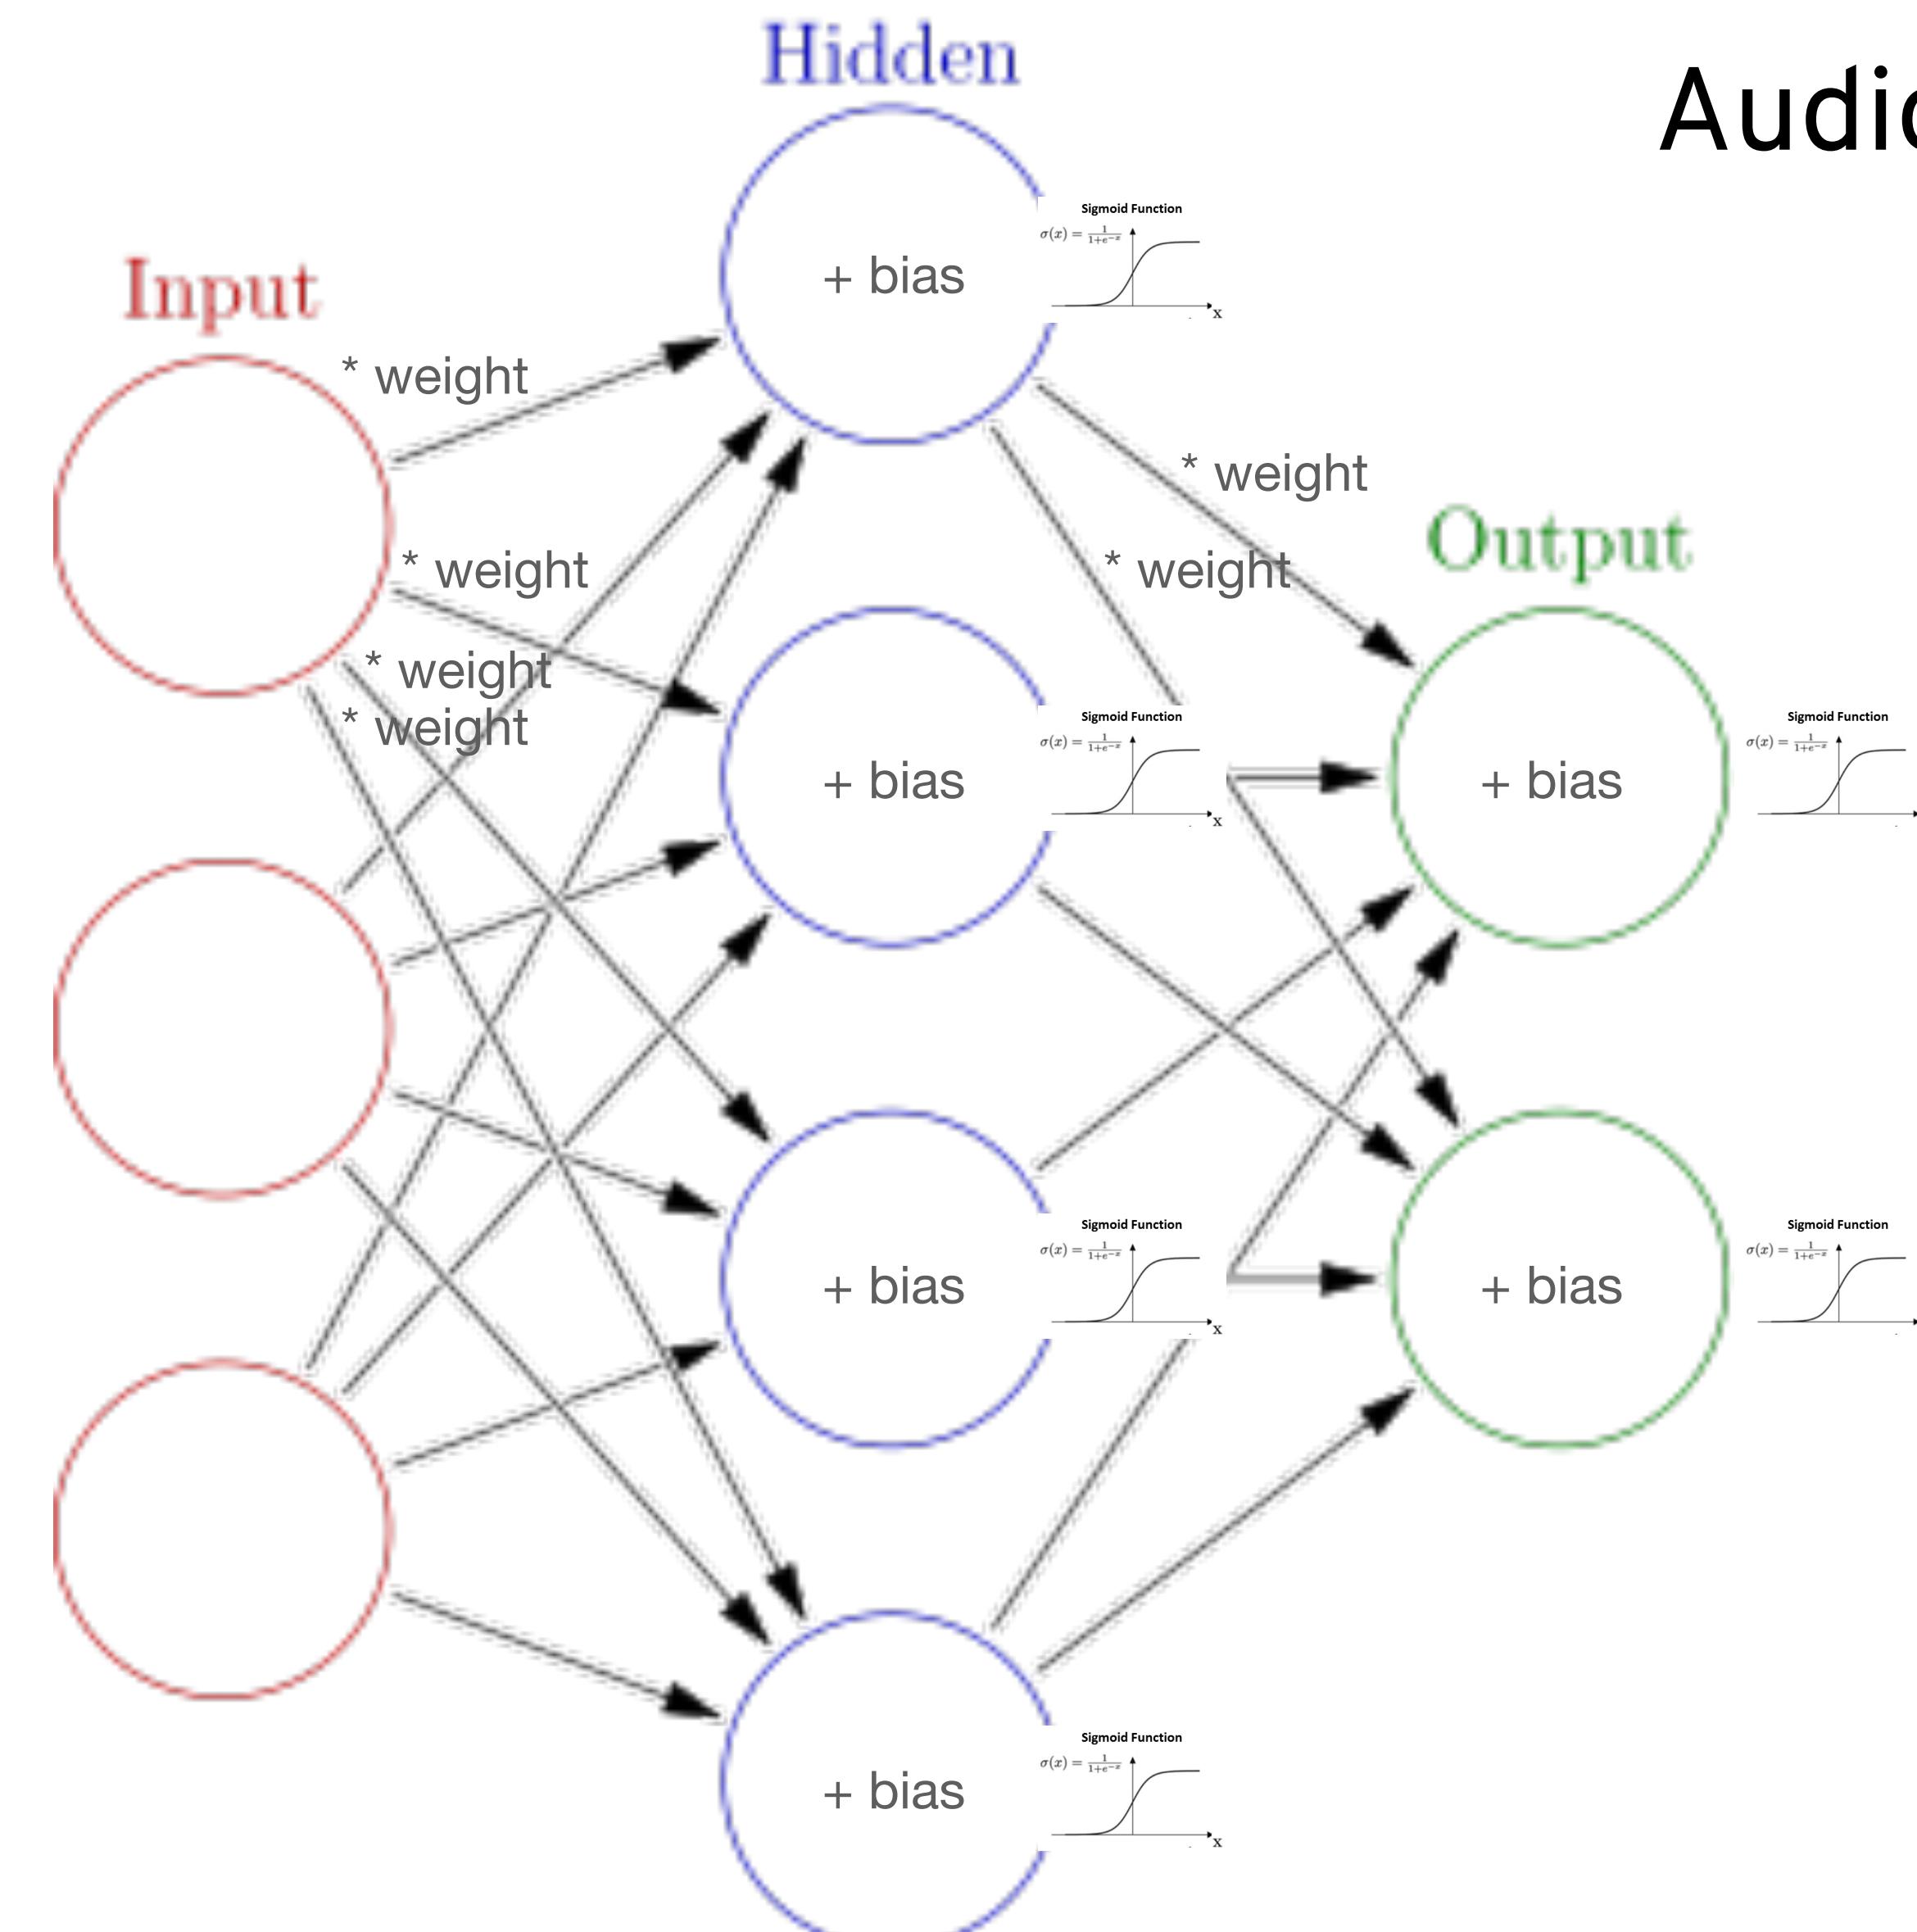
# Internal Parameters

weights & biases



@activation  
@outputactivation

# Audio Mixer Metaphor

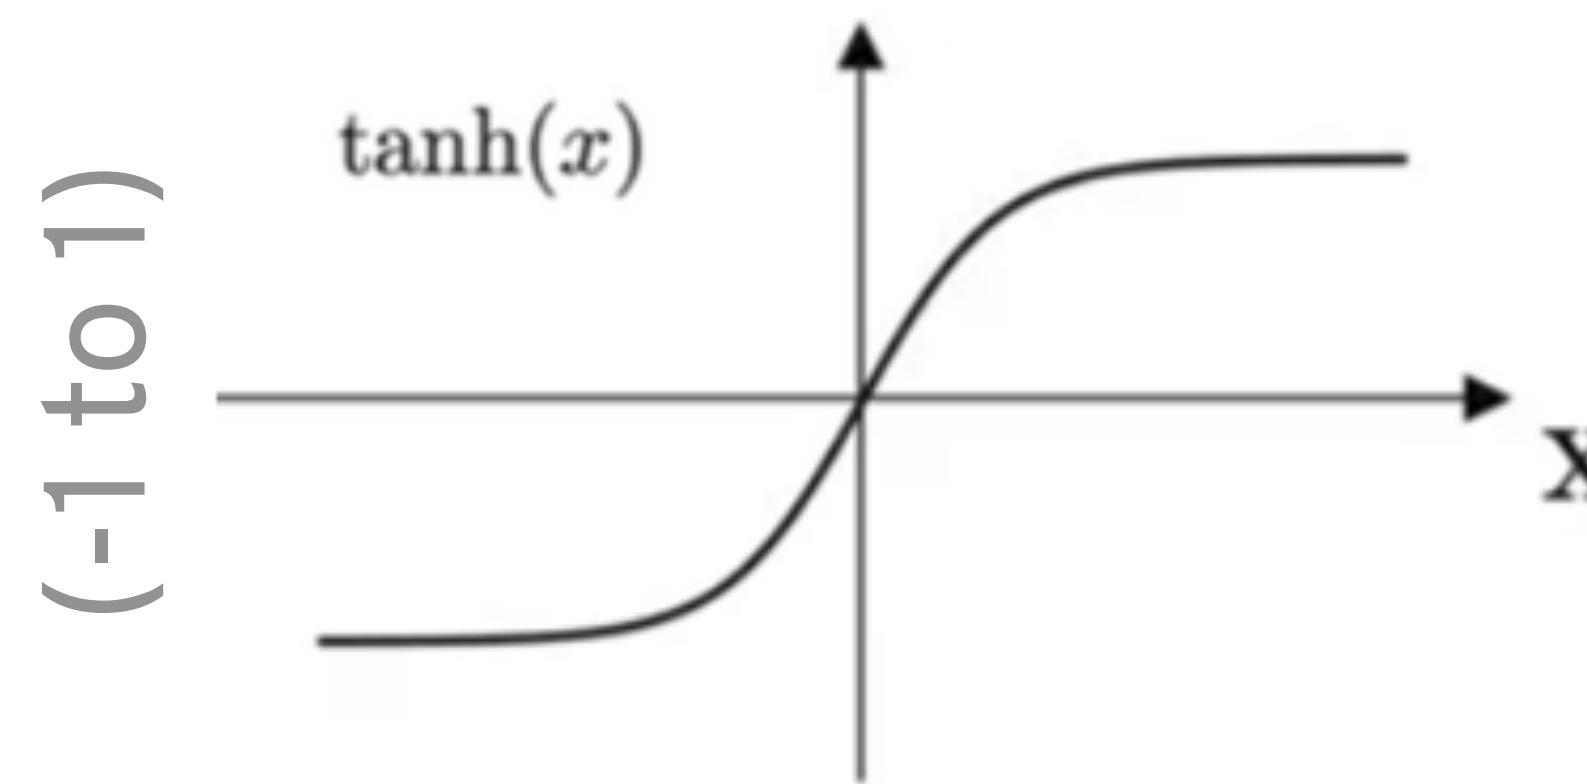


will reset network

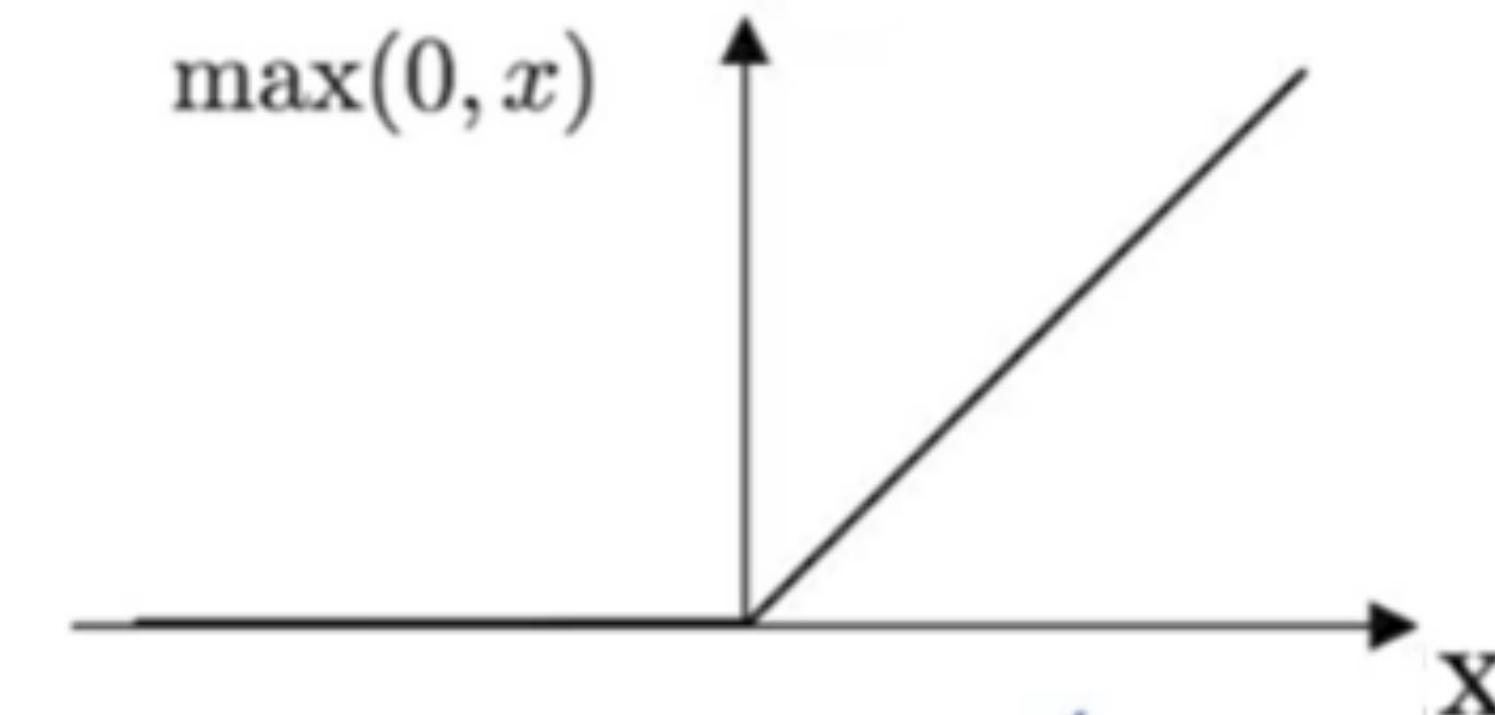
@activation

@outputactivation

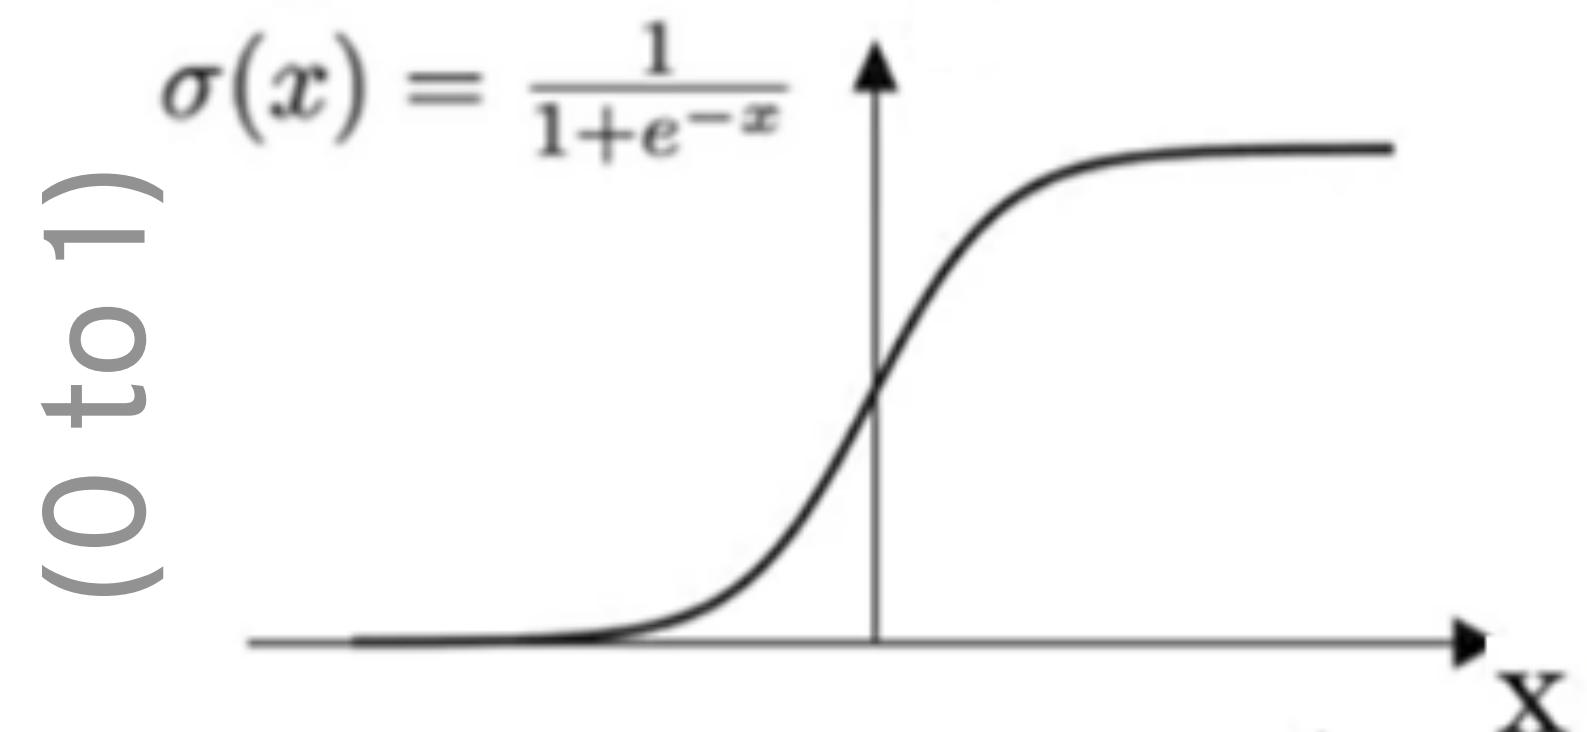
(3) Hyper Tangent Function



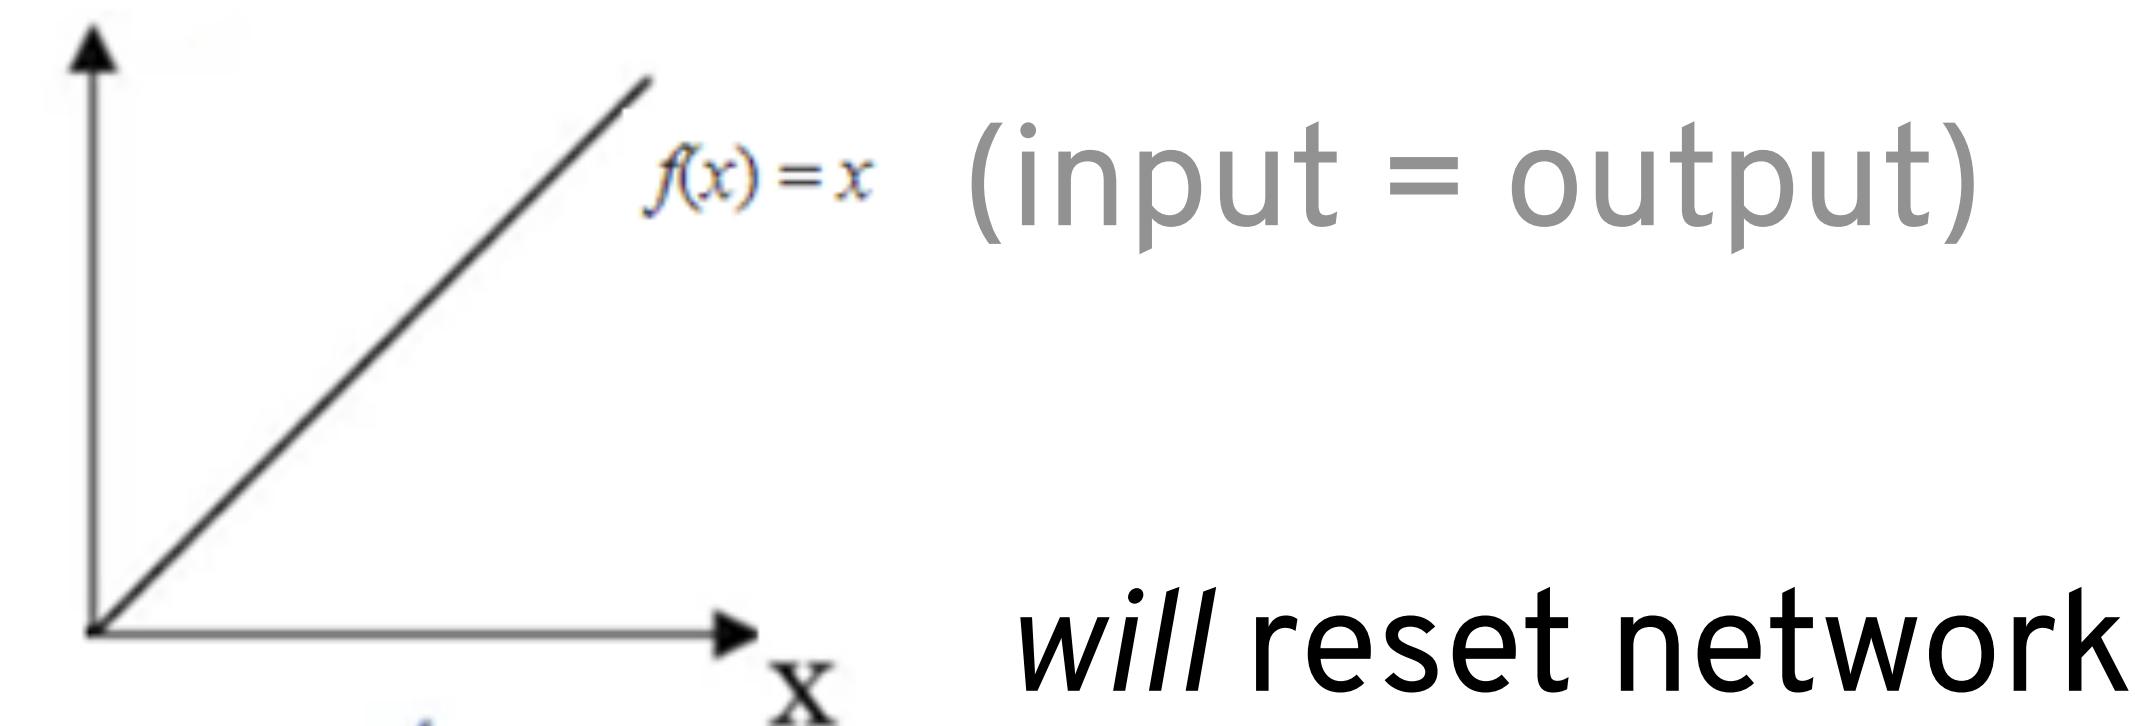
(2) ReLU Function



(1) Sigmoid Function



(0) Identity Function



`@learnrate`

a scalar for indicating how much the neural network should adjust its internal parameters during training (during back-propagation)—starting around 0.1 and decreasing over time is usually a good start

`@maxiter`

number of times to use all the examples on each “fit” message—also called an “epoch”

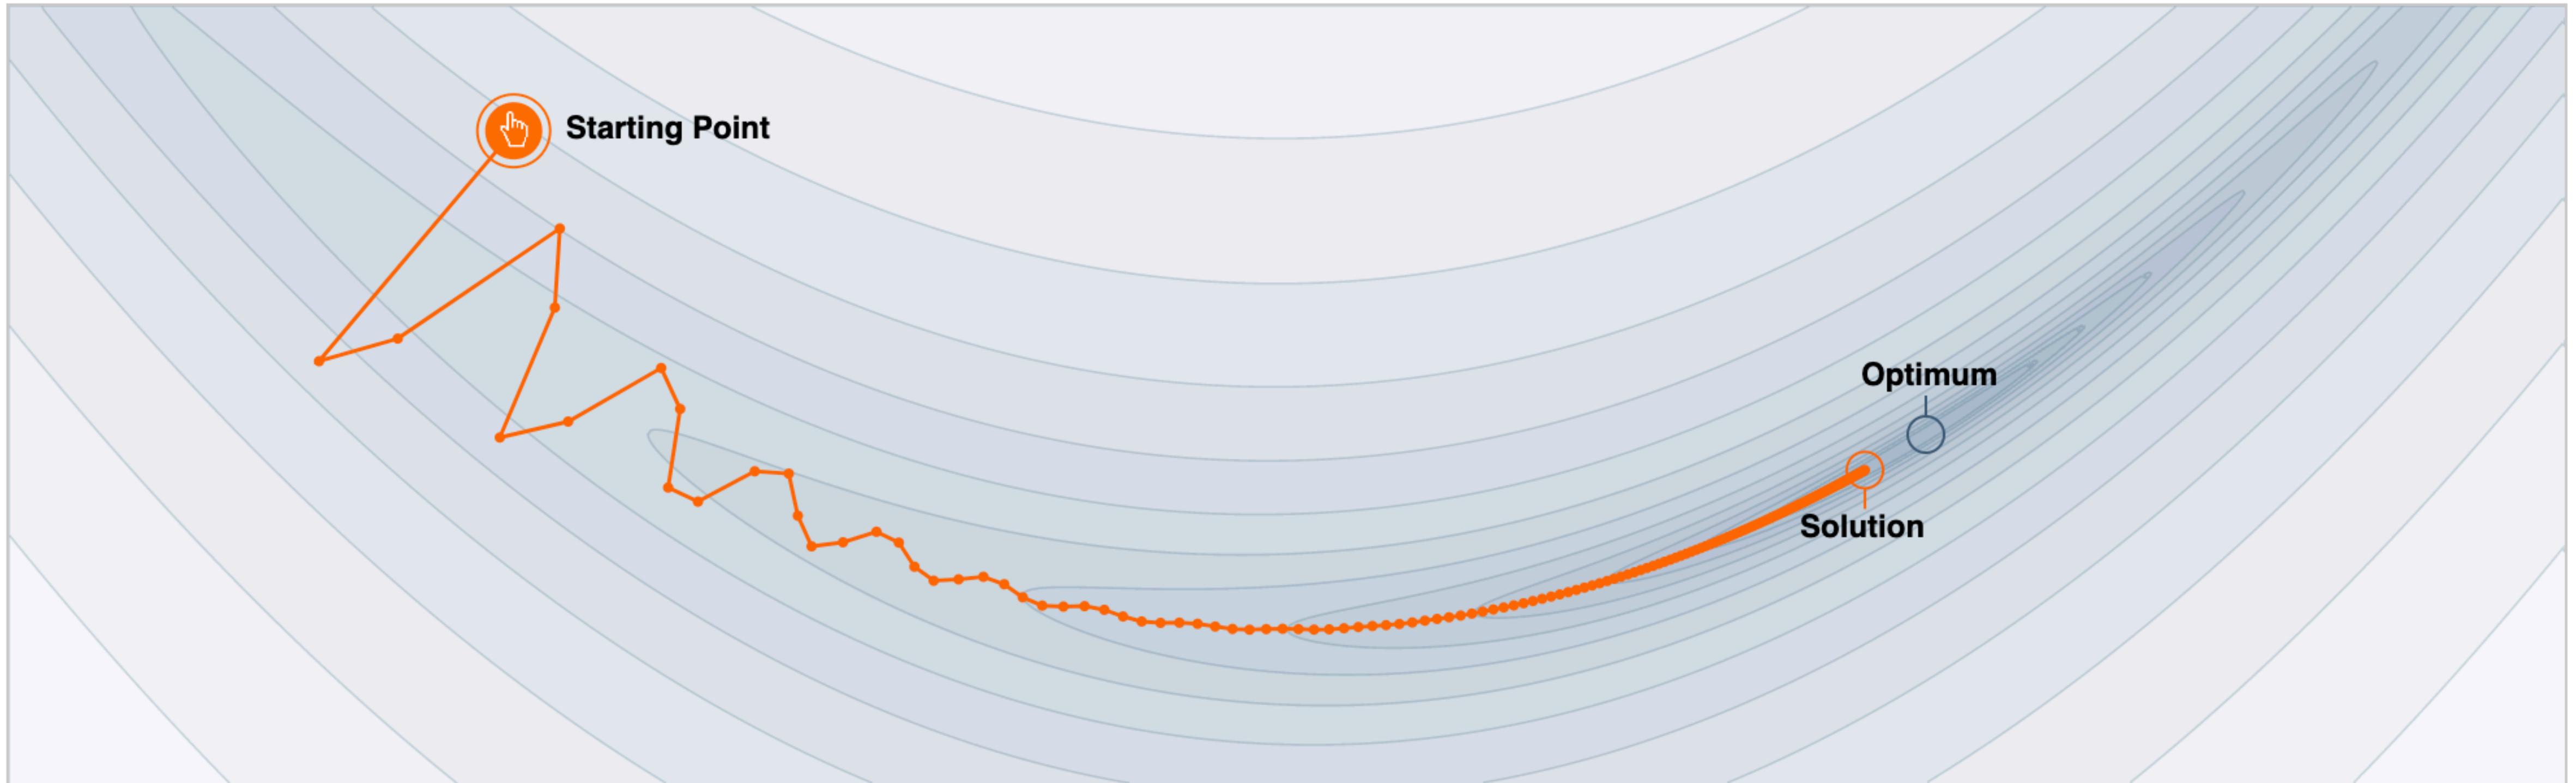
`@batchsize`

number of input examples feed-forward before checking their output pairs and adjusting

`@validation`

percentage (as a decimal) of training data to hold back (not train on) then test with—when there’s *lots* of data, may be helpful to prevent overfitting

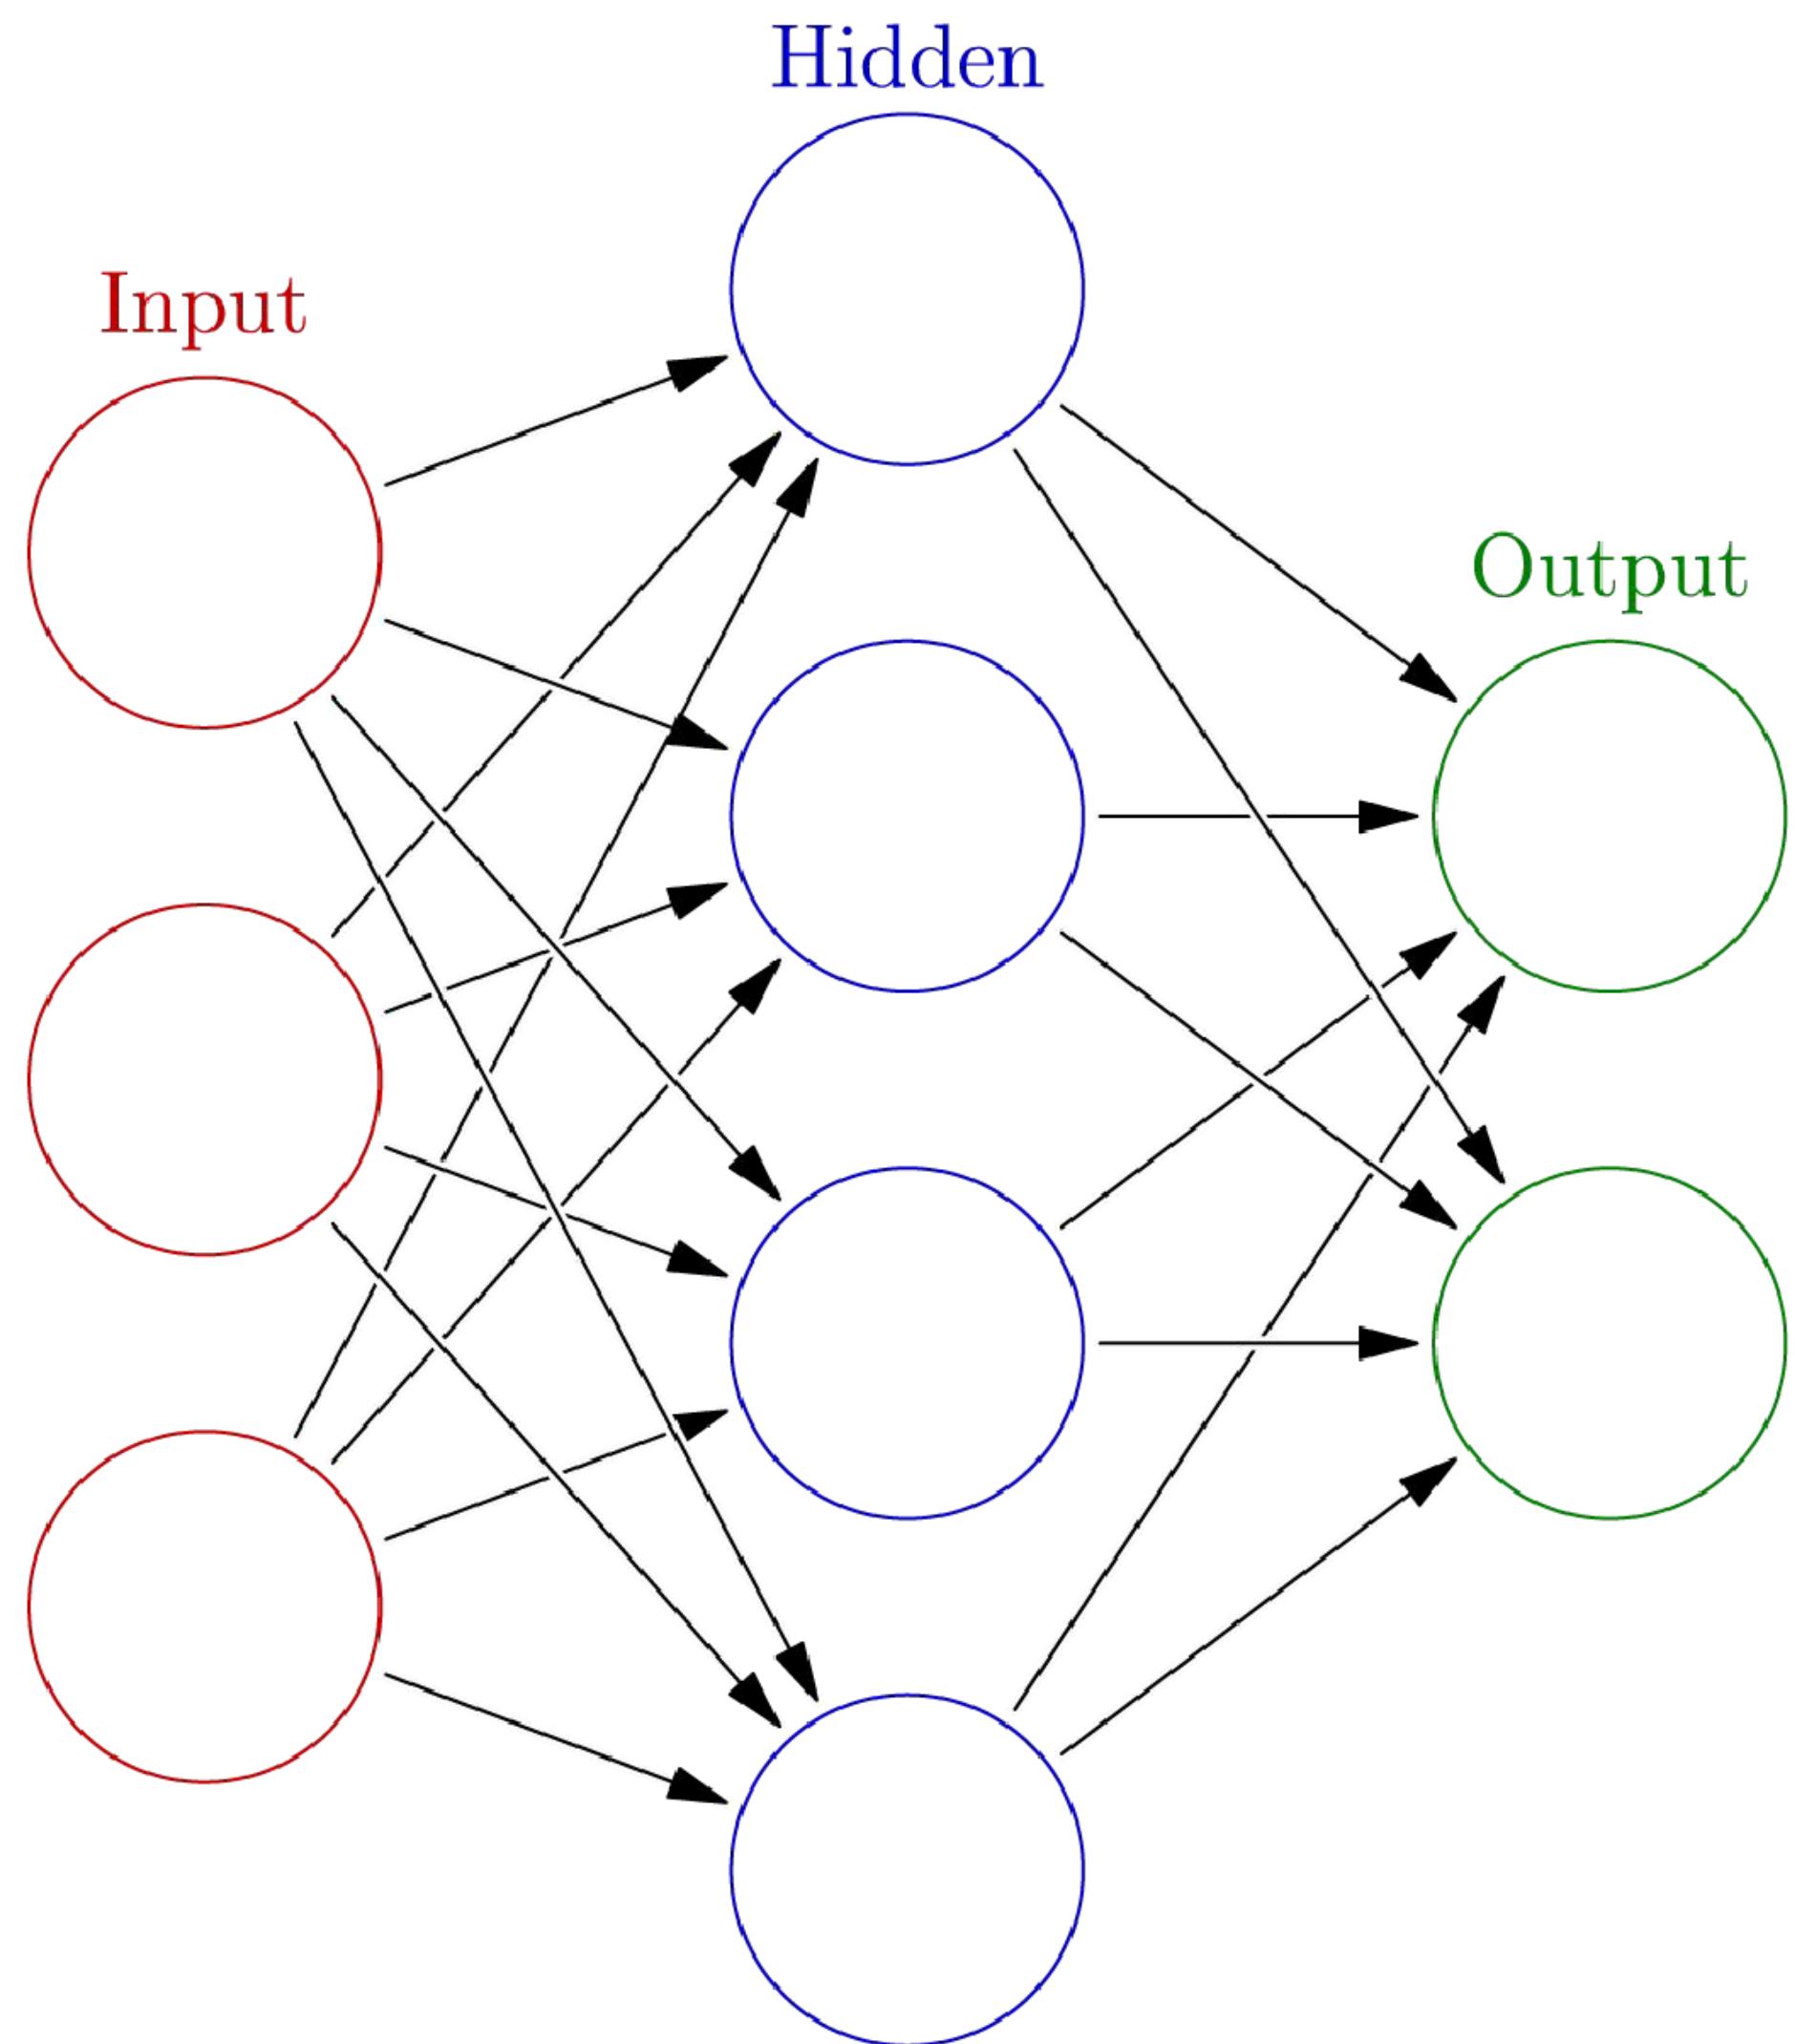
@momentum apply smoothing adjustments



# *Regression with Control Data*



# *Neural Network*



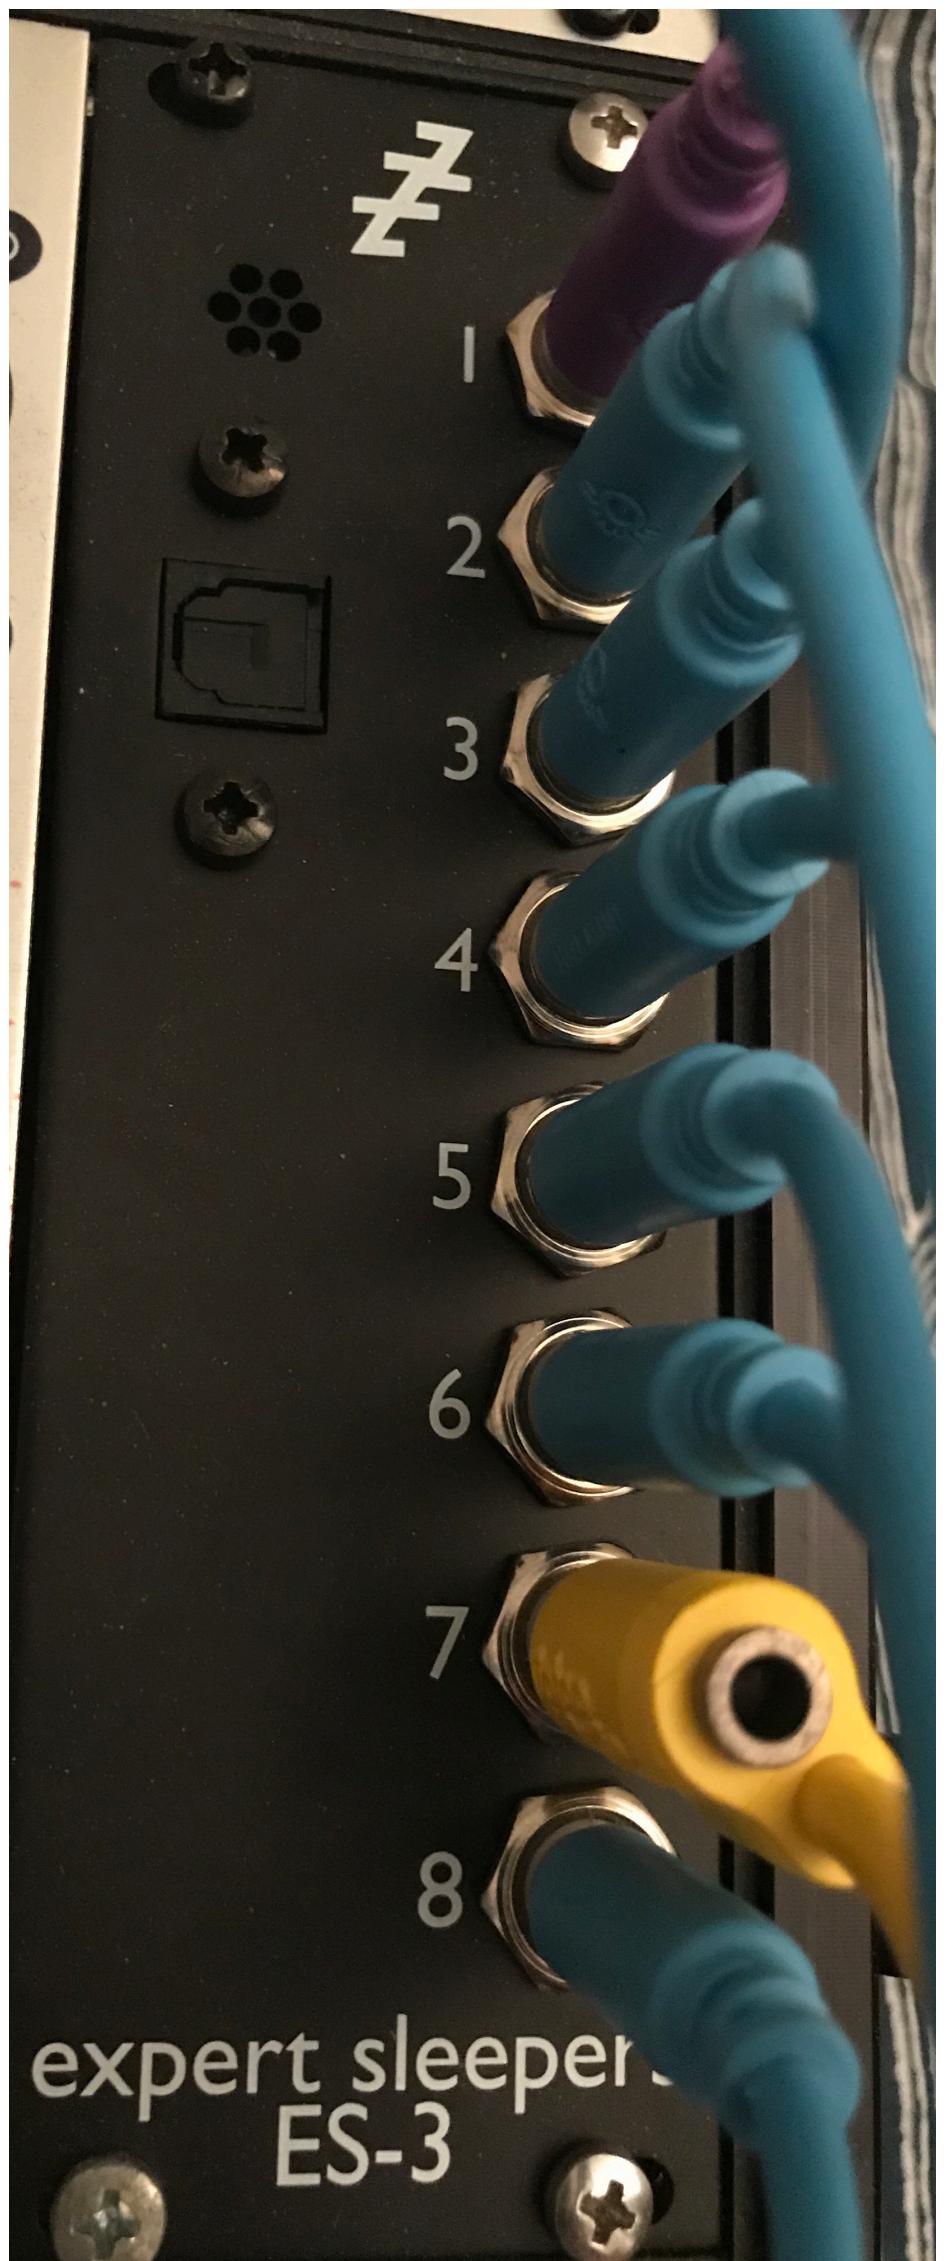
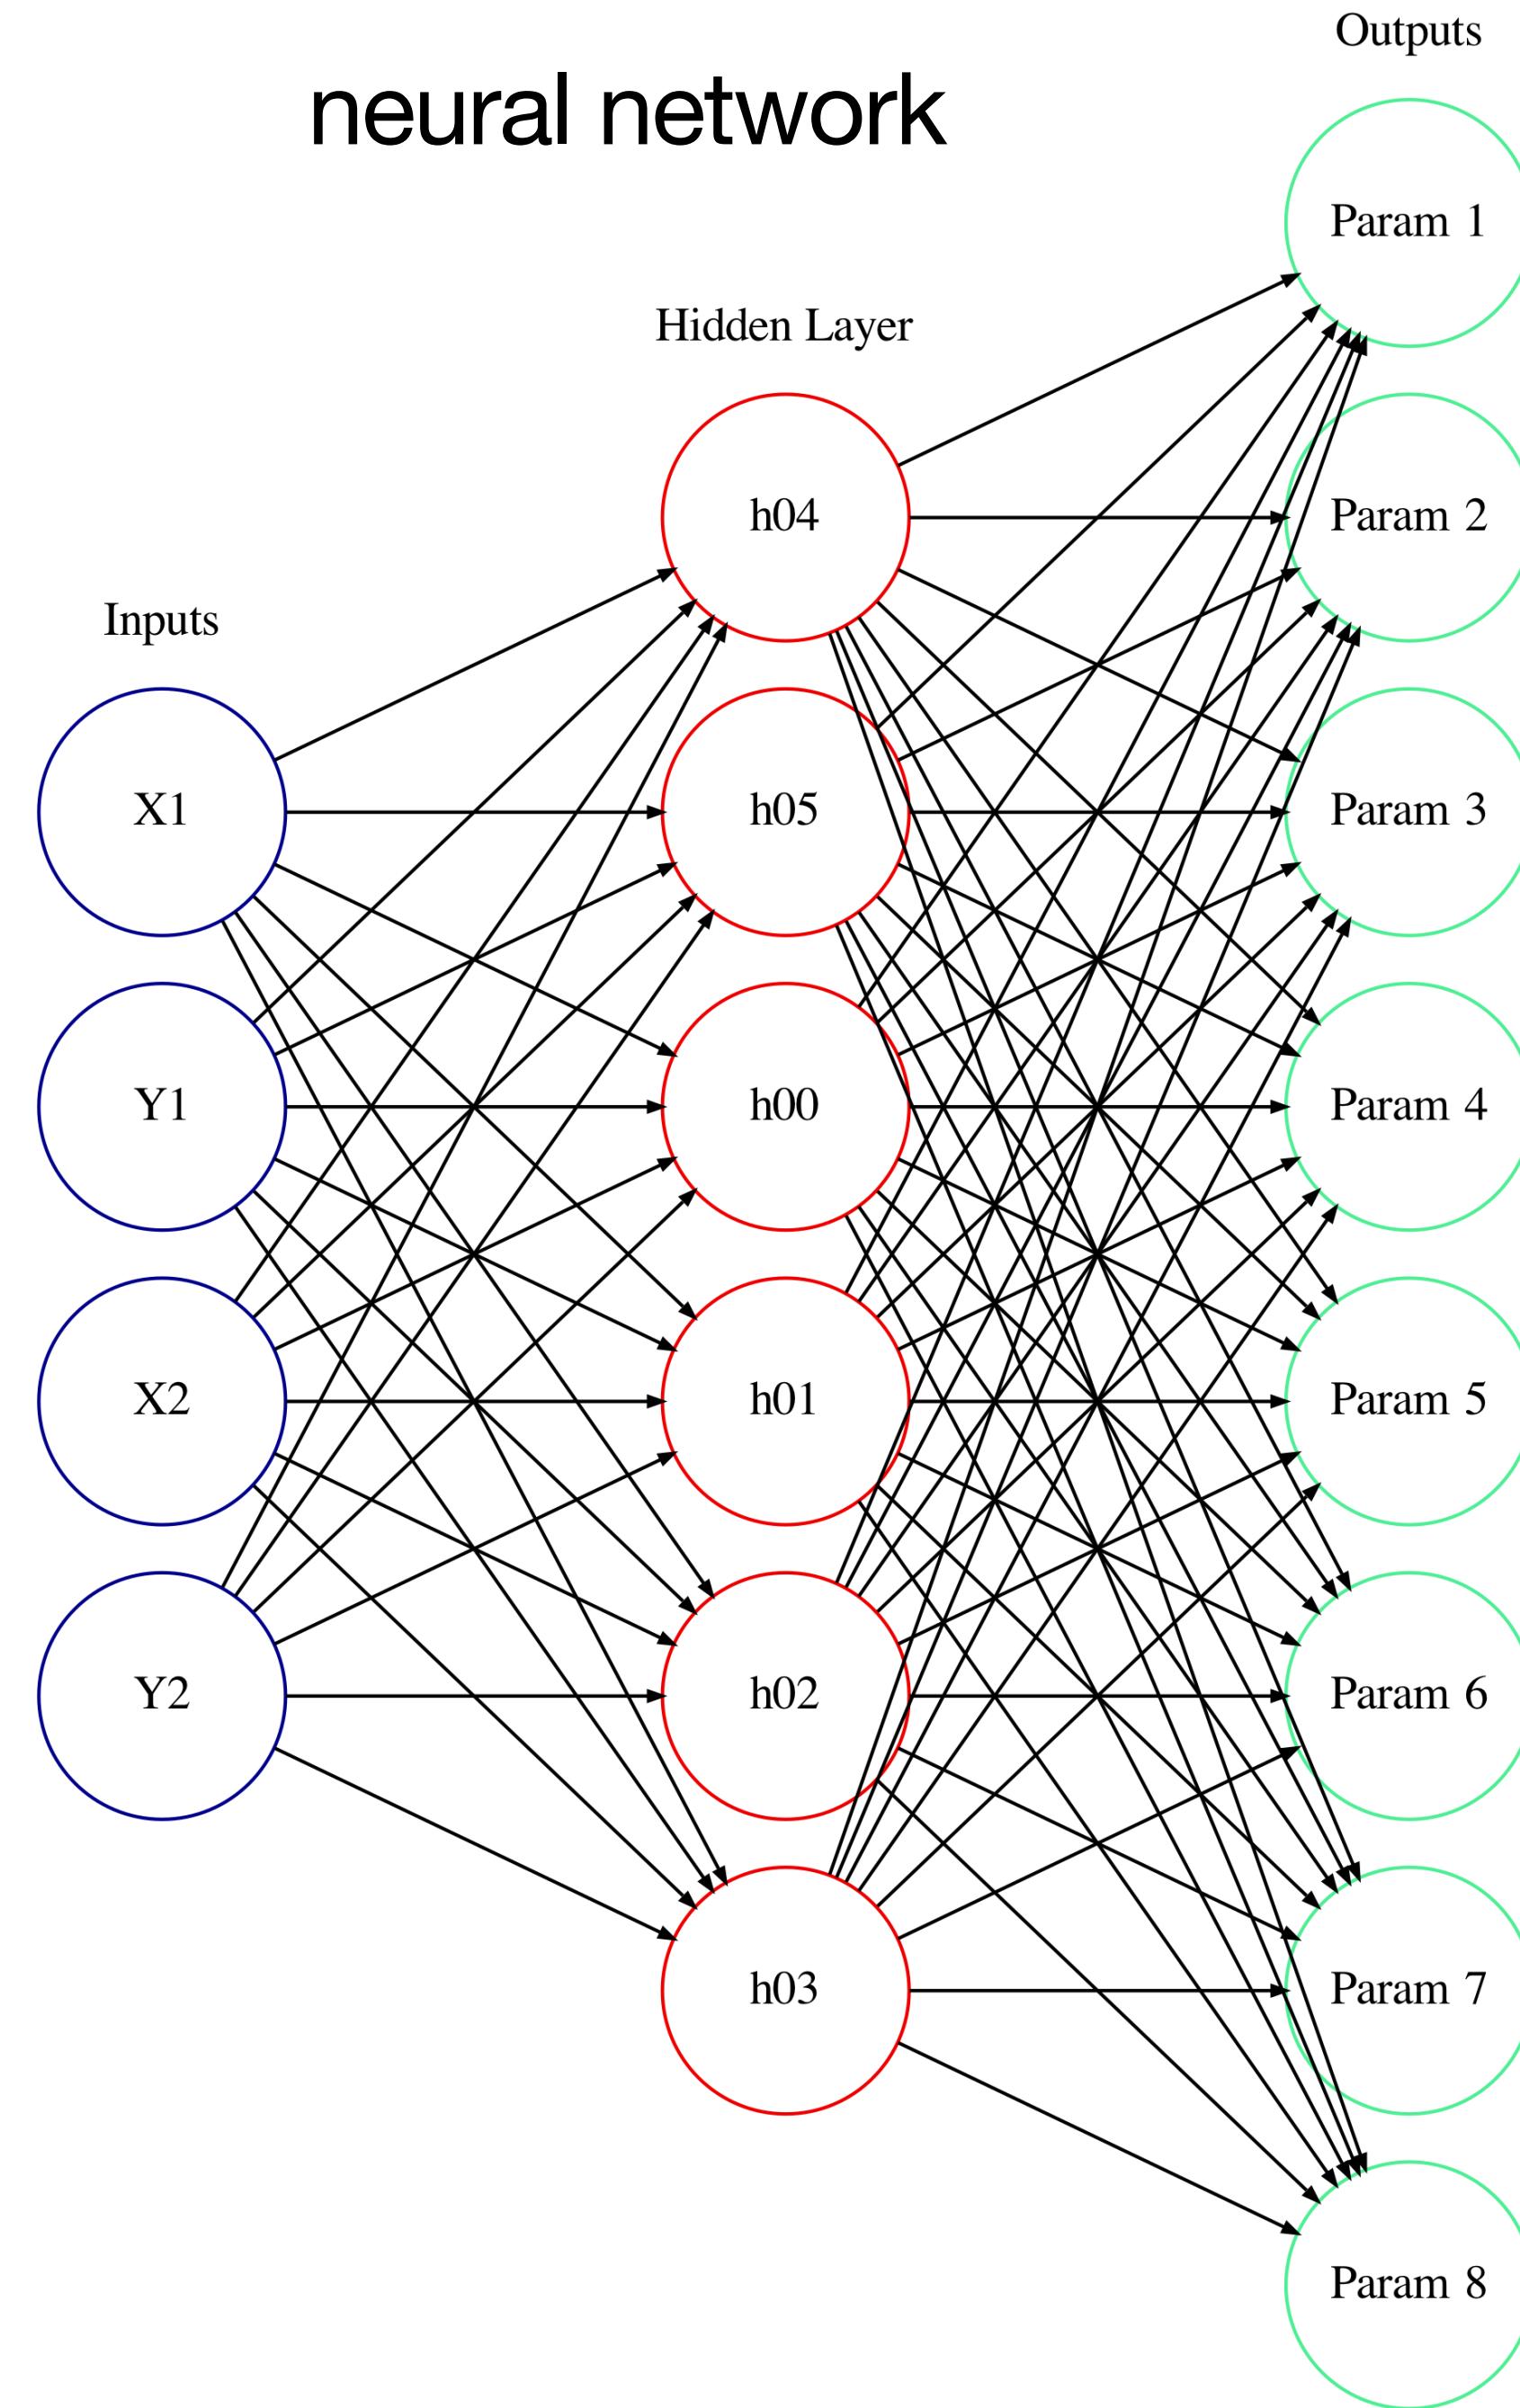
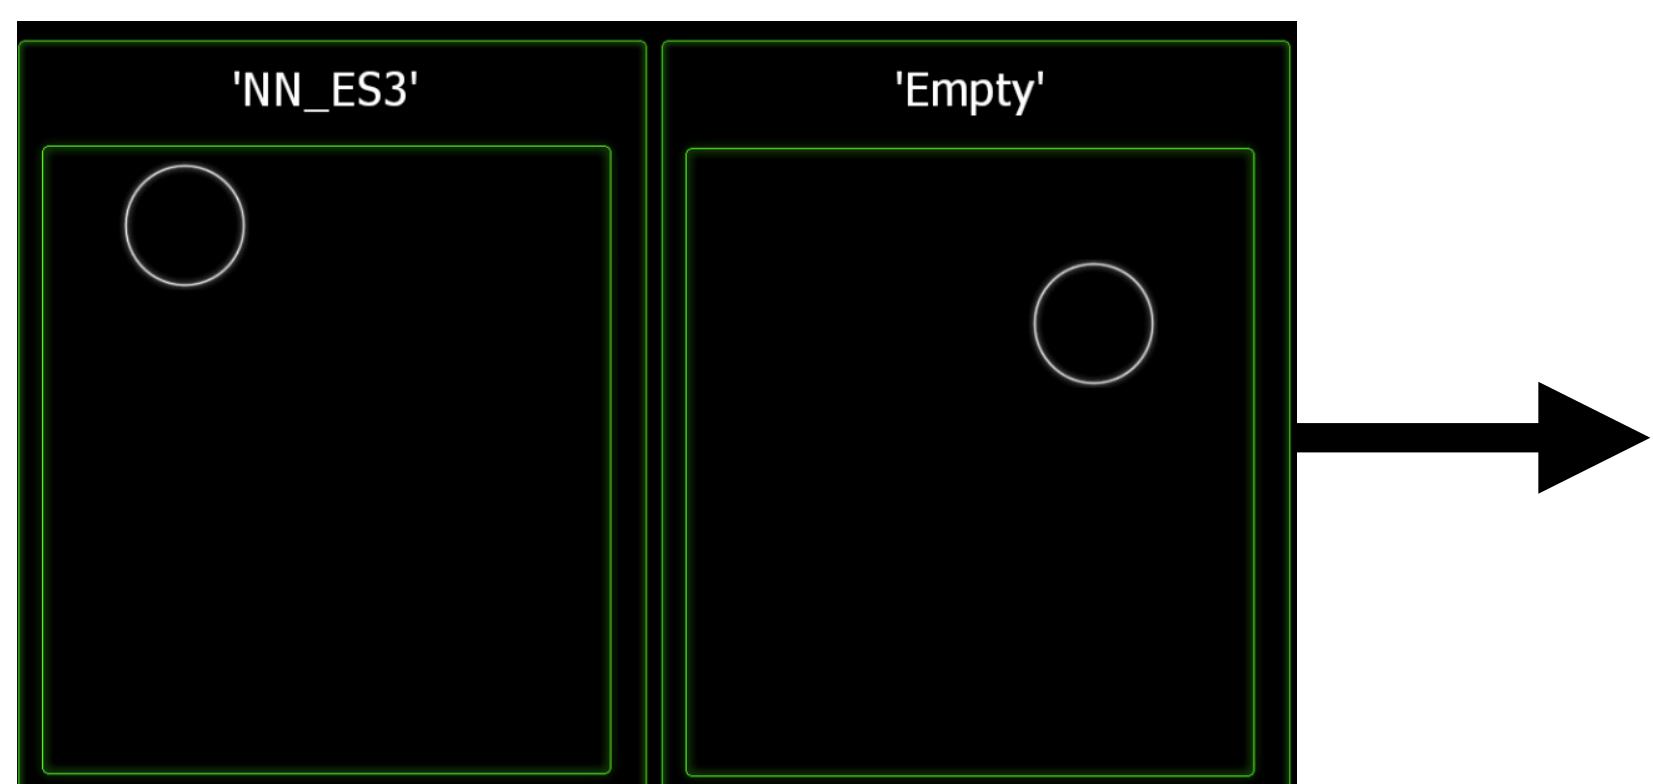
**Regression:**  
a neural network predicts  
**output values (to anything)**  
**from**  
**input values (from anything)**







# neural network





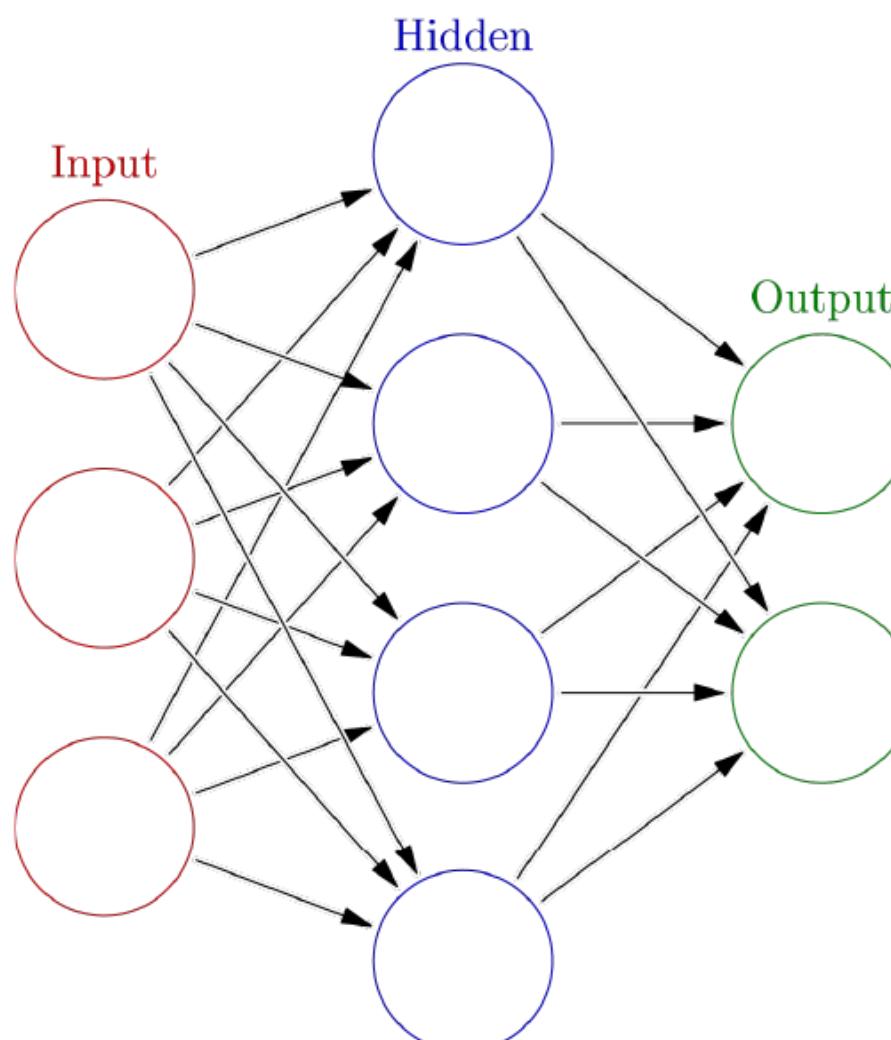
# *Neural Network Training a Regressor*

identifier

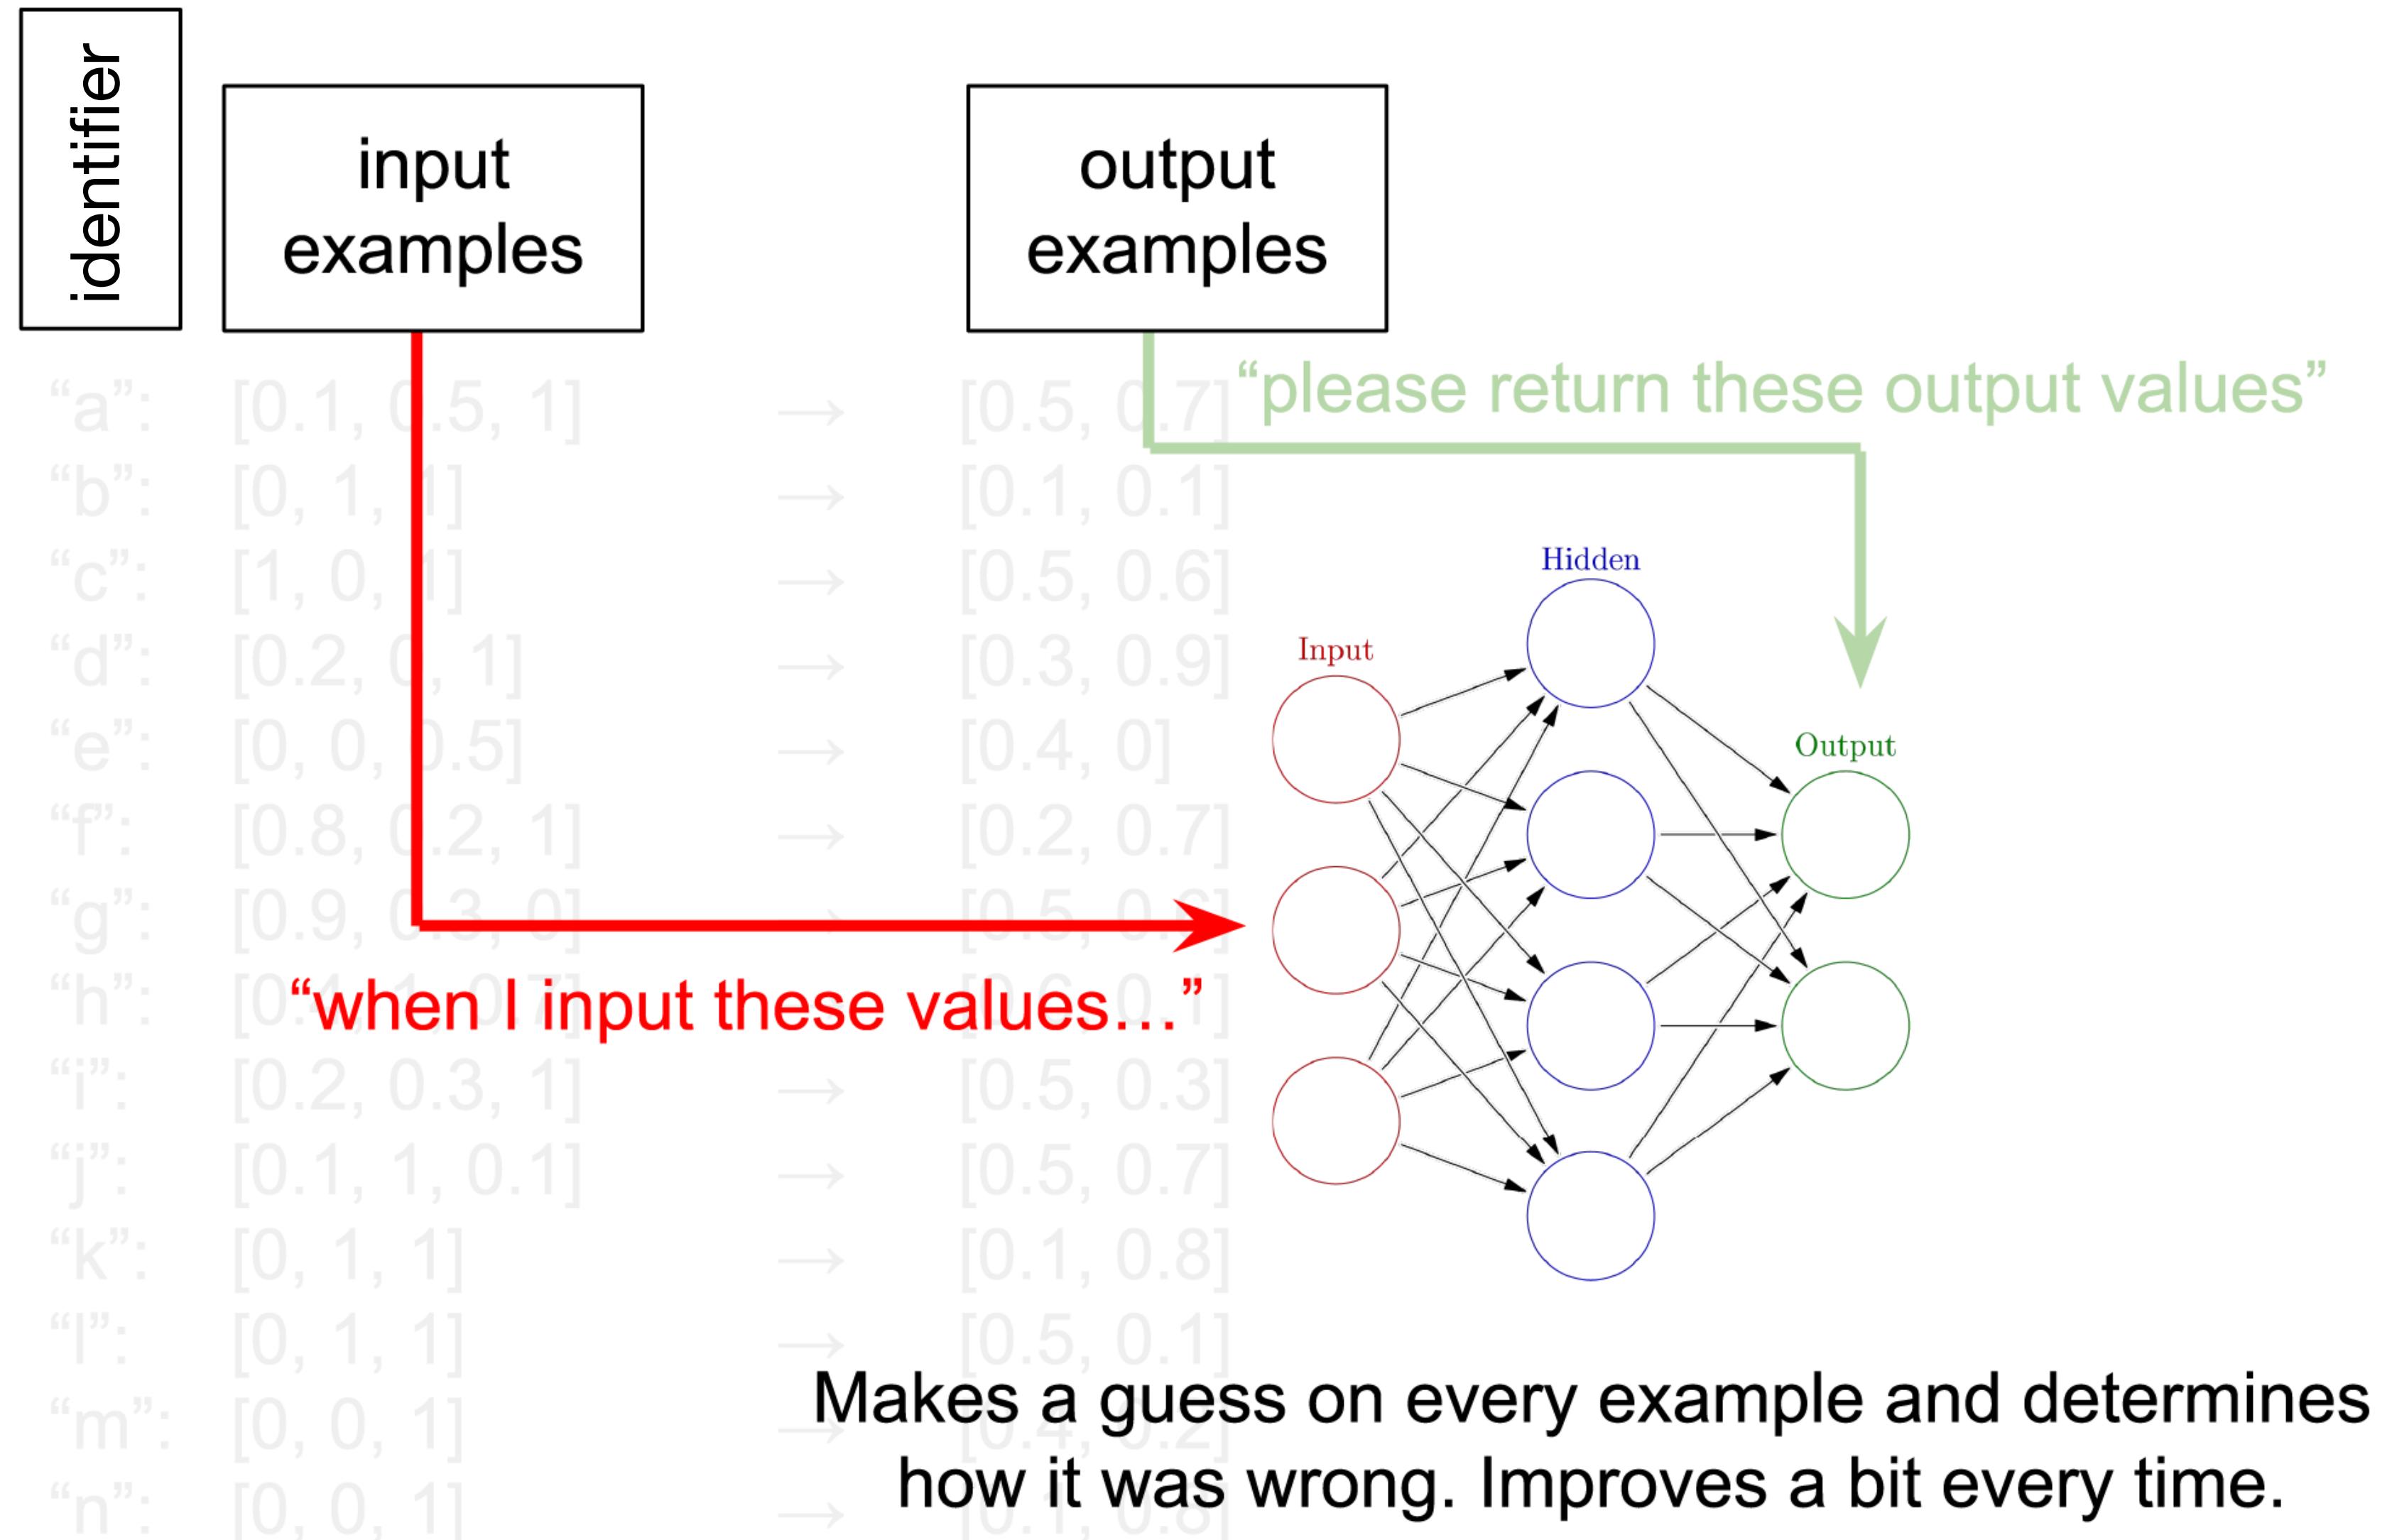
input  
examples

“a”:	[0.1, 0.5, 1]	→	[0.5, 0.7]
“b”:	[0, 1, 1]	→	[0.1, 0.1]
“c”:	[1, 0, 1]	→	[0.5, 0.6]
“d”:	[0.2, 0, 1]	→	[0.3, 0.9]
“e”:	[0, 0, 0.5]	→	[0.4, 0]
“f”:	[0.8, 0.2, 1]	→	[0.2, 0.7]
“g”:	[0.9, 0.3, 0]	→	[0.5, 0.6]
“h”:	[0.4, 1, 0.7]	→	[0.6, 0.1]
“i”:	[0.2, 0.3, 1]	→	[0.5, 0.3]
“j”:	[0.1, 1, 0.1]	→	[0.5, 0.7]
“k”:	[0, 1, 1]	→	[0.1, 0.8]
“l”:	[0, 1, 1]	→	[0.5, 0.1]
“m”:	[0, 0, 1]	→	[0.4, 0.2]
“n”:	[0, 0, 1]	→	[0.1, 0.8]

output  
examples



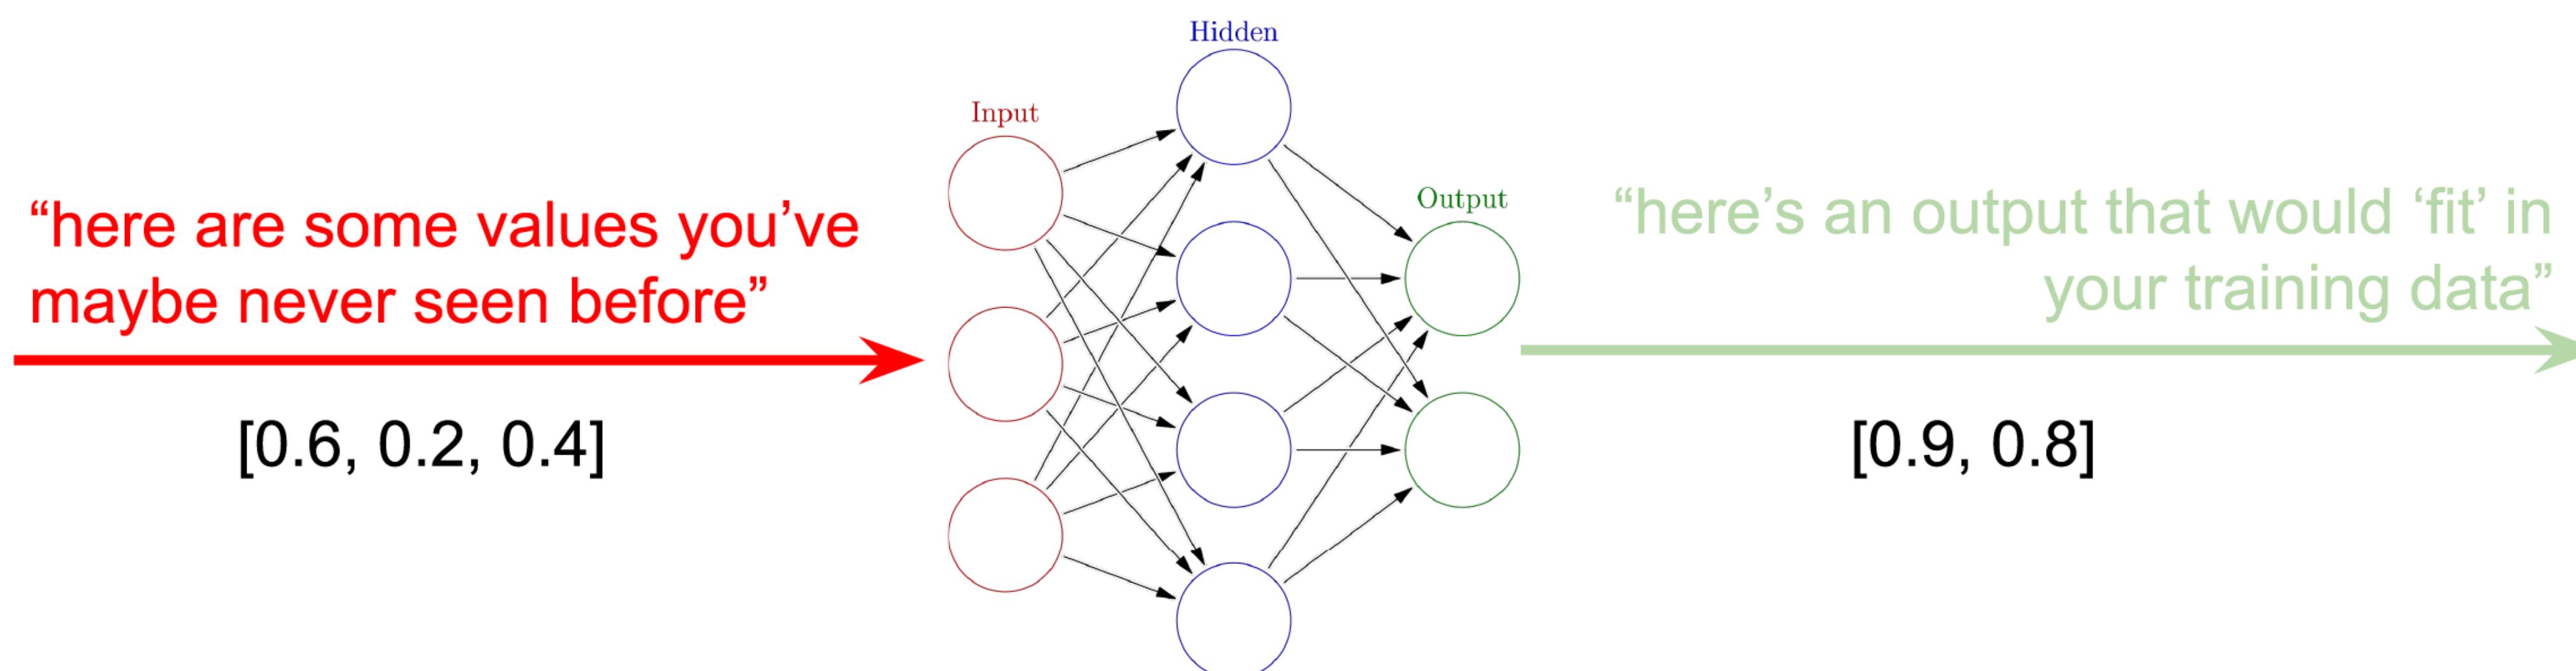
# *Neural Network Training a Regressor*



# Neural Network *Predicting with Regression*

input  
examples

output  
examples



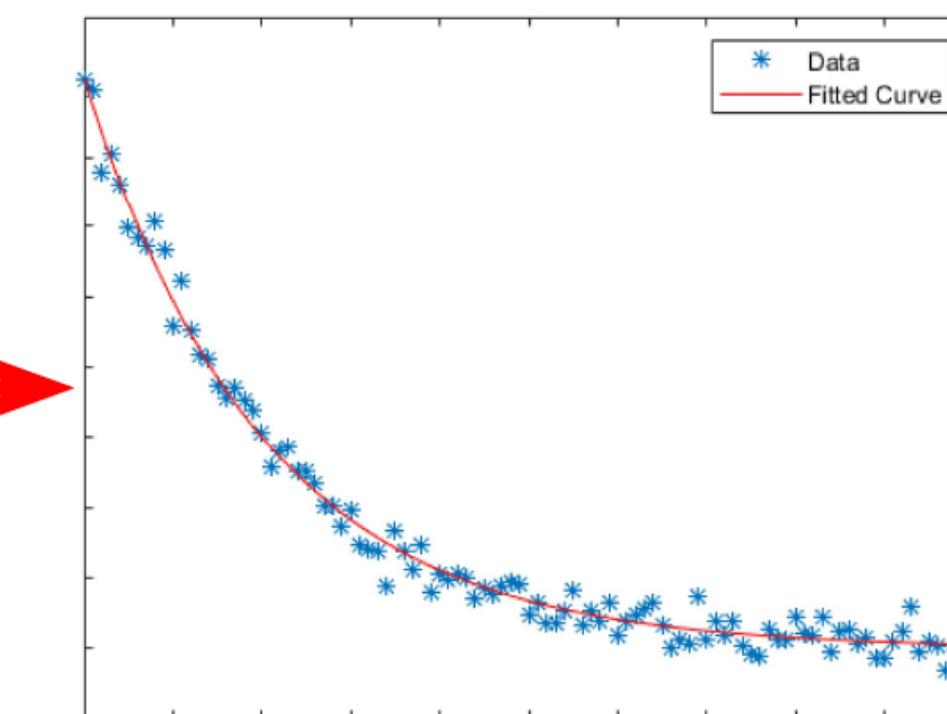
# Neural Network *Predicting with Regression*

input  
examples

output  
examples

“here are some values you’ve  
maybe never seen before”

[0.6, 0.2, 0.4]



“here’s an output that would ‘fit’ in  
your training data”

[0.9, 0.8]

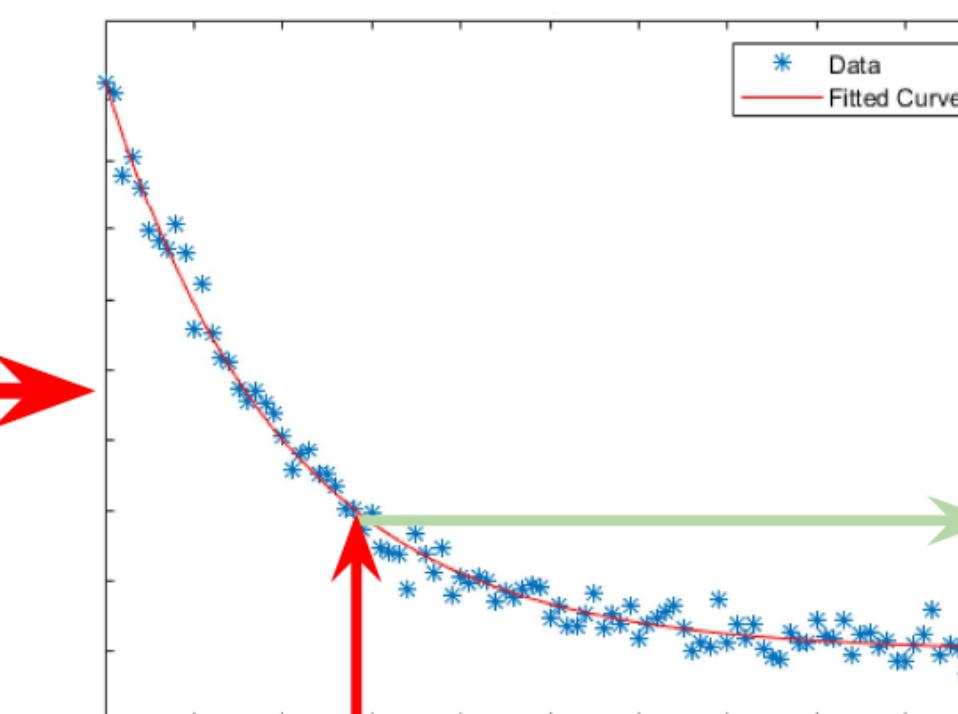
# Neural Network *Predicting with Regression*

input  
examples

output  
examples

“here are some values you’ve  
maybe never seen before”

[0.6, 0.2, 0.4]

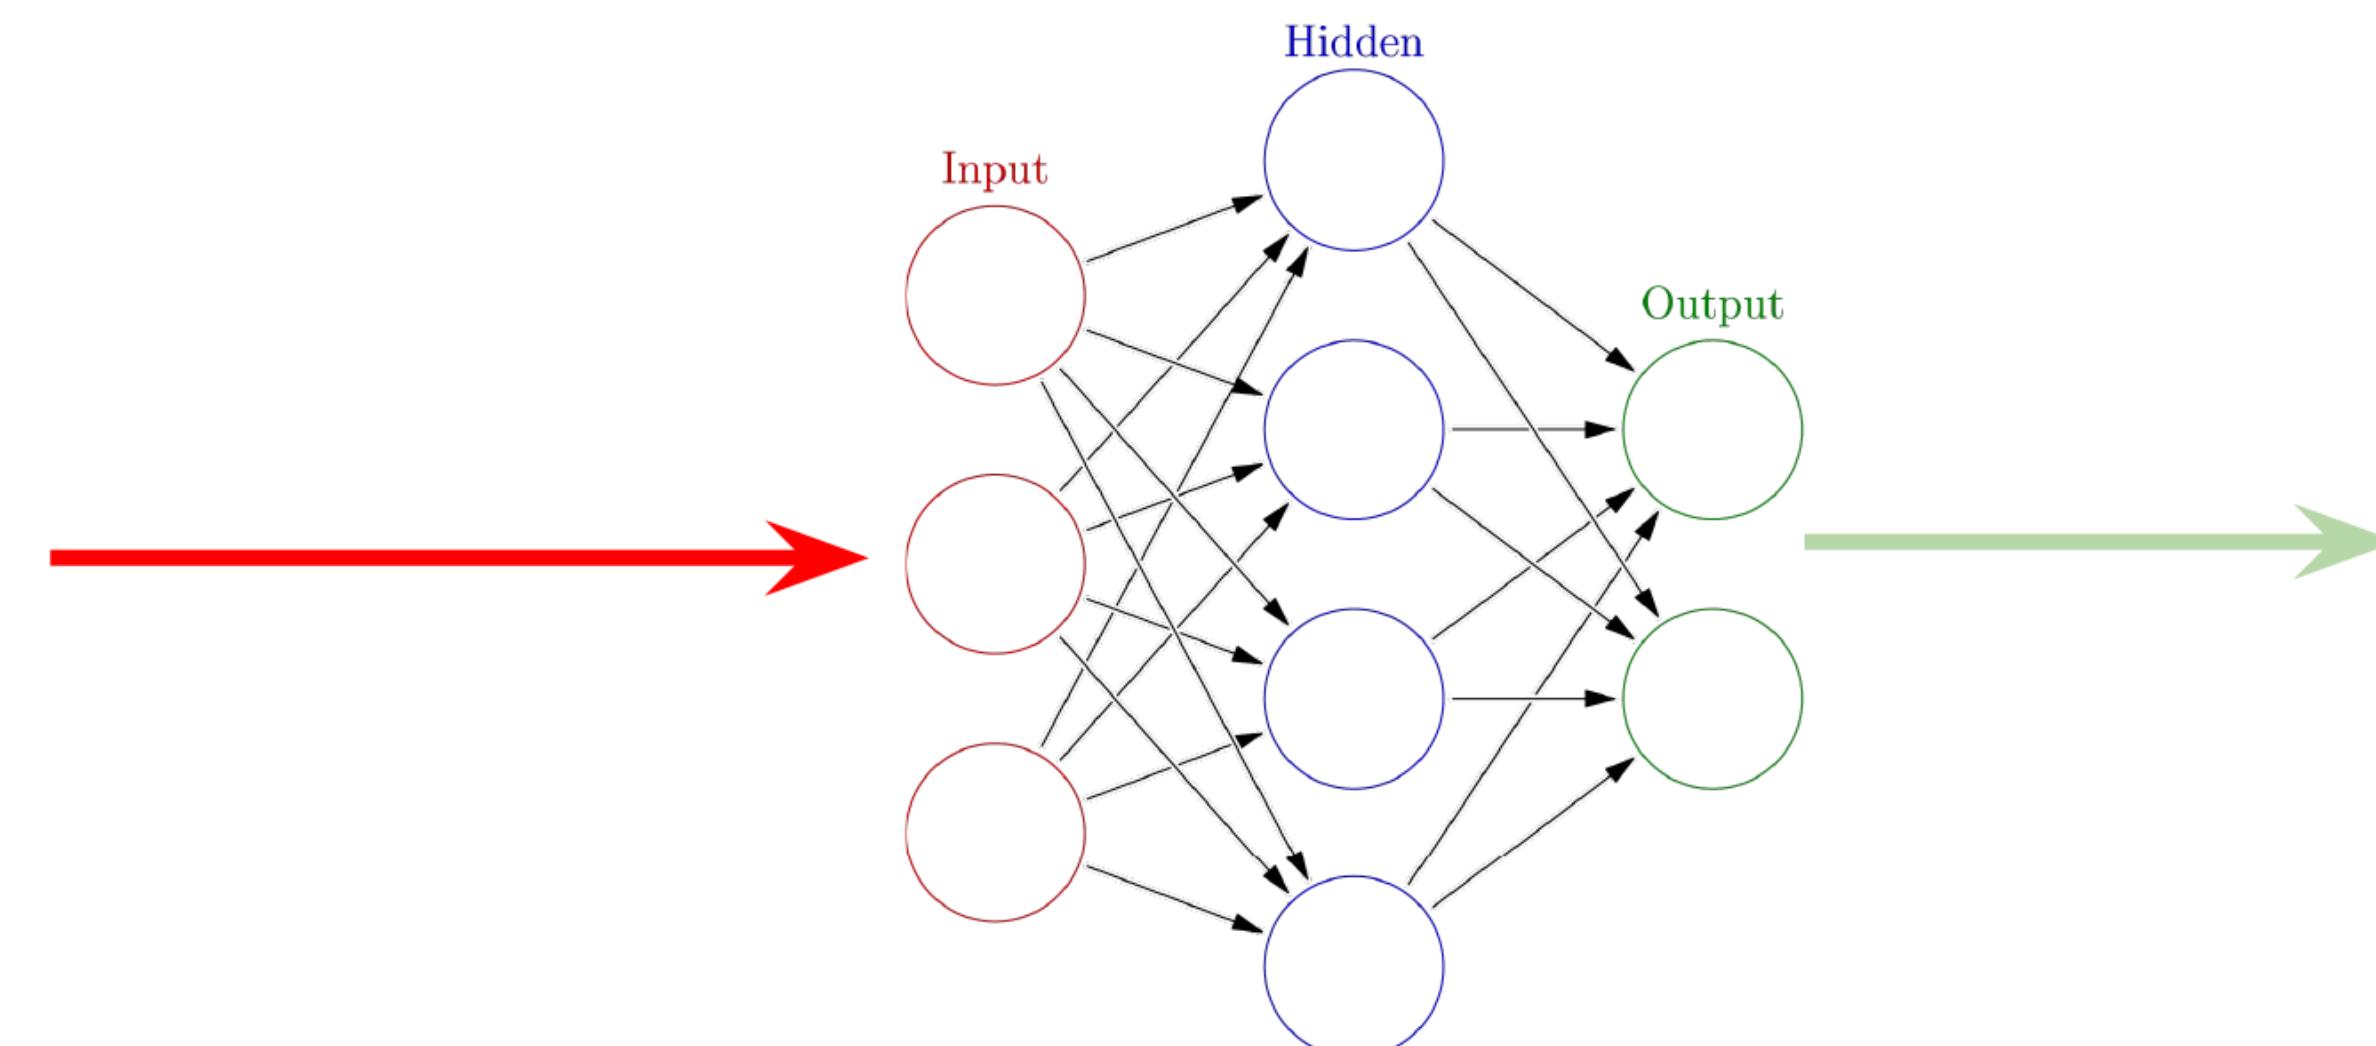


“here’s an output that would ‘fit’ in  
your training data”

[0.9, 0.8]

# Neural Network Predicting with Regression

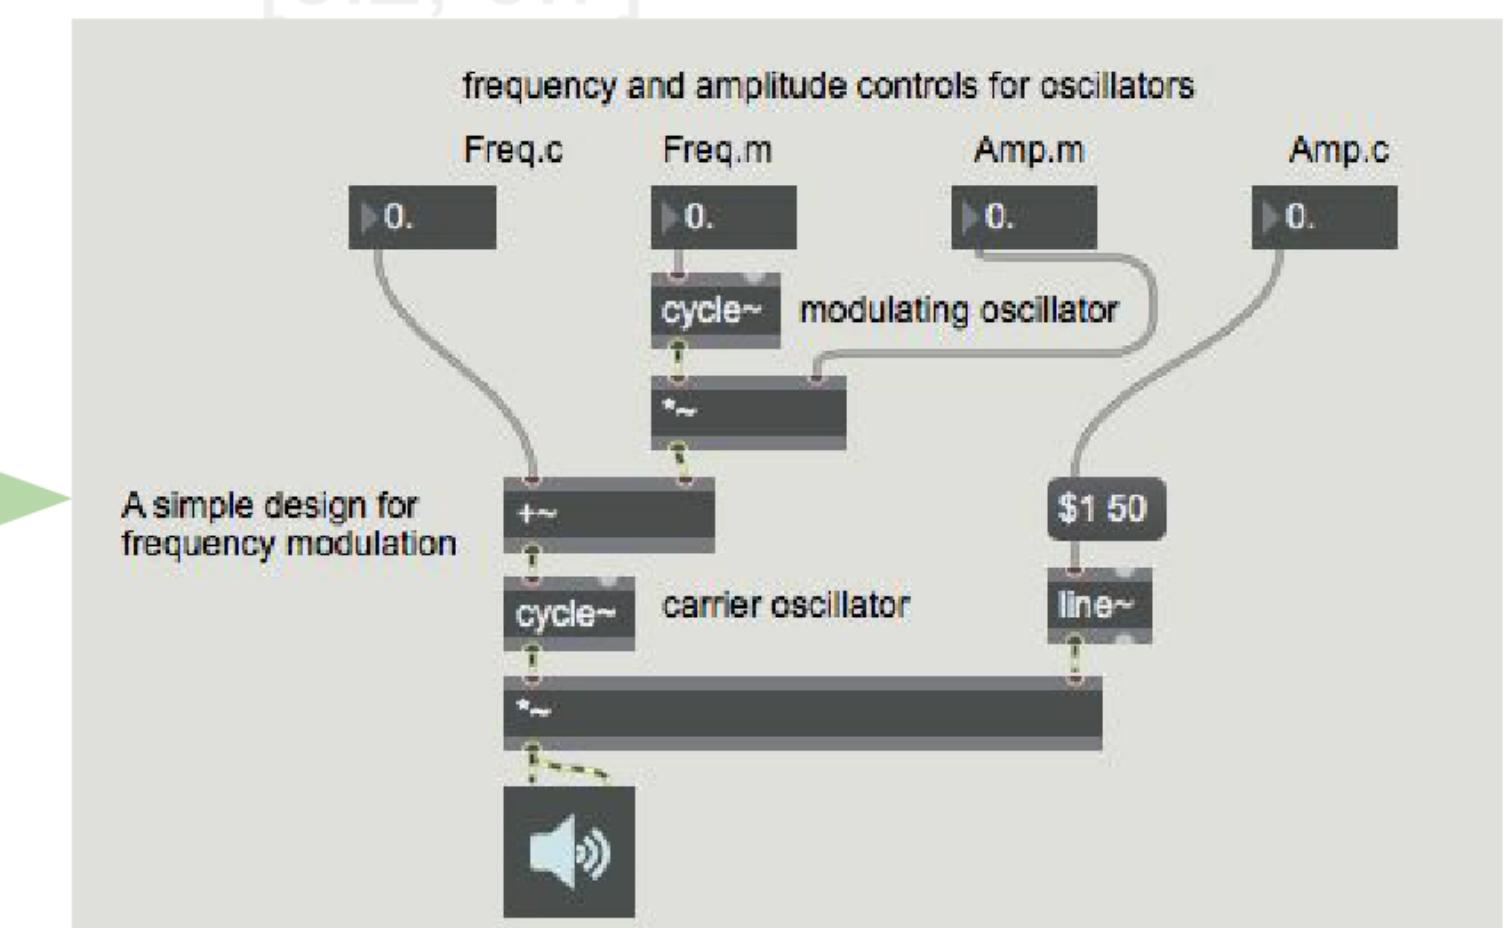
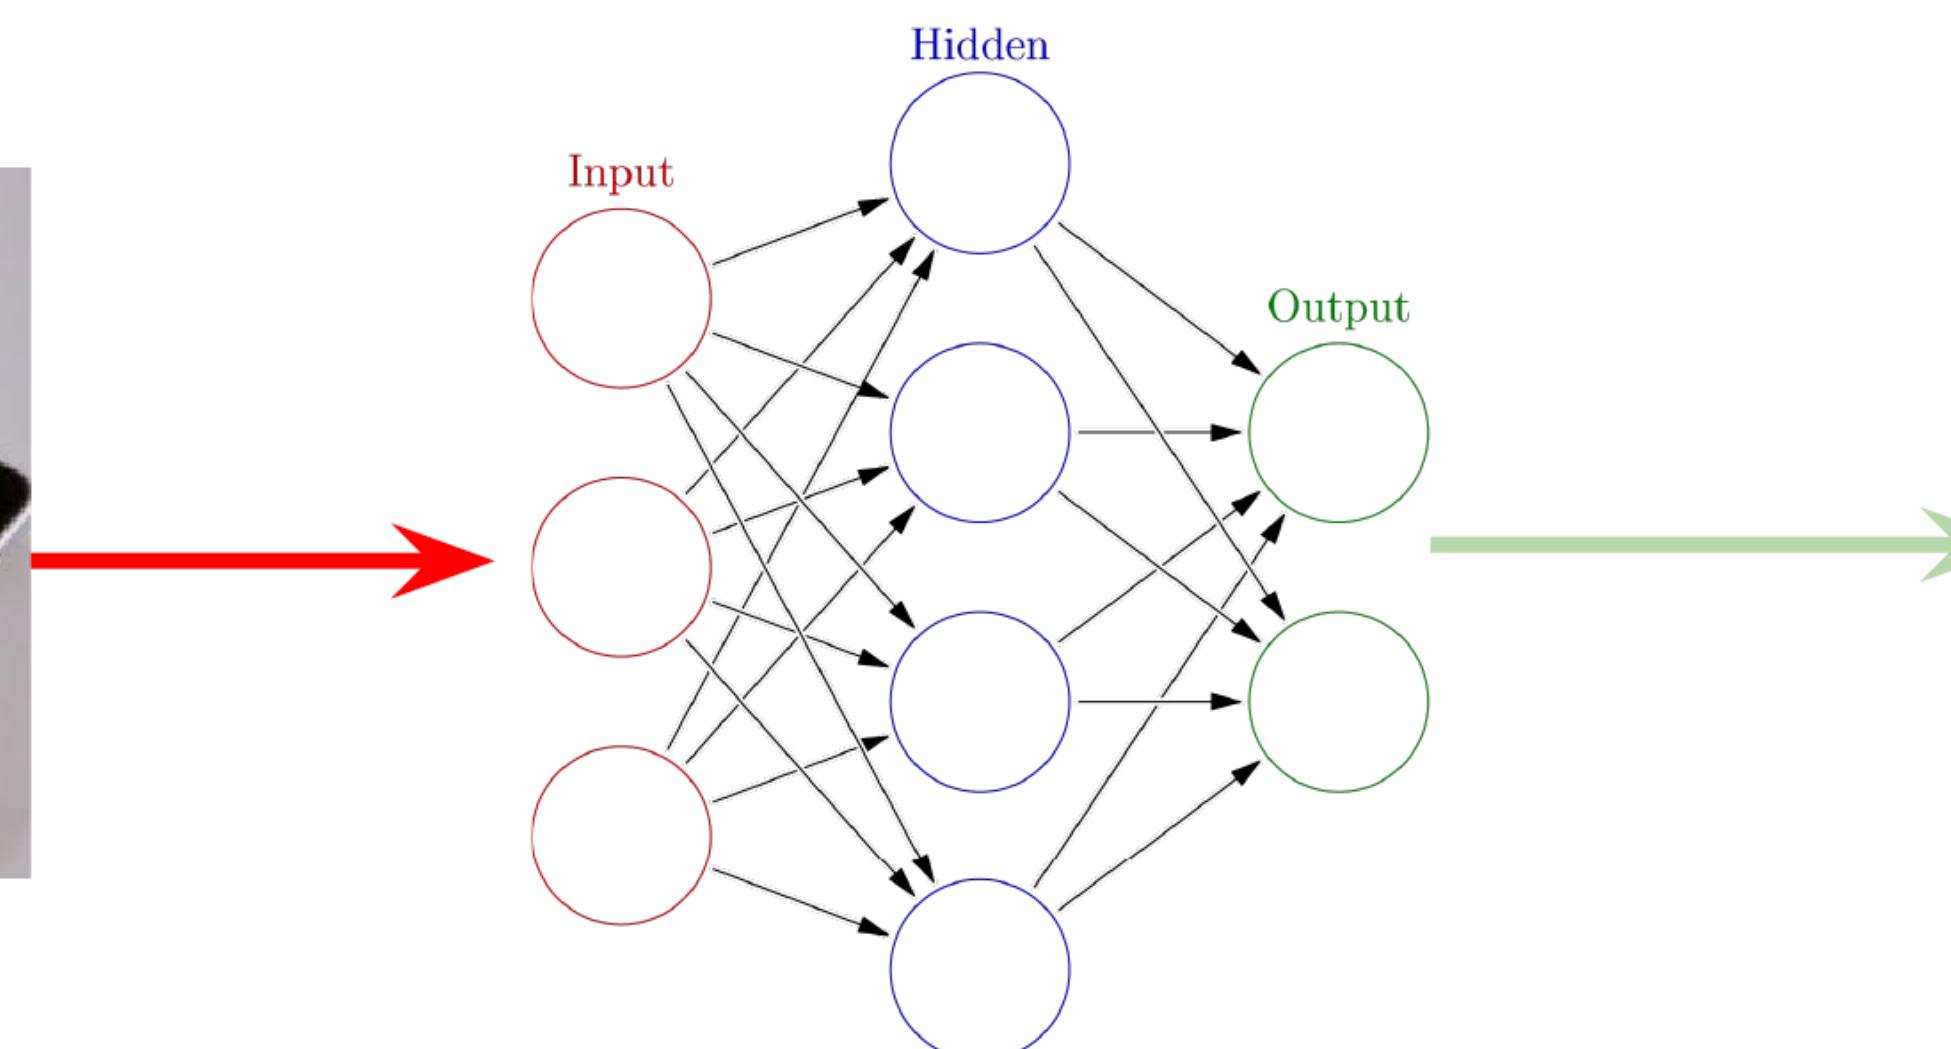
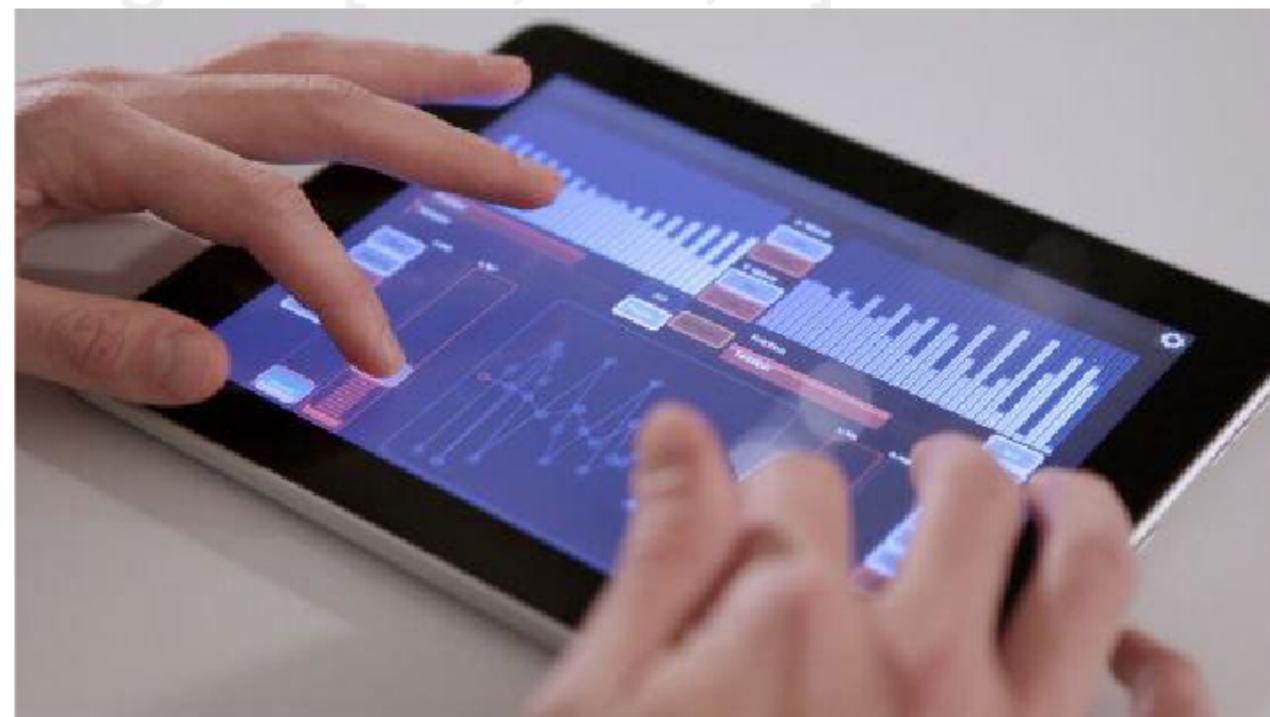
“a”: [0.1, 0.5, 1]  
“b”: [0, 1, 1]  
“c”: [1, 0, 1]  
“d”: [0.2, 0, 1]  
“e”: [0, 0, 0.5]  
“f”: [0.8, 0.2, 1]  
“g”: [0.9, 0.3, 0]  
“h”: [0.4, 1, 0.7]  
“i”: [0.2, 0.3, 1]  
“j”: [0.1, 1, 0.1]  
“k”: [0, 1, 1]  
“l”: [0, 1, 1]  
“m”: [0, 0, 1]  
“n”: [0.4, 1, 0.7]  
“o”: [0.2, 0.3, 1]  
“p”: [0.1, 1, 0.1]  
“q”: [0, 1, 1]  
“r”: [0, 1, 1]



[0.5, 0.7]  
[0.1, 0.1]  
[0.5, 0.6]  
[0.3, 0.9]  
[0.4, 0]  
[0.2, 0.7]  
[0.5, 0.6]  
[0.6, 0.1]  
[0.5, 0.3]  
[0.5, 0.7]  
[0.1, 0.8]  
[0.5, 0.1]  
[0.4, 0.2]  
[0.1, 0.8]  
[0.5, 0.6]  
[0.3, 0.9]  
[0.4, 0.3]  
[0.2, 0.7]

# Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]  
“b”: [0, 1, 1]  
“c”: [1, 0, 1]  
“d”: [0.2, 0, 1]  
“e”: [0, 0, 0.5]  
“f”: [0.8, 0.2, 1]  
“g”: [0.9, 0.3, 0]

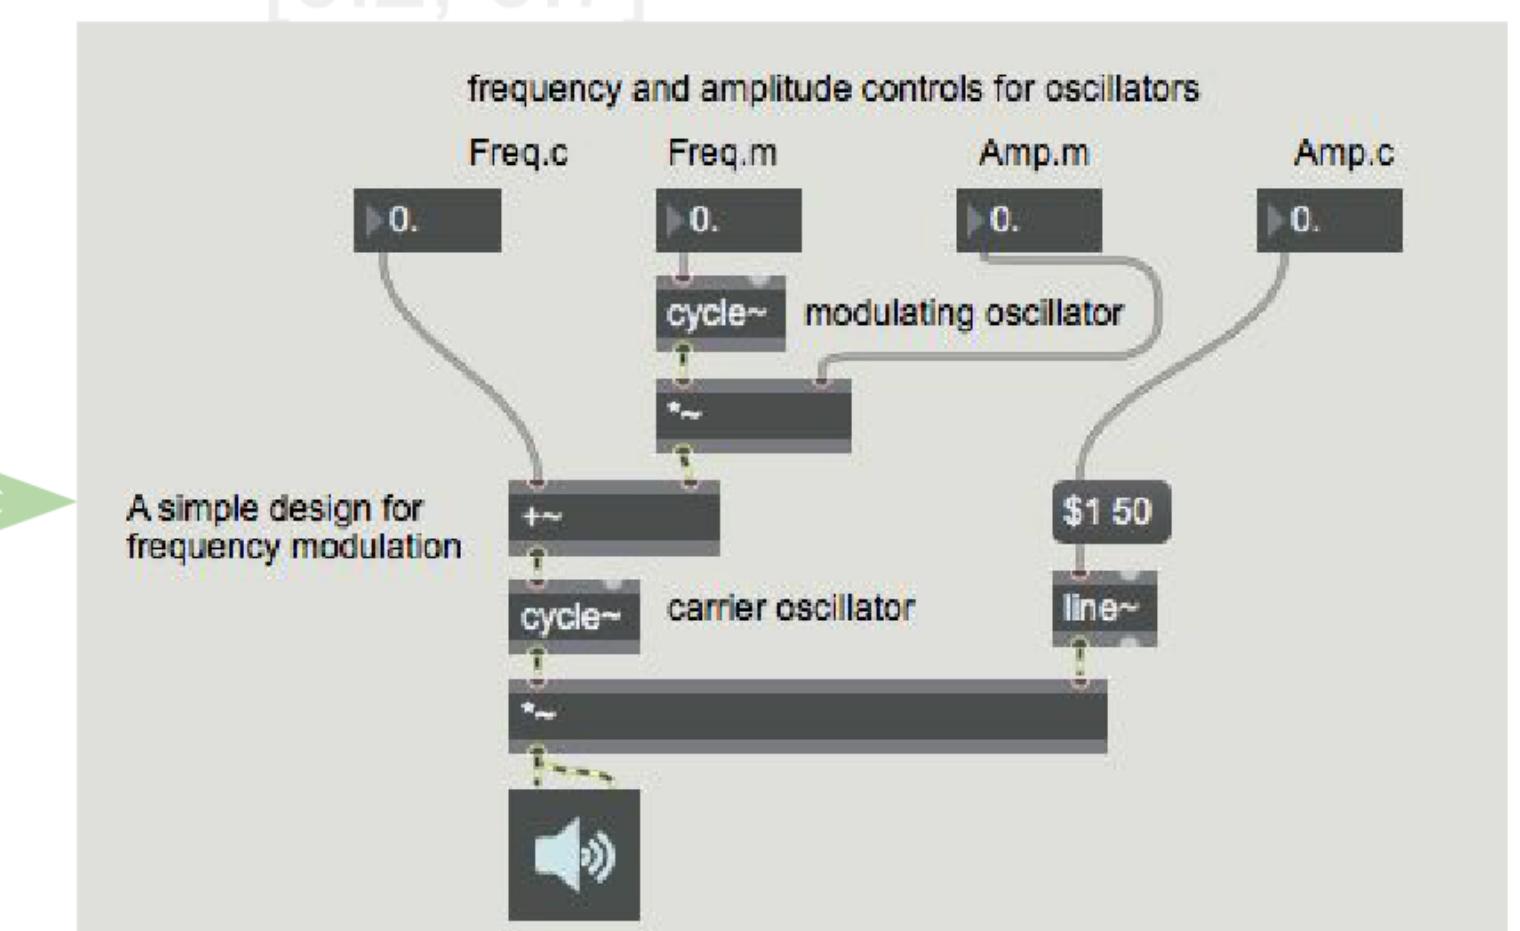
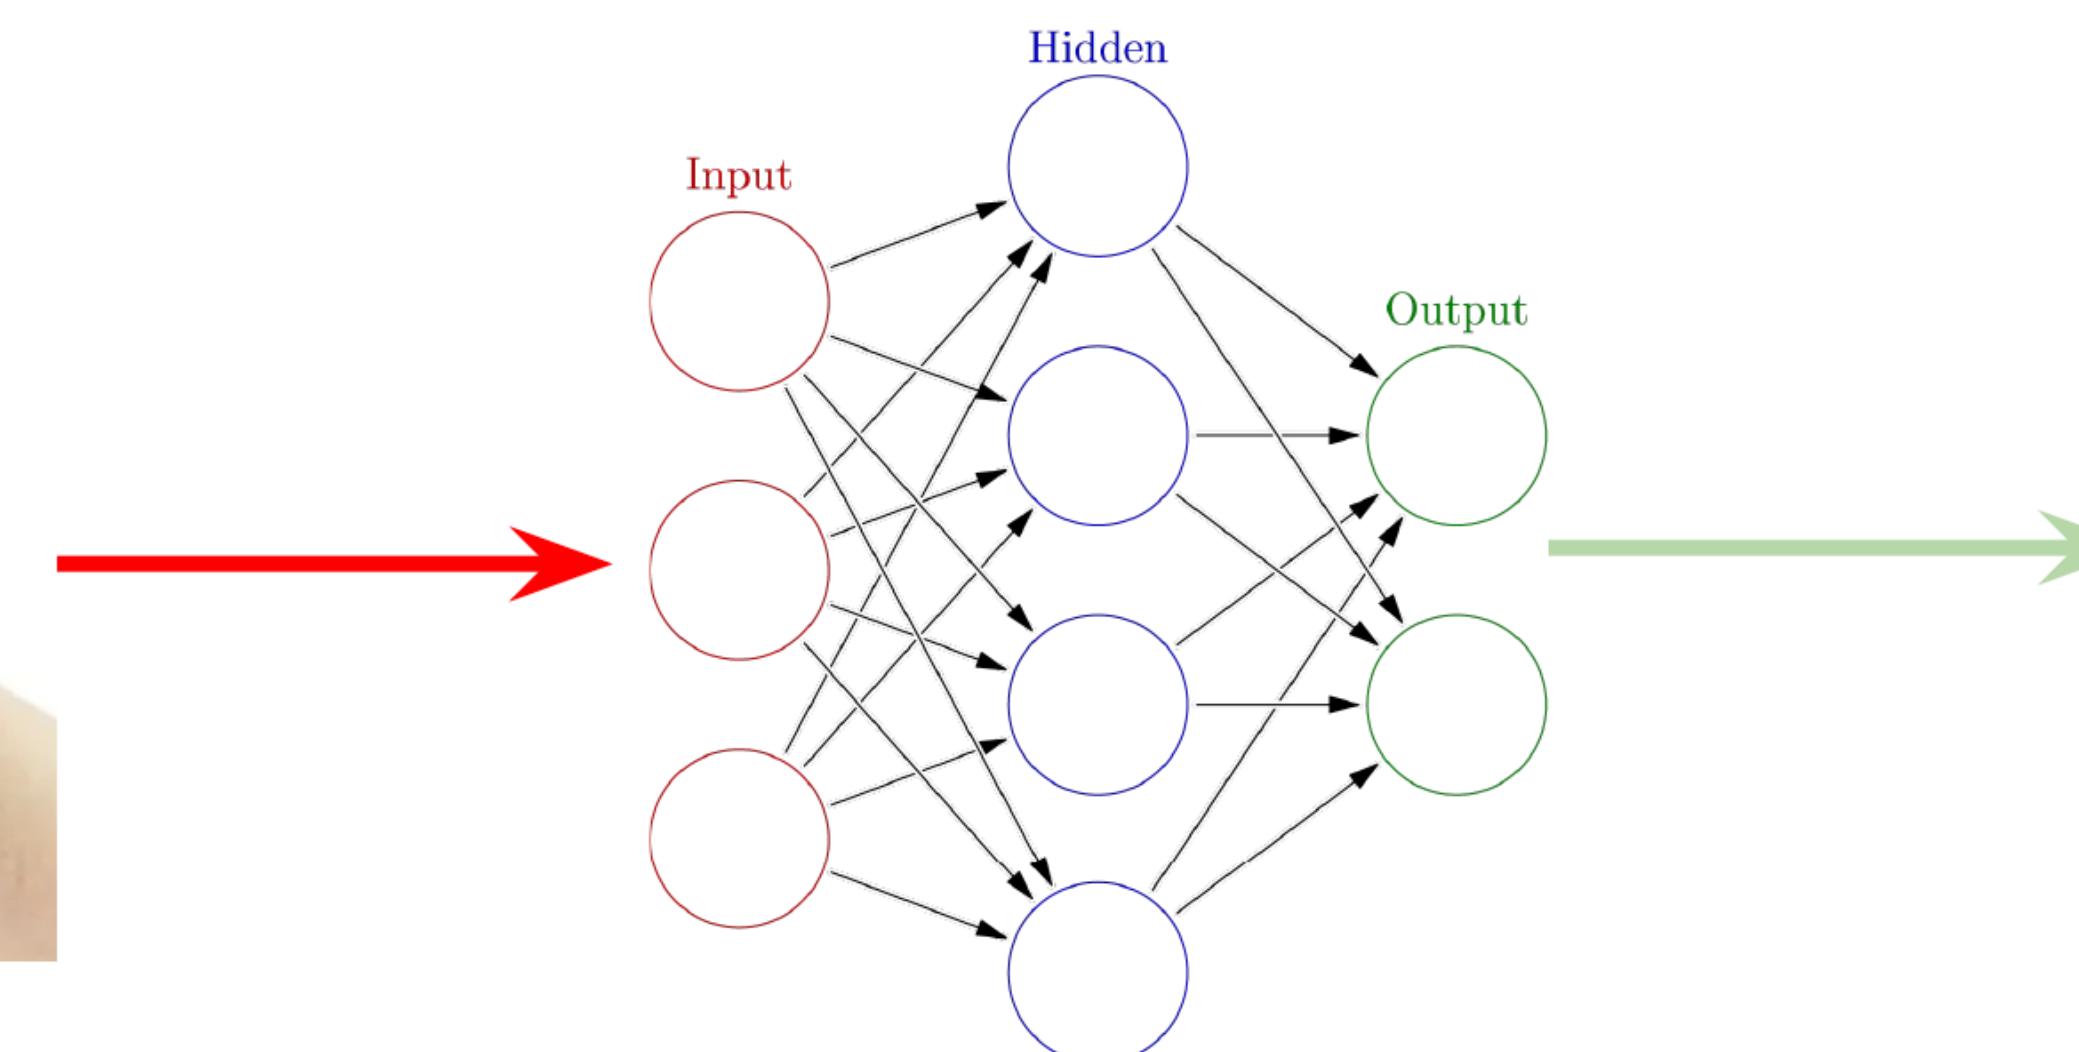
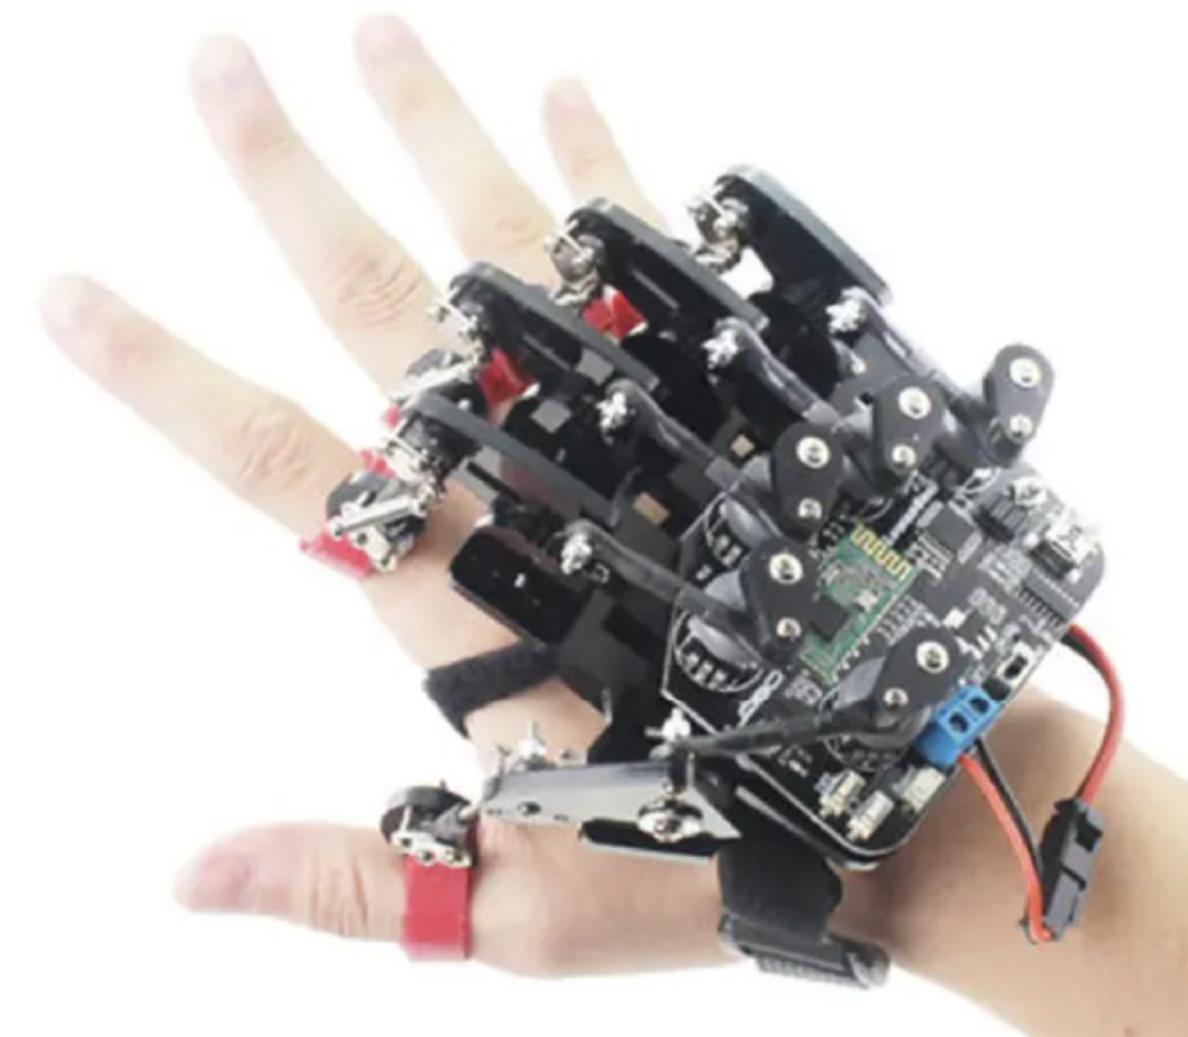


“n”: [0.4, 1, 0.7]  
“o”: [0.2, 0.3, 1]  
“p”: [0.1, 1, 0.1]  
“q”: [0, 1, 1]  
“r”: [0, 1, 1]

[0.5, 0.7]  
[0.1, 0.1]  
[0.5, 0.6]  
[0.3, 0.9]  
[0.4, 0]  
[0.2, 0.7]

[0.5, 0.7]  
[0.1, 0.8]  
[0.5, 0.1]  
[0.4, 0.2]  
[0.1, 0.8]

# *Neural Network Predicting with Regression*



# Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]

“b”: [0, 1, 1]

“c”: [1, 0, 1]

“d”: [0.2, 0, 1]

“e”: [0, 0, 0.5]

“f”: [0.8, 0.2, 1]



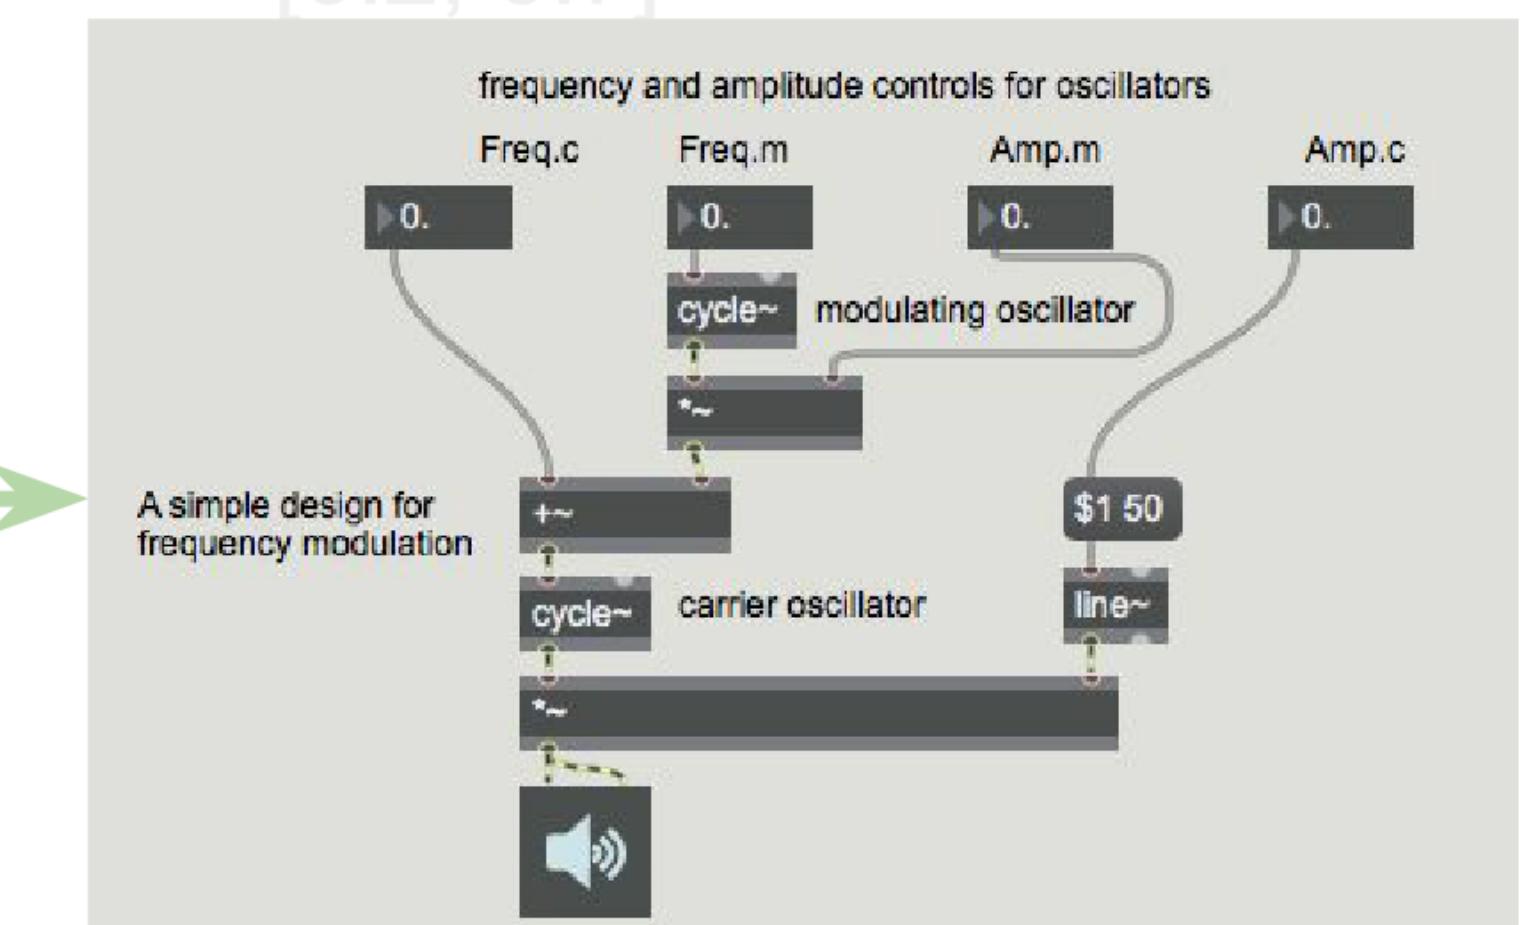
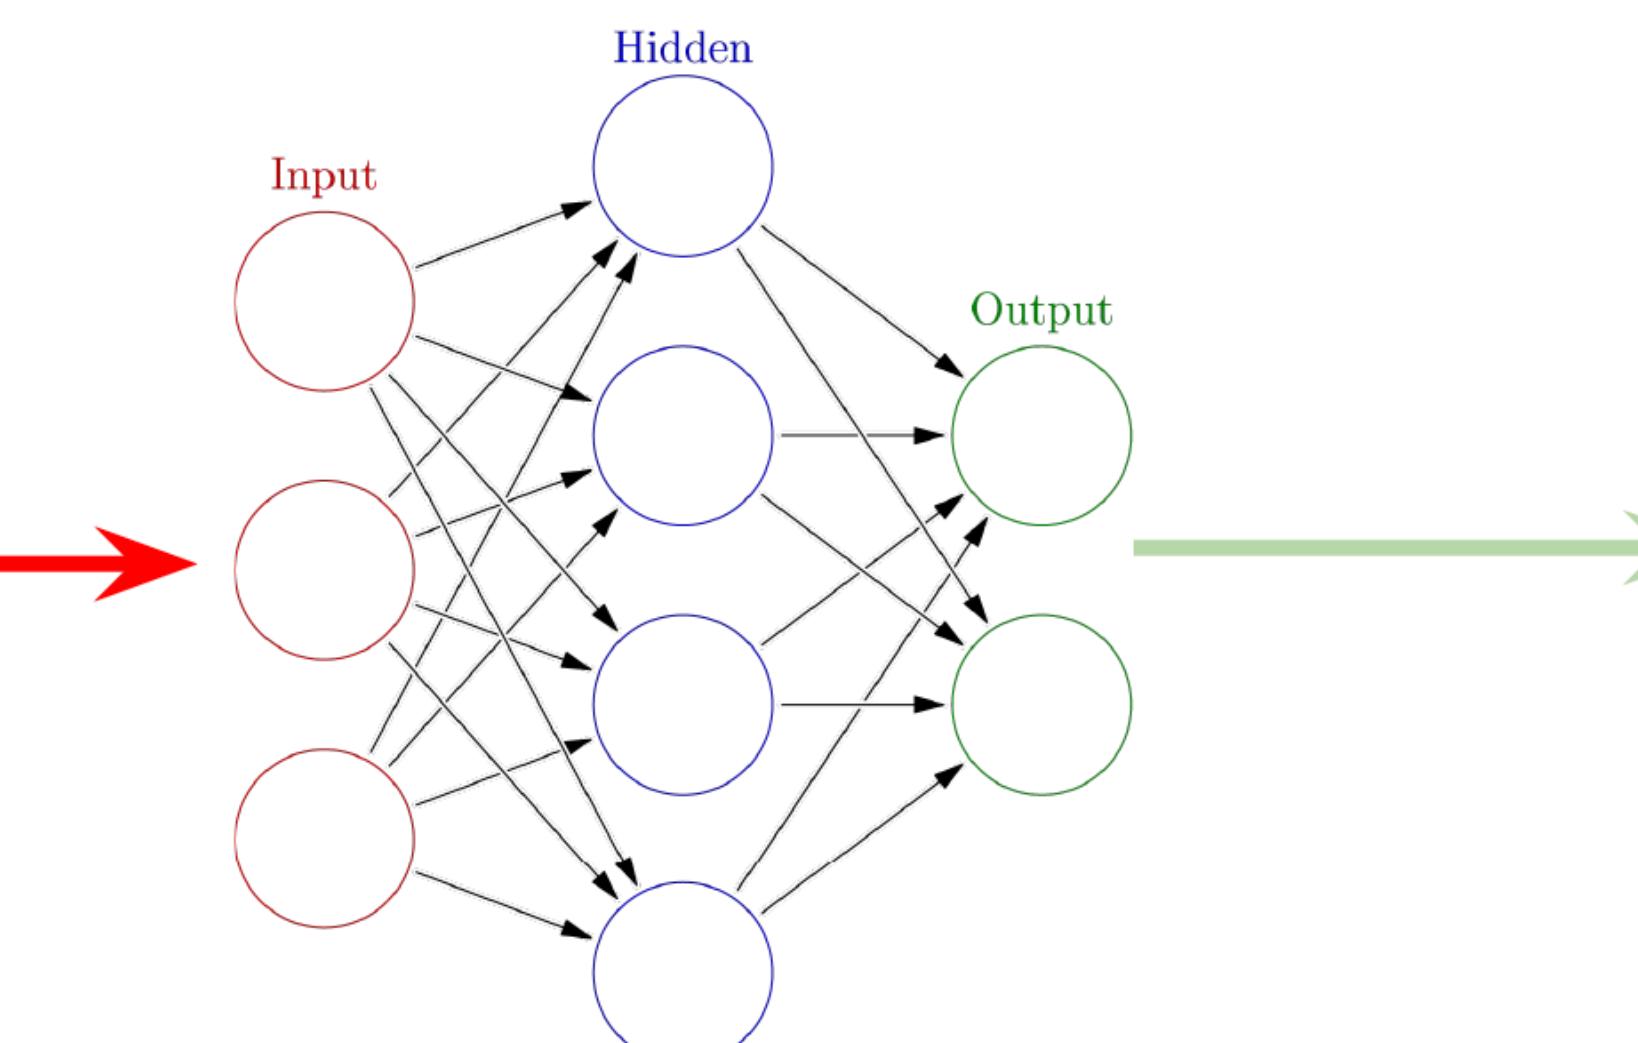
e.g.:

**pitch**

**loudness**

**spectral centroid**

**spectral flatness**



“g”: [0.1, 1, 1]

“h”: [0, 1, 1]

“i”: [0, 0, 1]

“j”: [0.4, 1, 0.7]

“k”: [0, 1, 1]

“l”: [0, 1, 1]

“m”: [0, 0, 1]

“n”: [0.4, 1, 0.7]

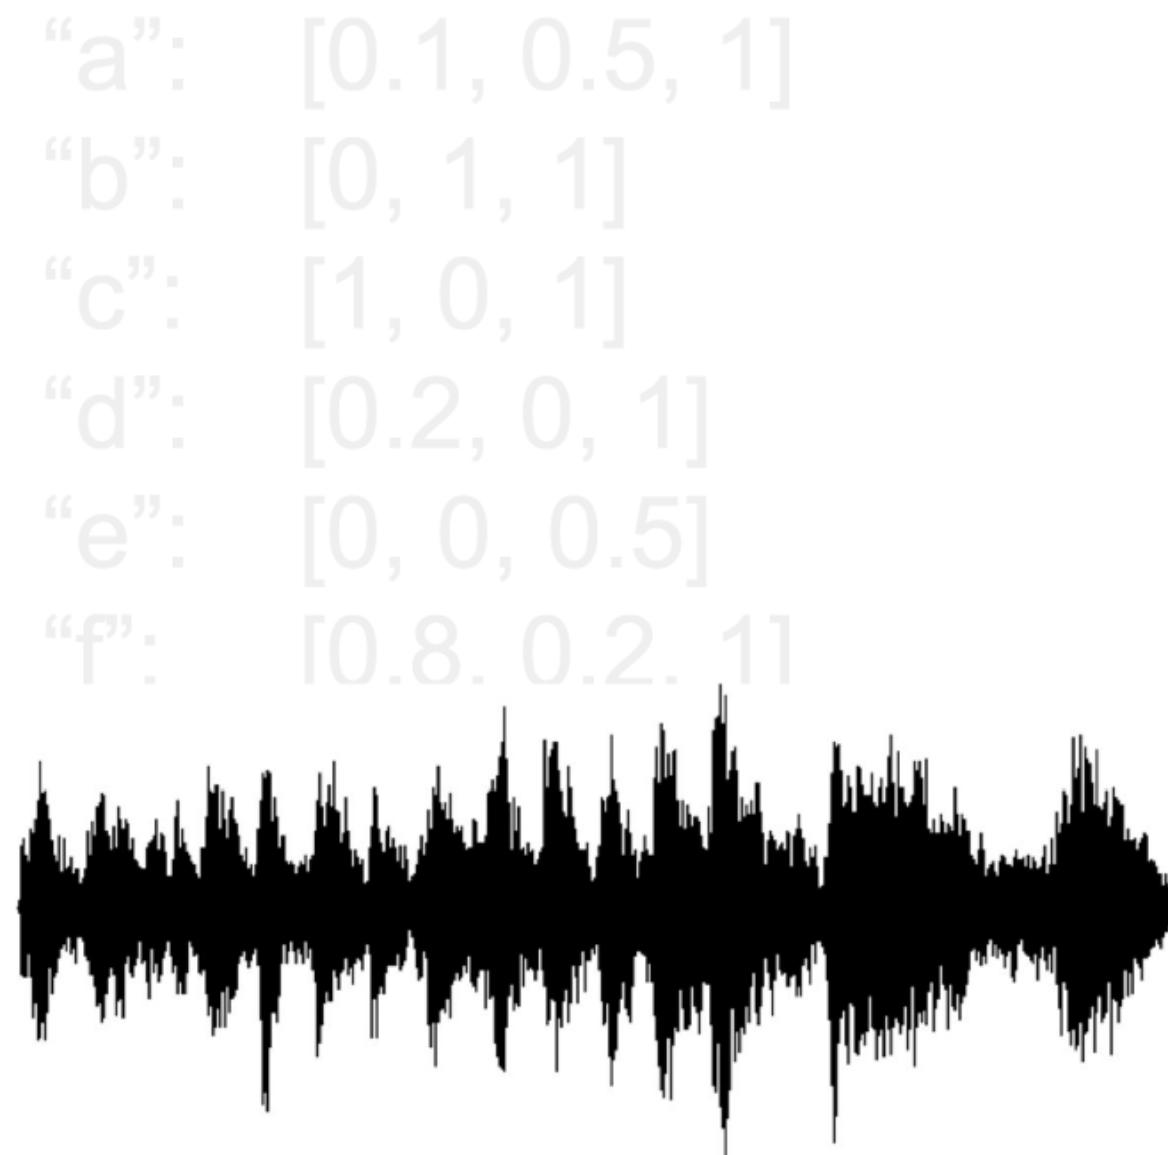
“o”: [0.1, 0.8]

“p”: [0.5, 0.1]

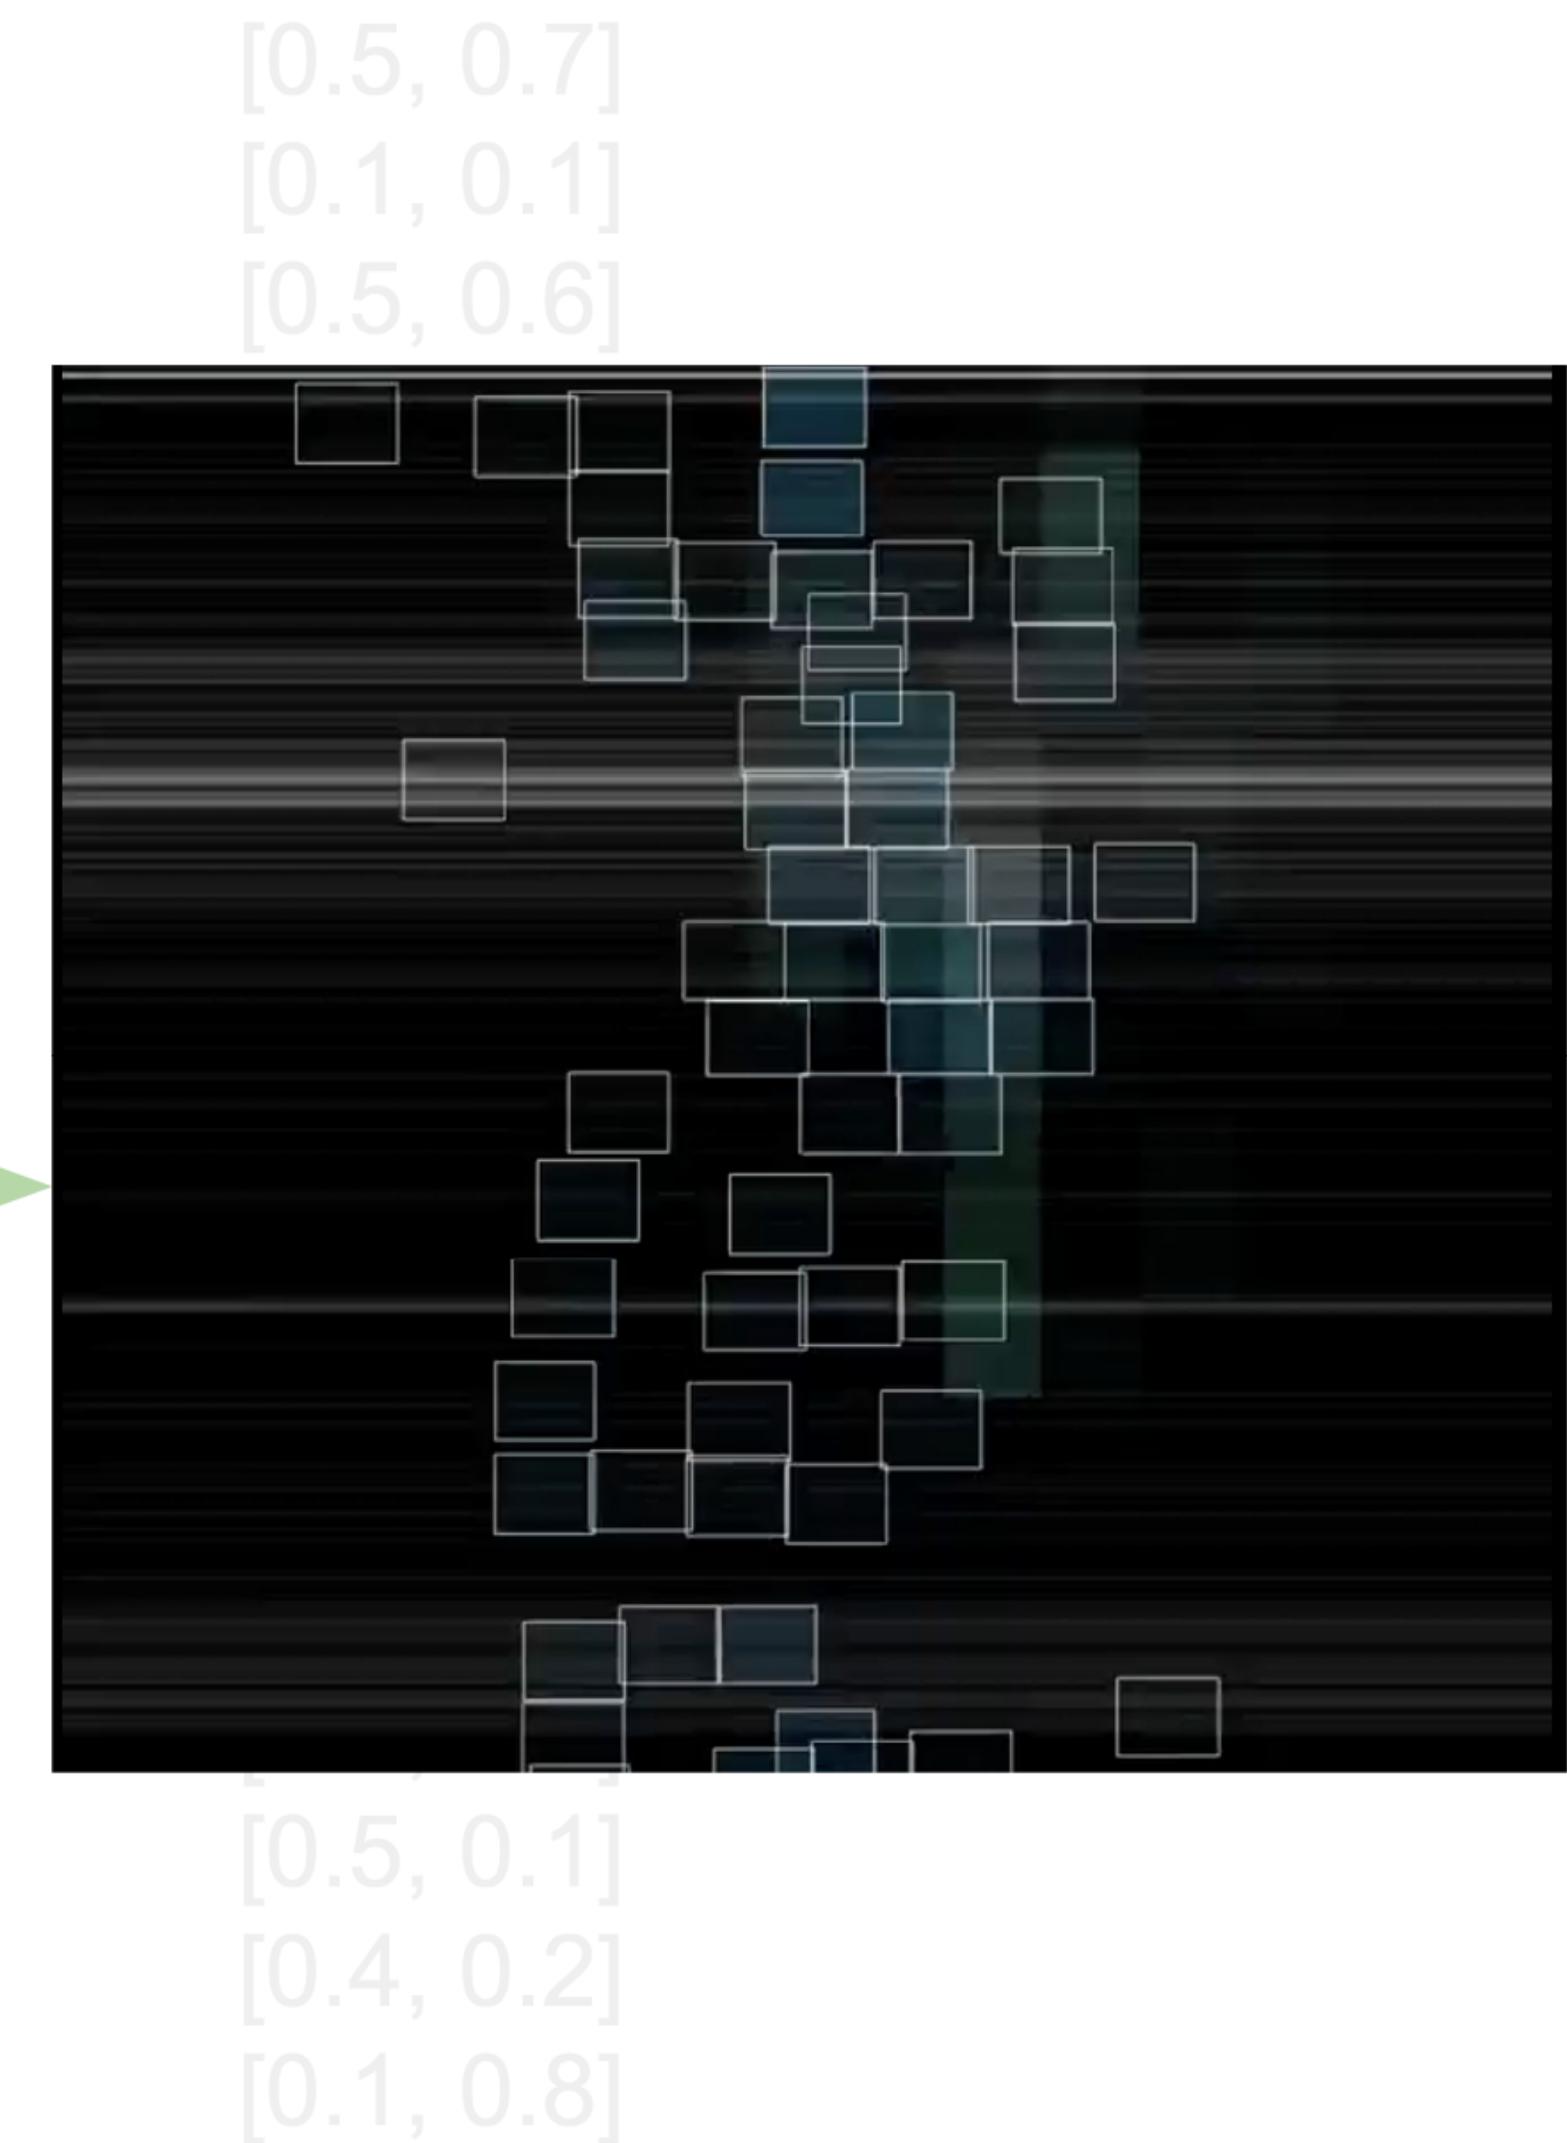
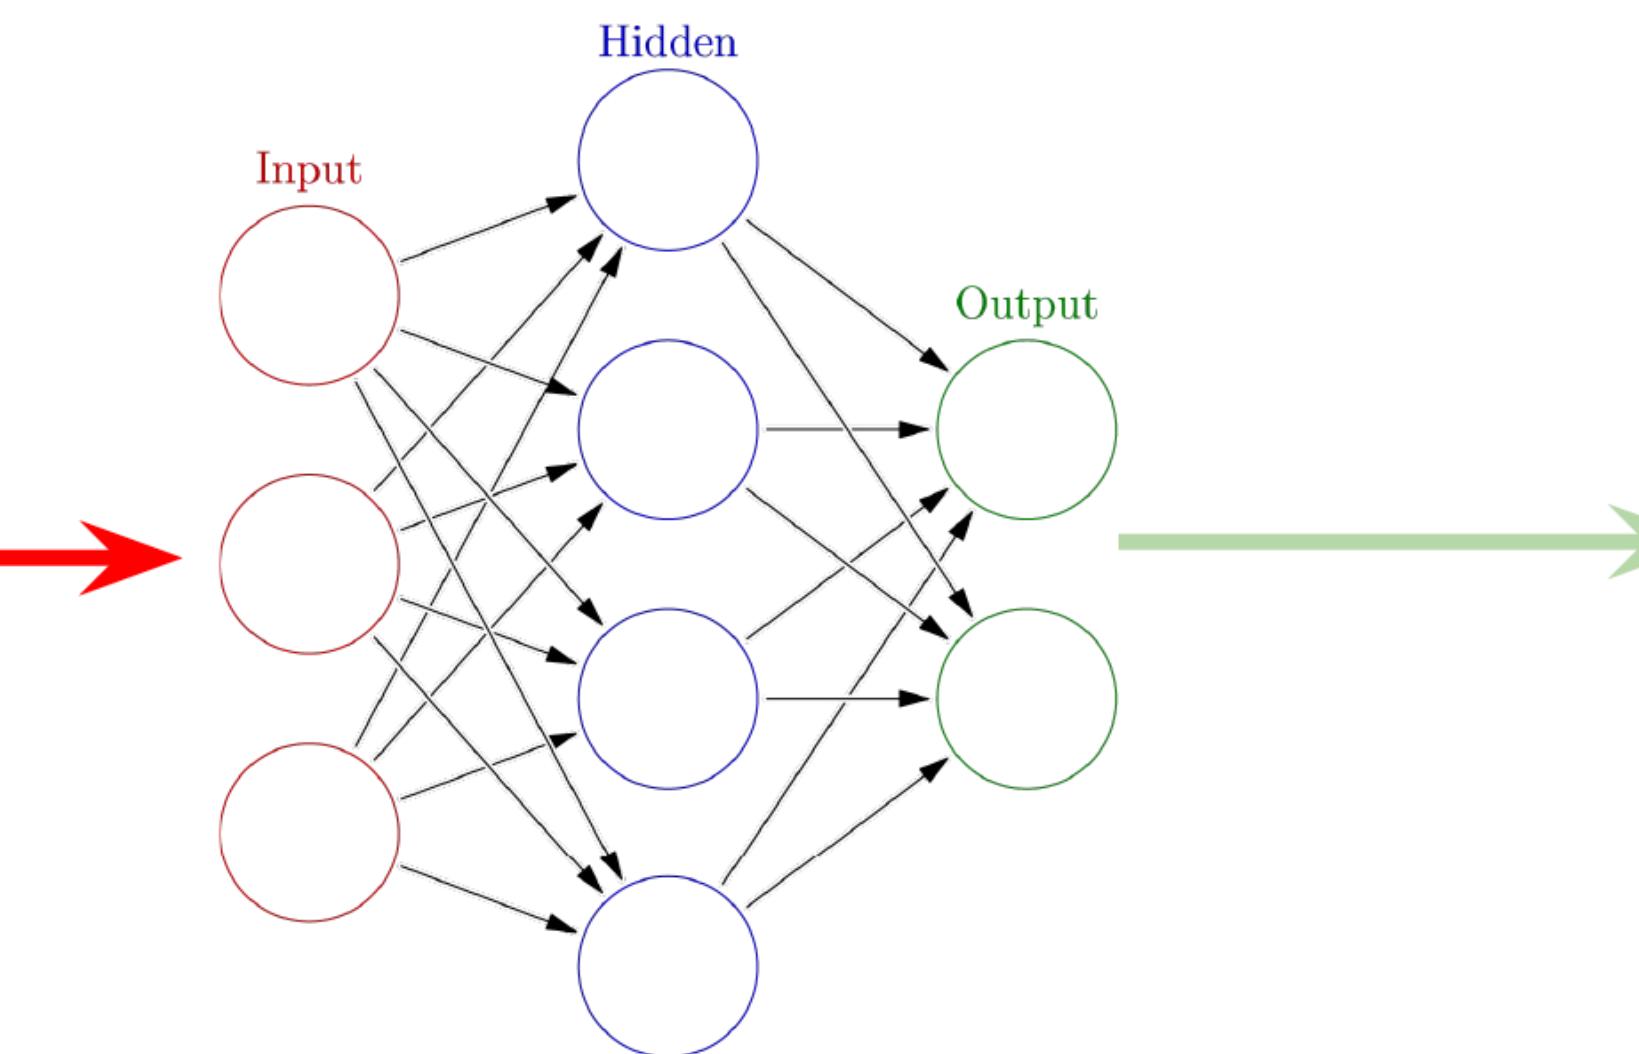
“q”: [0.4, 0.2]

“r”: [0.1, 0.8]

# Neural Network Predicting with Regression

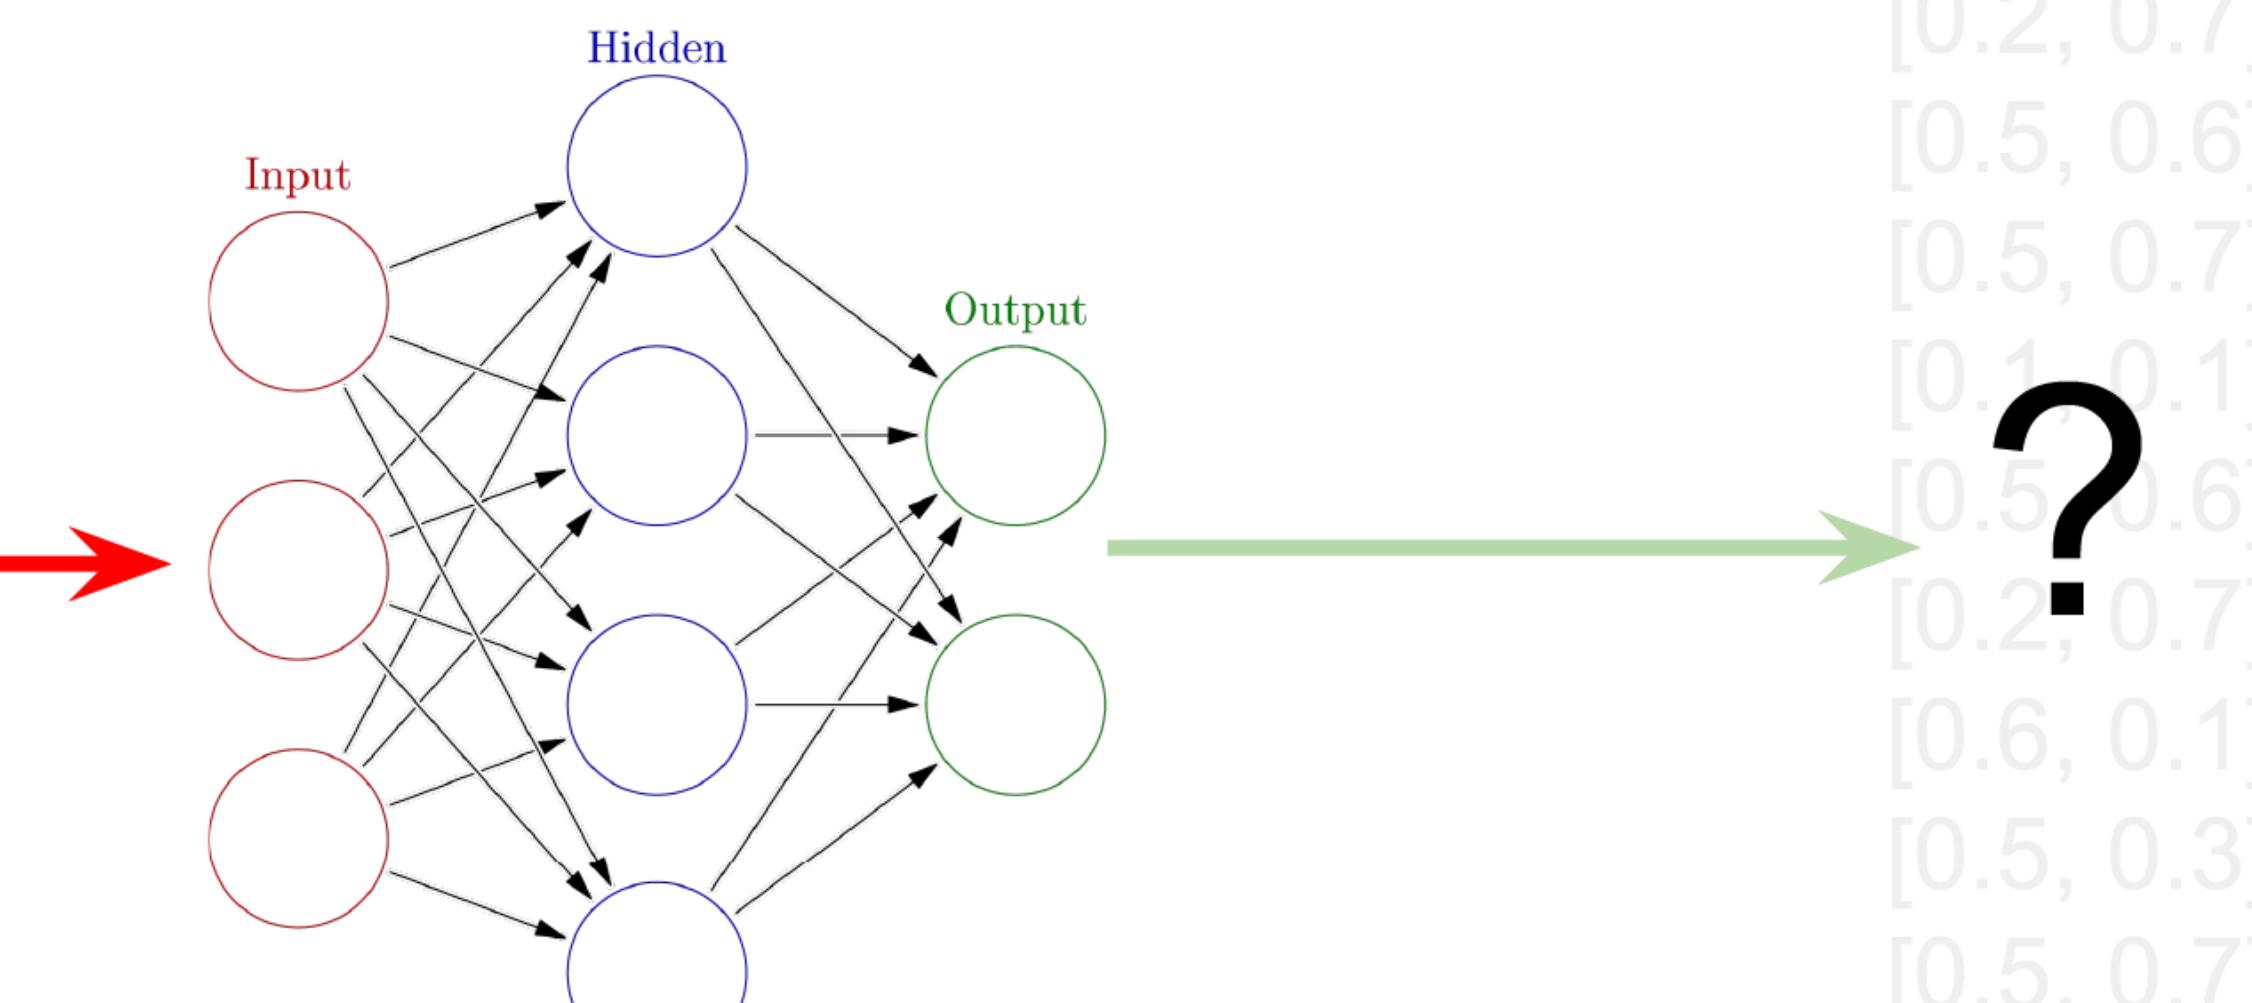


“a”: [0.1, 0.5, 1]  
“b”: [0, 1, 1]  
“c”: [1, 0, 1]  
“d”: [0.2, 0, 1]  
“e”: [0, 0, 0.5]  
“f”: [0.8, 0.2, 1]  
“j”: [0, 0, 0.5]  
“k”: [0, 1, 1]  
“l”: [0, 1, 1]  
“m”: [0, 0, 1]  
“n”: [0.4, 1, 0.7]  
“o”: [0.2, 0.3, 1]  
“p”: [0.1, 1, 0.1]  
“q”: [0, 1, 1]  
“r”: [0, 1, 1]



# Neural Network Predicting with Regression

“a”: [0.1, 0.5, 1]  
“b”: [0, 1, 1]  
“c”: [1, 0, 1]  
“d”: [0.2, 0, 1]  
“e”: [0, 0, 0.5]  
“f”: [0.8, 0.2, 1]  
“g”: [0.9, 0.3, 0]  
“h”: [0.4, 1, 0.7]  
“i”: [0.2, 0.3, 1]  
“j”: [0.1, 0.1, 0.1]  
“k”: [0, 1, 1]  
“l”: [0, 1, 1]  
“m”: [0, 0, 1]  
“n”: [0.4, 1, 0.7]  
“o”: [0.2, 0.3, 1]  
“p”: [0.1, 1, 0.1]  
“q”: [0, 1, 1]  
“r”: [0, 1, 1]



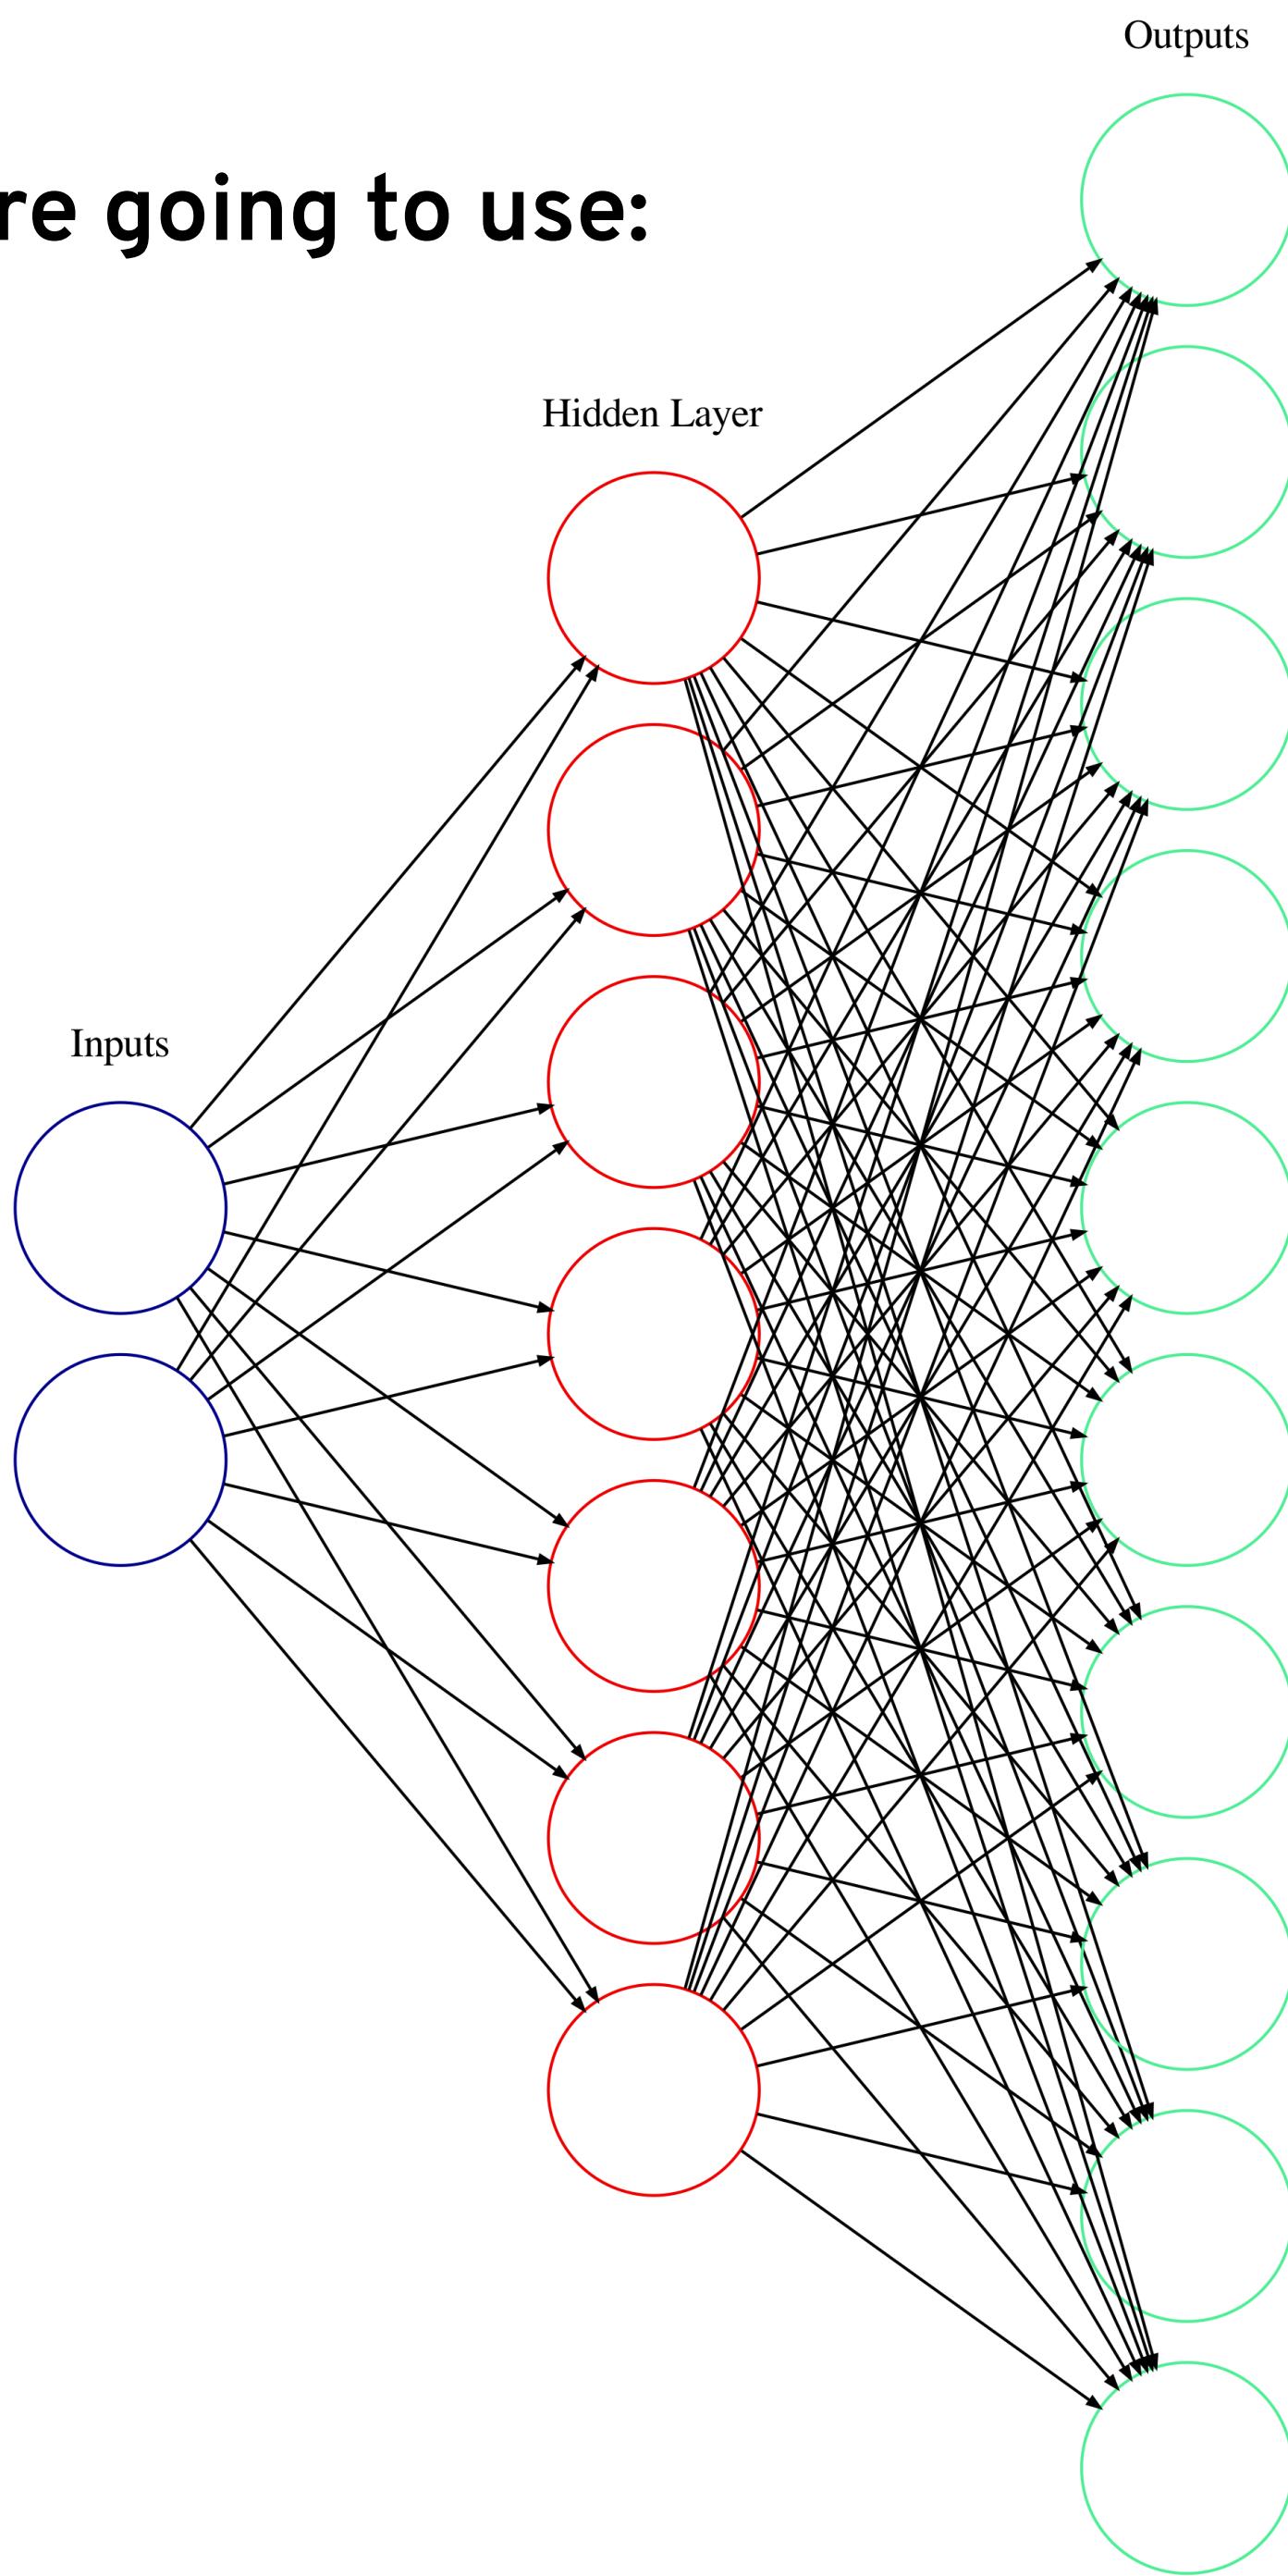
[0.5, 0.7]  
[0.1, 0.1]  
[0.5, 0.6]  
[0.3, 0.9]  
[0.4, 0]  
[0.2, 0.7]  
[0.5, 0.6]  
[0.5, 0.7]  
[0.1, 0.1]  
[0.5, 0.6]  
[0.2, 0.7]  
[0.6, 0.1]  
[0.5, 0.3]  
[0.5, 0.7]  
[0.1, 0.8]  
[0.5, 0.1]  
[0.4, 0.2]  
[0.1, 0.8]

## Structure of the neural network we're going to use:

2 inputs

1 hidden layer of 7 nodes

10 outputs



# *Regression with Audio Analyses*

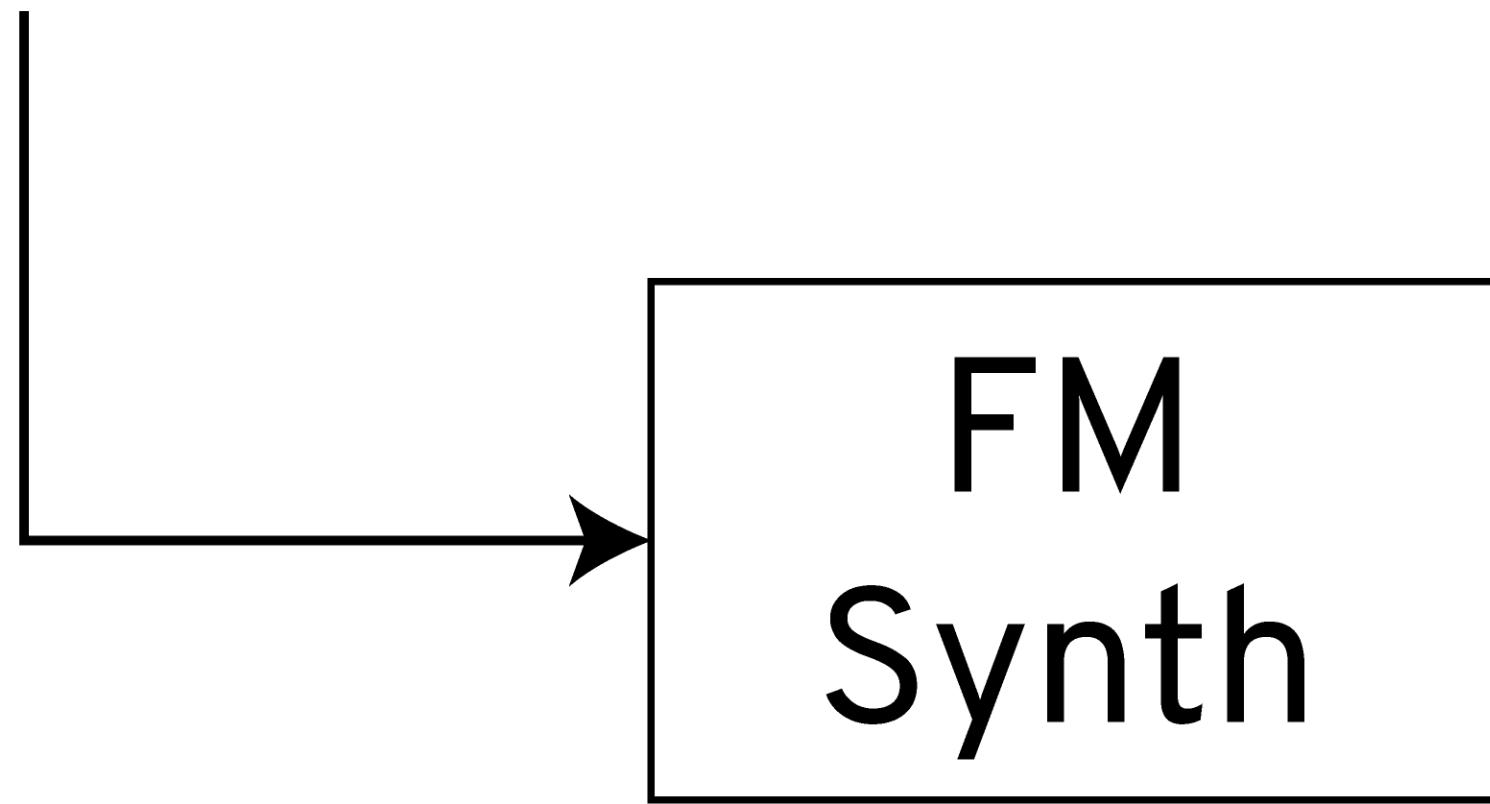


# Making the Datasets

## Freq. Mod. Params

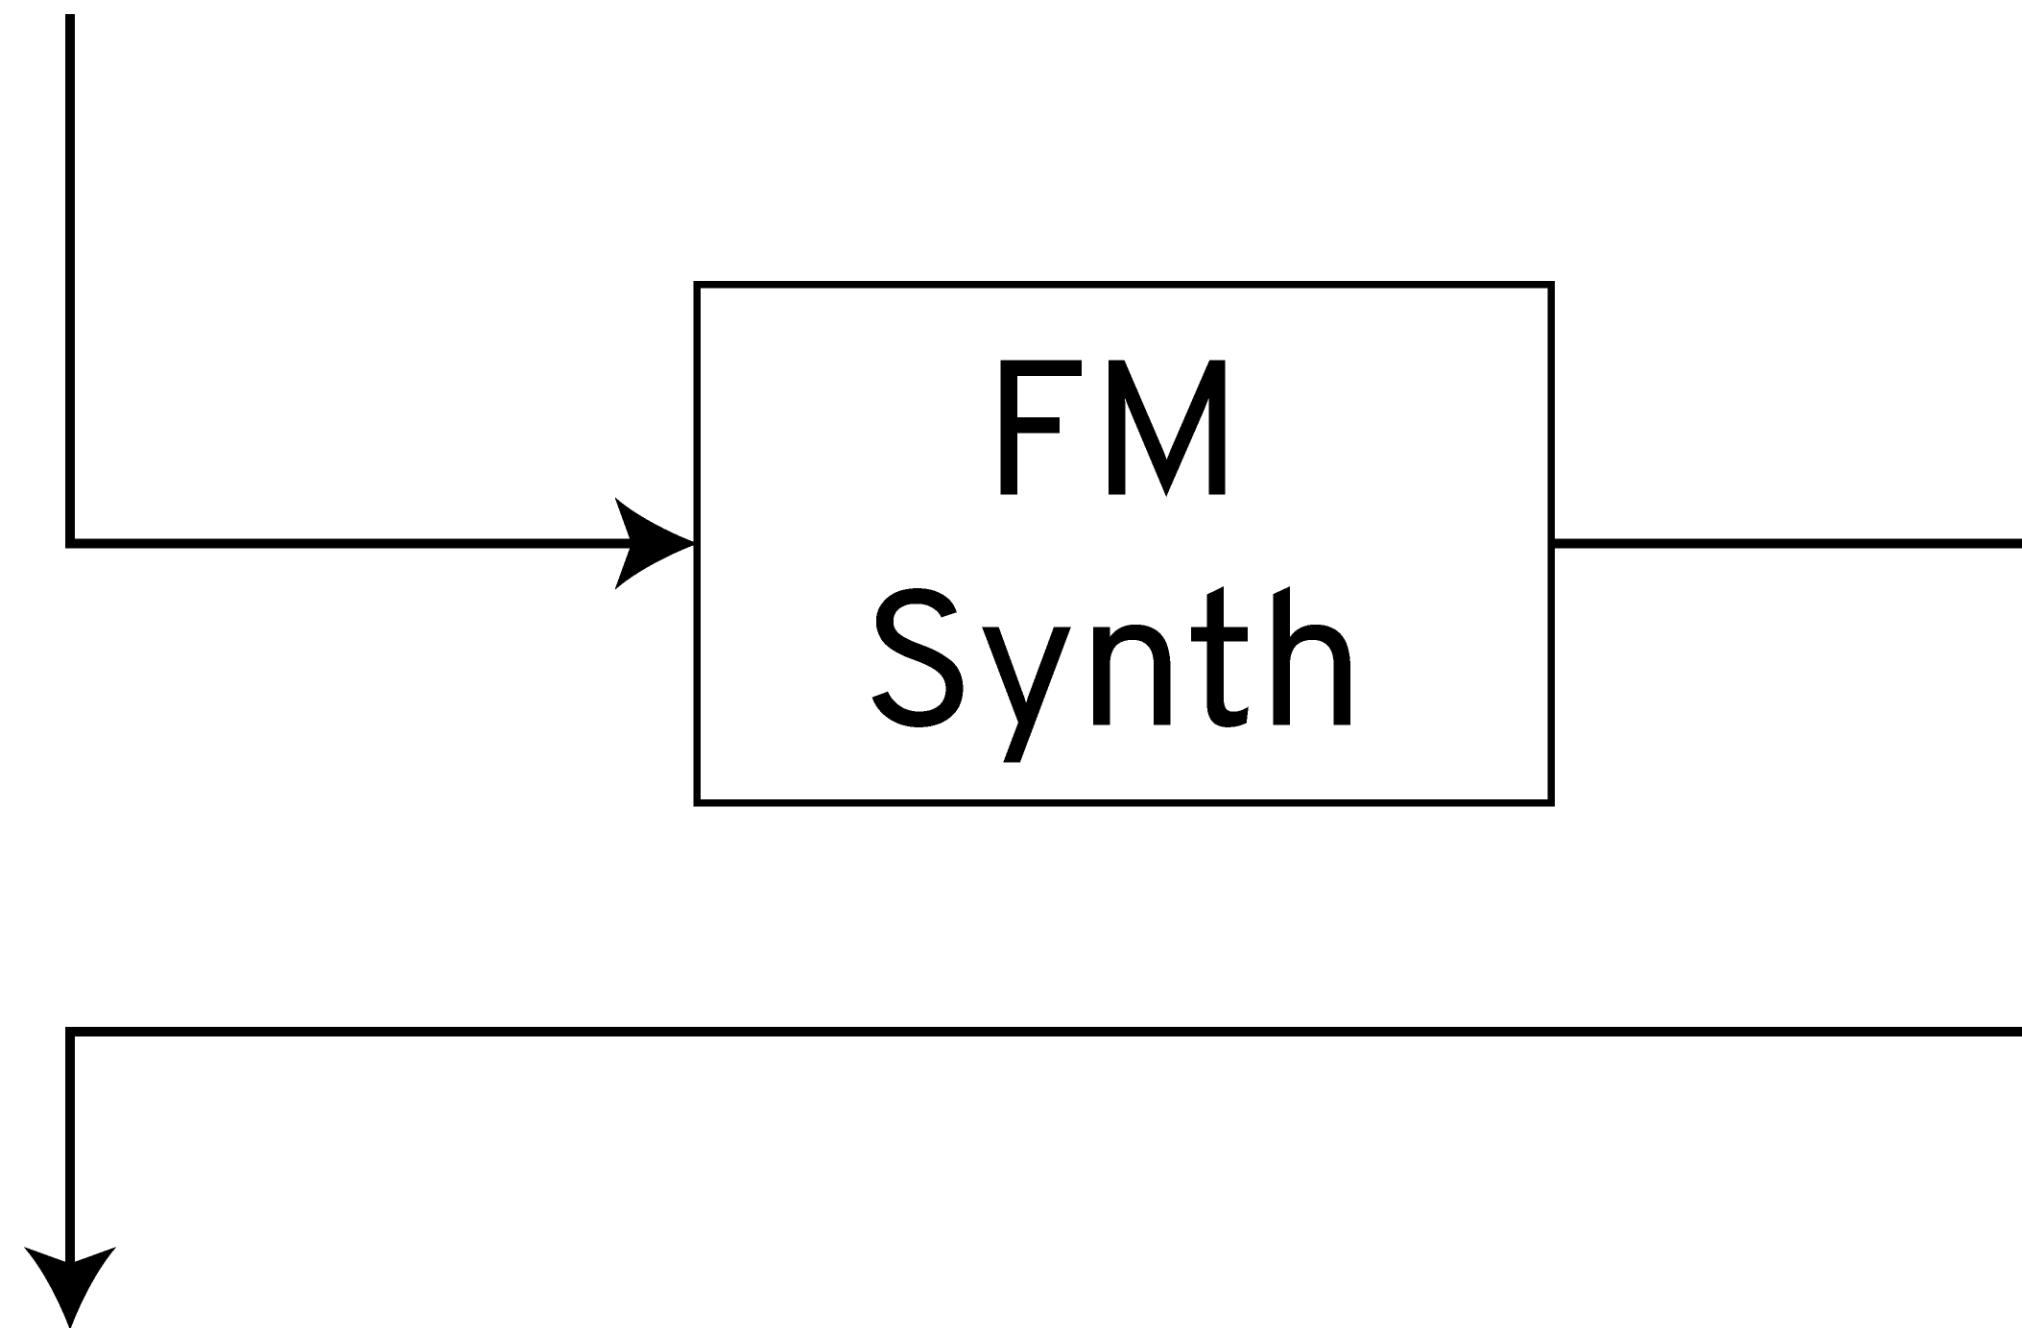
# Making the Datasets

Freq. Mod. Params



# Making the Datasets

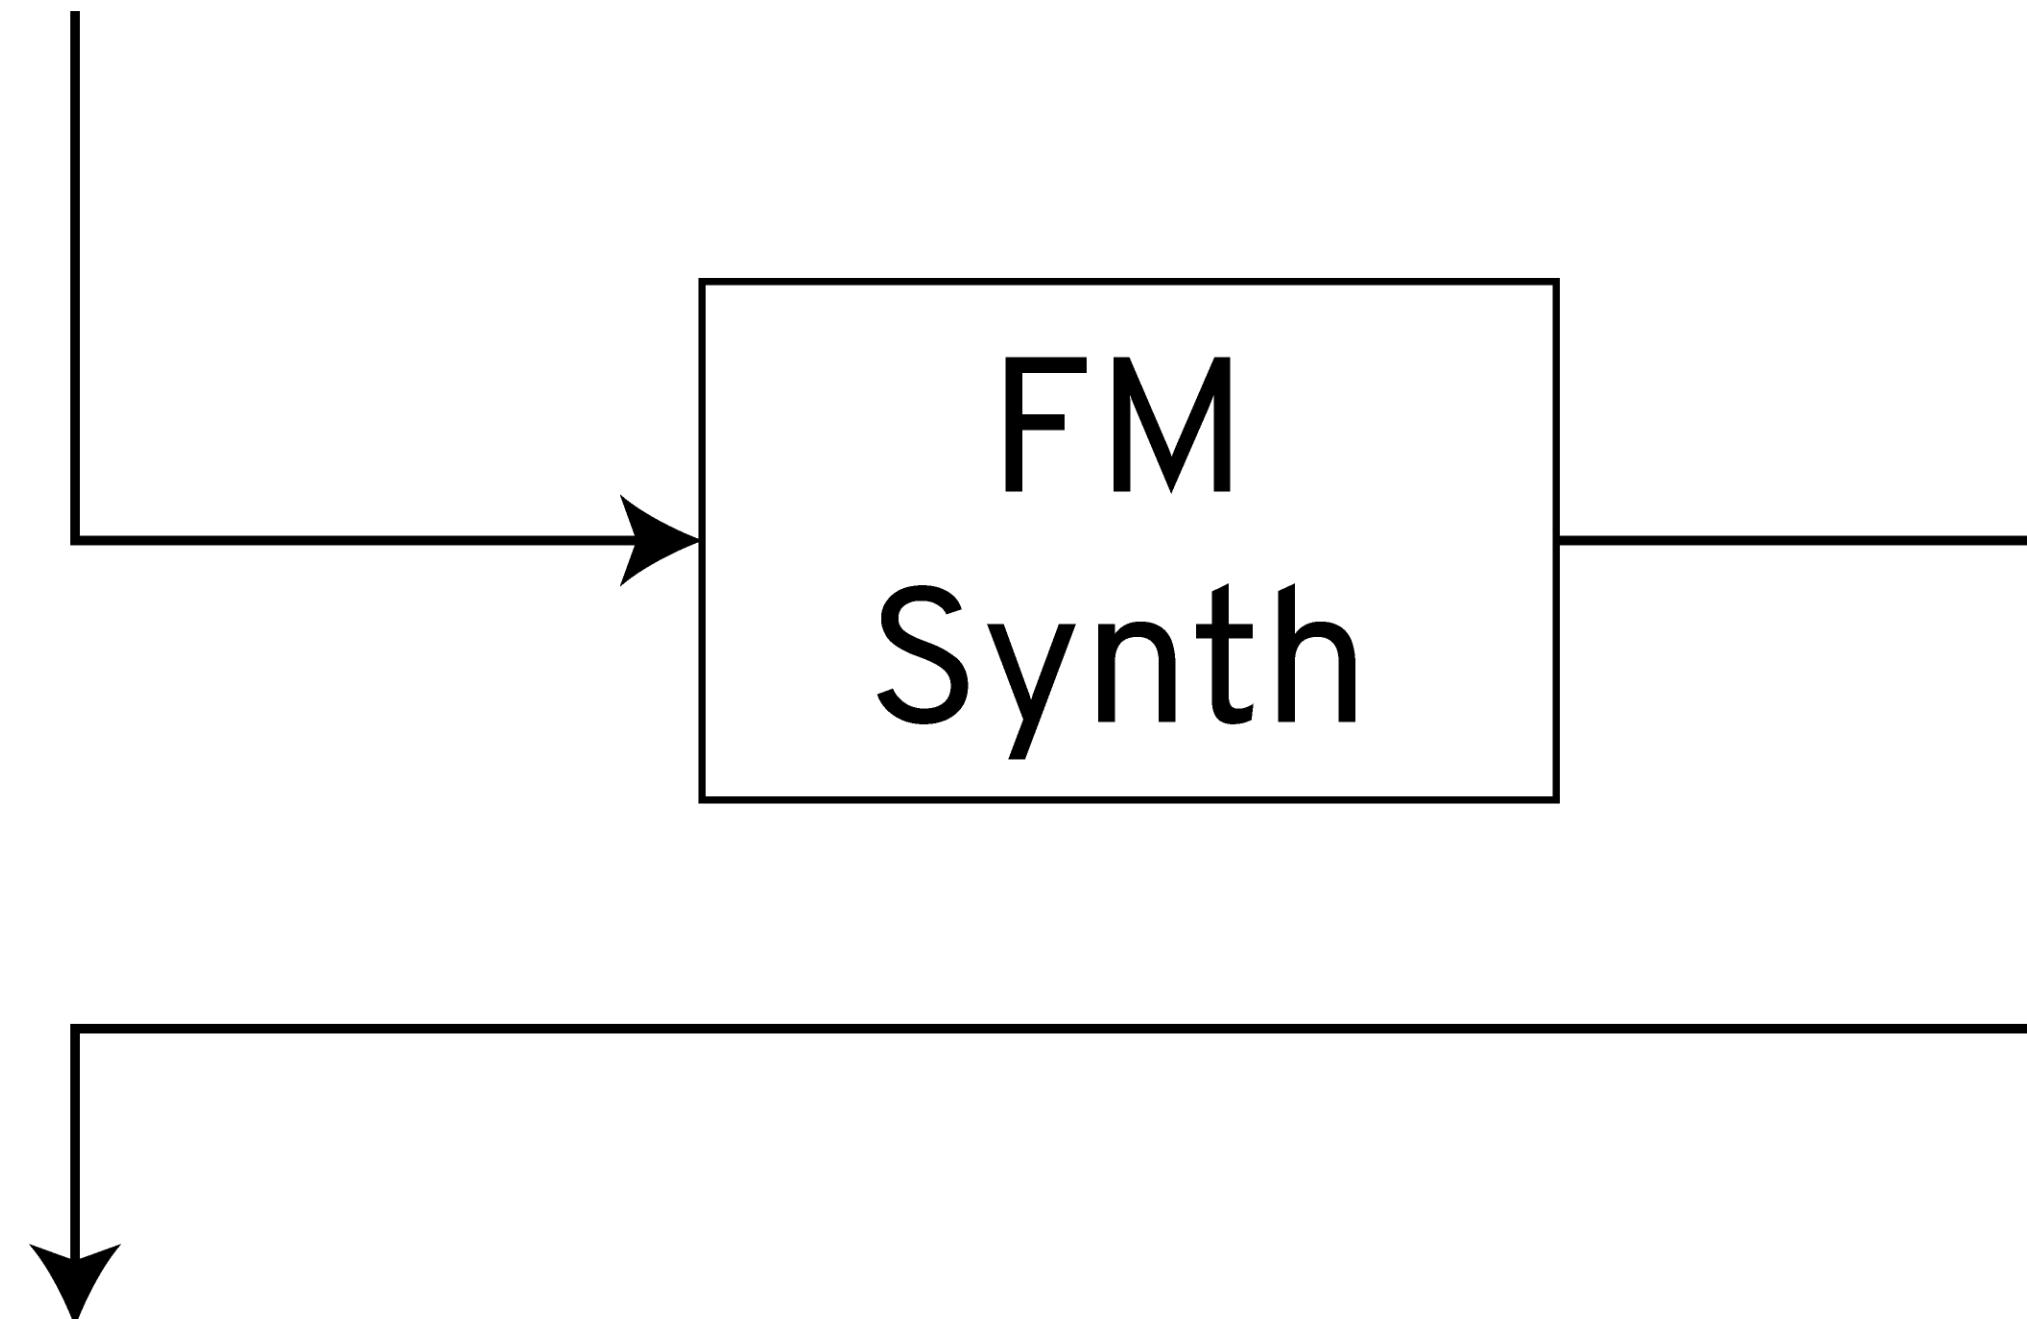
Freq. Mod. Params



Audio Signal

# Making the Datasets

Freq. Mod. Params



Audio Signal —————→ 13 MFCCs

# Making the Datasets

## Freq. Mod. Params

13 MFCCs

# Making the Datasets

## Freq. Mod. Params

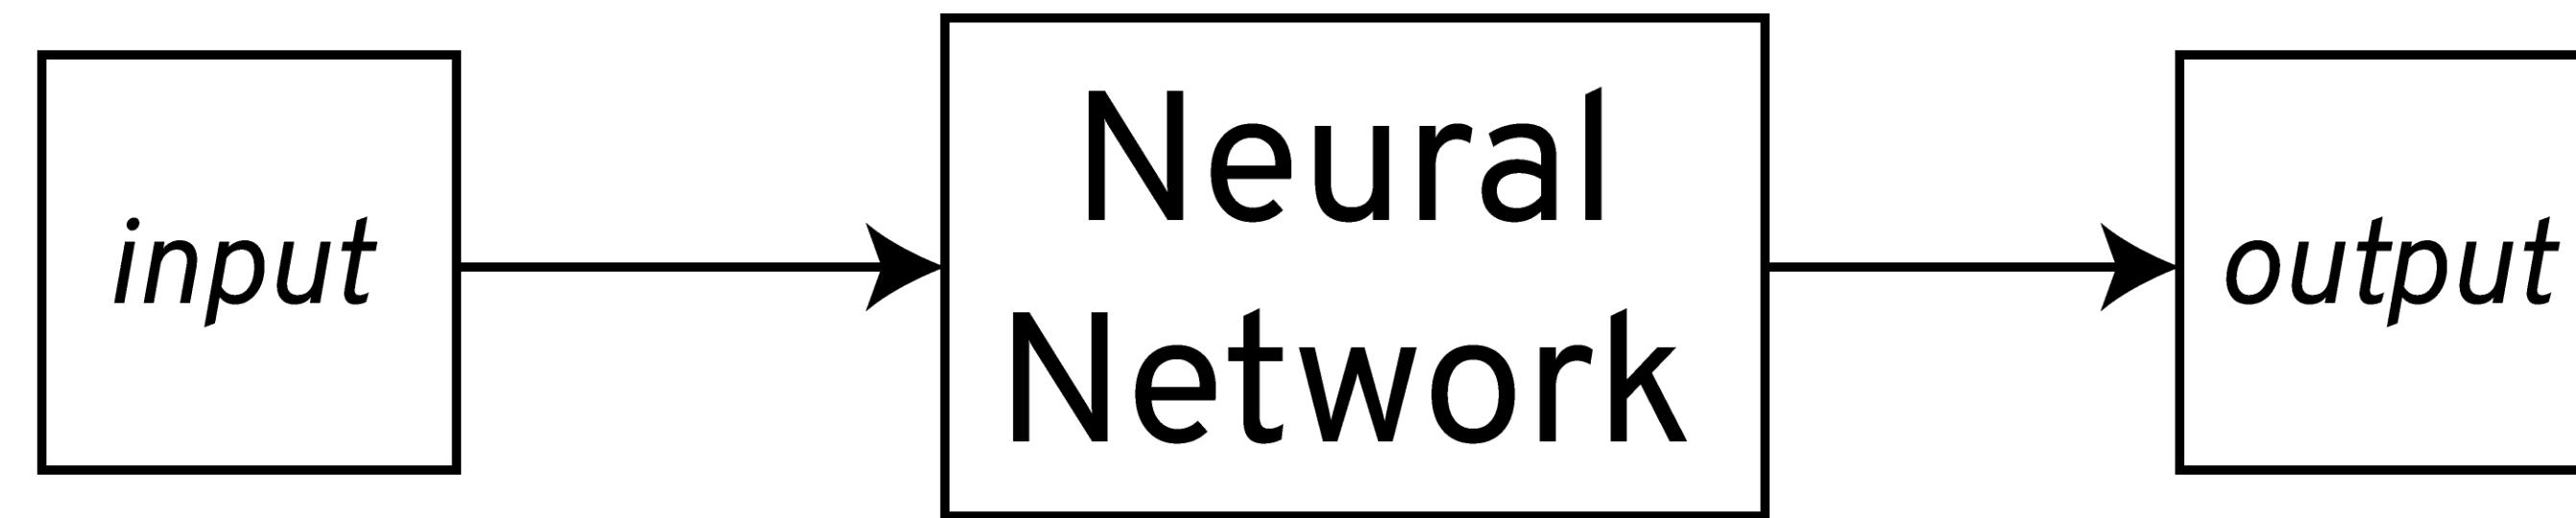
```
DataSet 64:  
rows: 100 cols: 3  
point-0      2411.1    206.48   3.9116  
point-1      69.214     65.334   3.862  
point-2      214.7     67.041   6.1554  
...  
point-97     2114.2    82.574   7.8388  
point-98     307.08    307.08   15.905  
point-99     3797.7    60.559   15.826
```

```
DataSet 63:  
rows: 100 cols: 13  
point-0      183.98    -204.99   -97.85    ...    17.116    6.7544   -2.6597  
point-1      238.64    92.367    67.025    ...    -2.1895   0.26965   3.3148  
point-2      279.43    96.95     44.27     ...    9.4531    2.0989   -1.2892  
...  
point-97     162.65    -185.08   -142.91    ...    -19.588   9.1036   14.504  
point-98     171.71    -210.44   64.047    ...    0.6732    -15.17   -8.3311  
point-99     4.409     -235.47   69.209    ...    -12.299   -6.0117   28.718
```

13 MFCCs

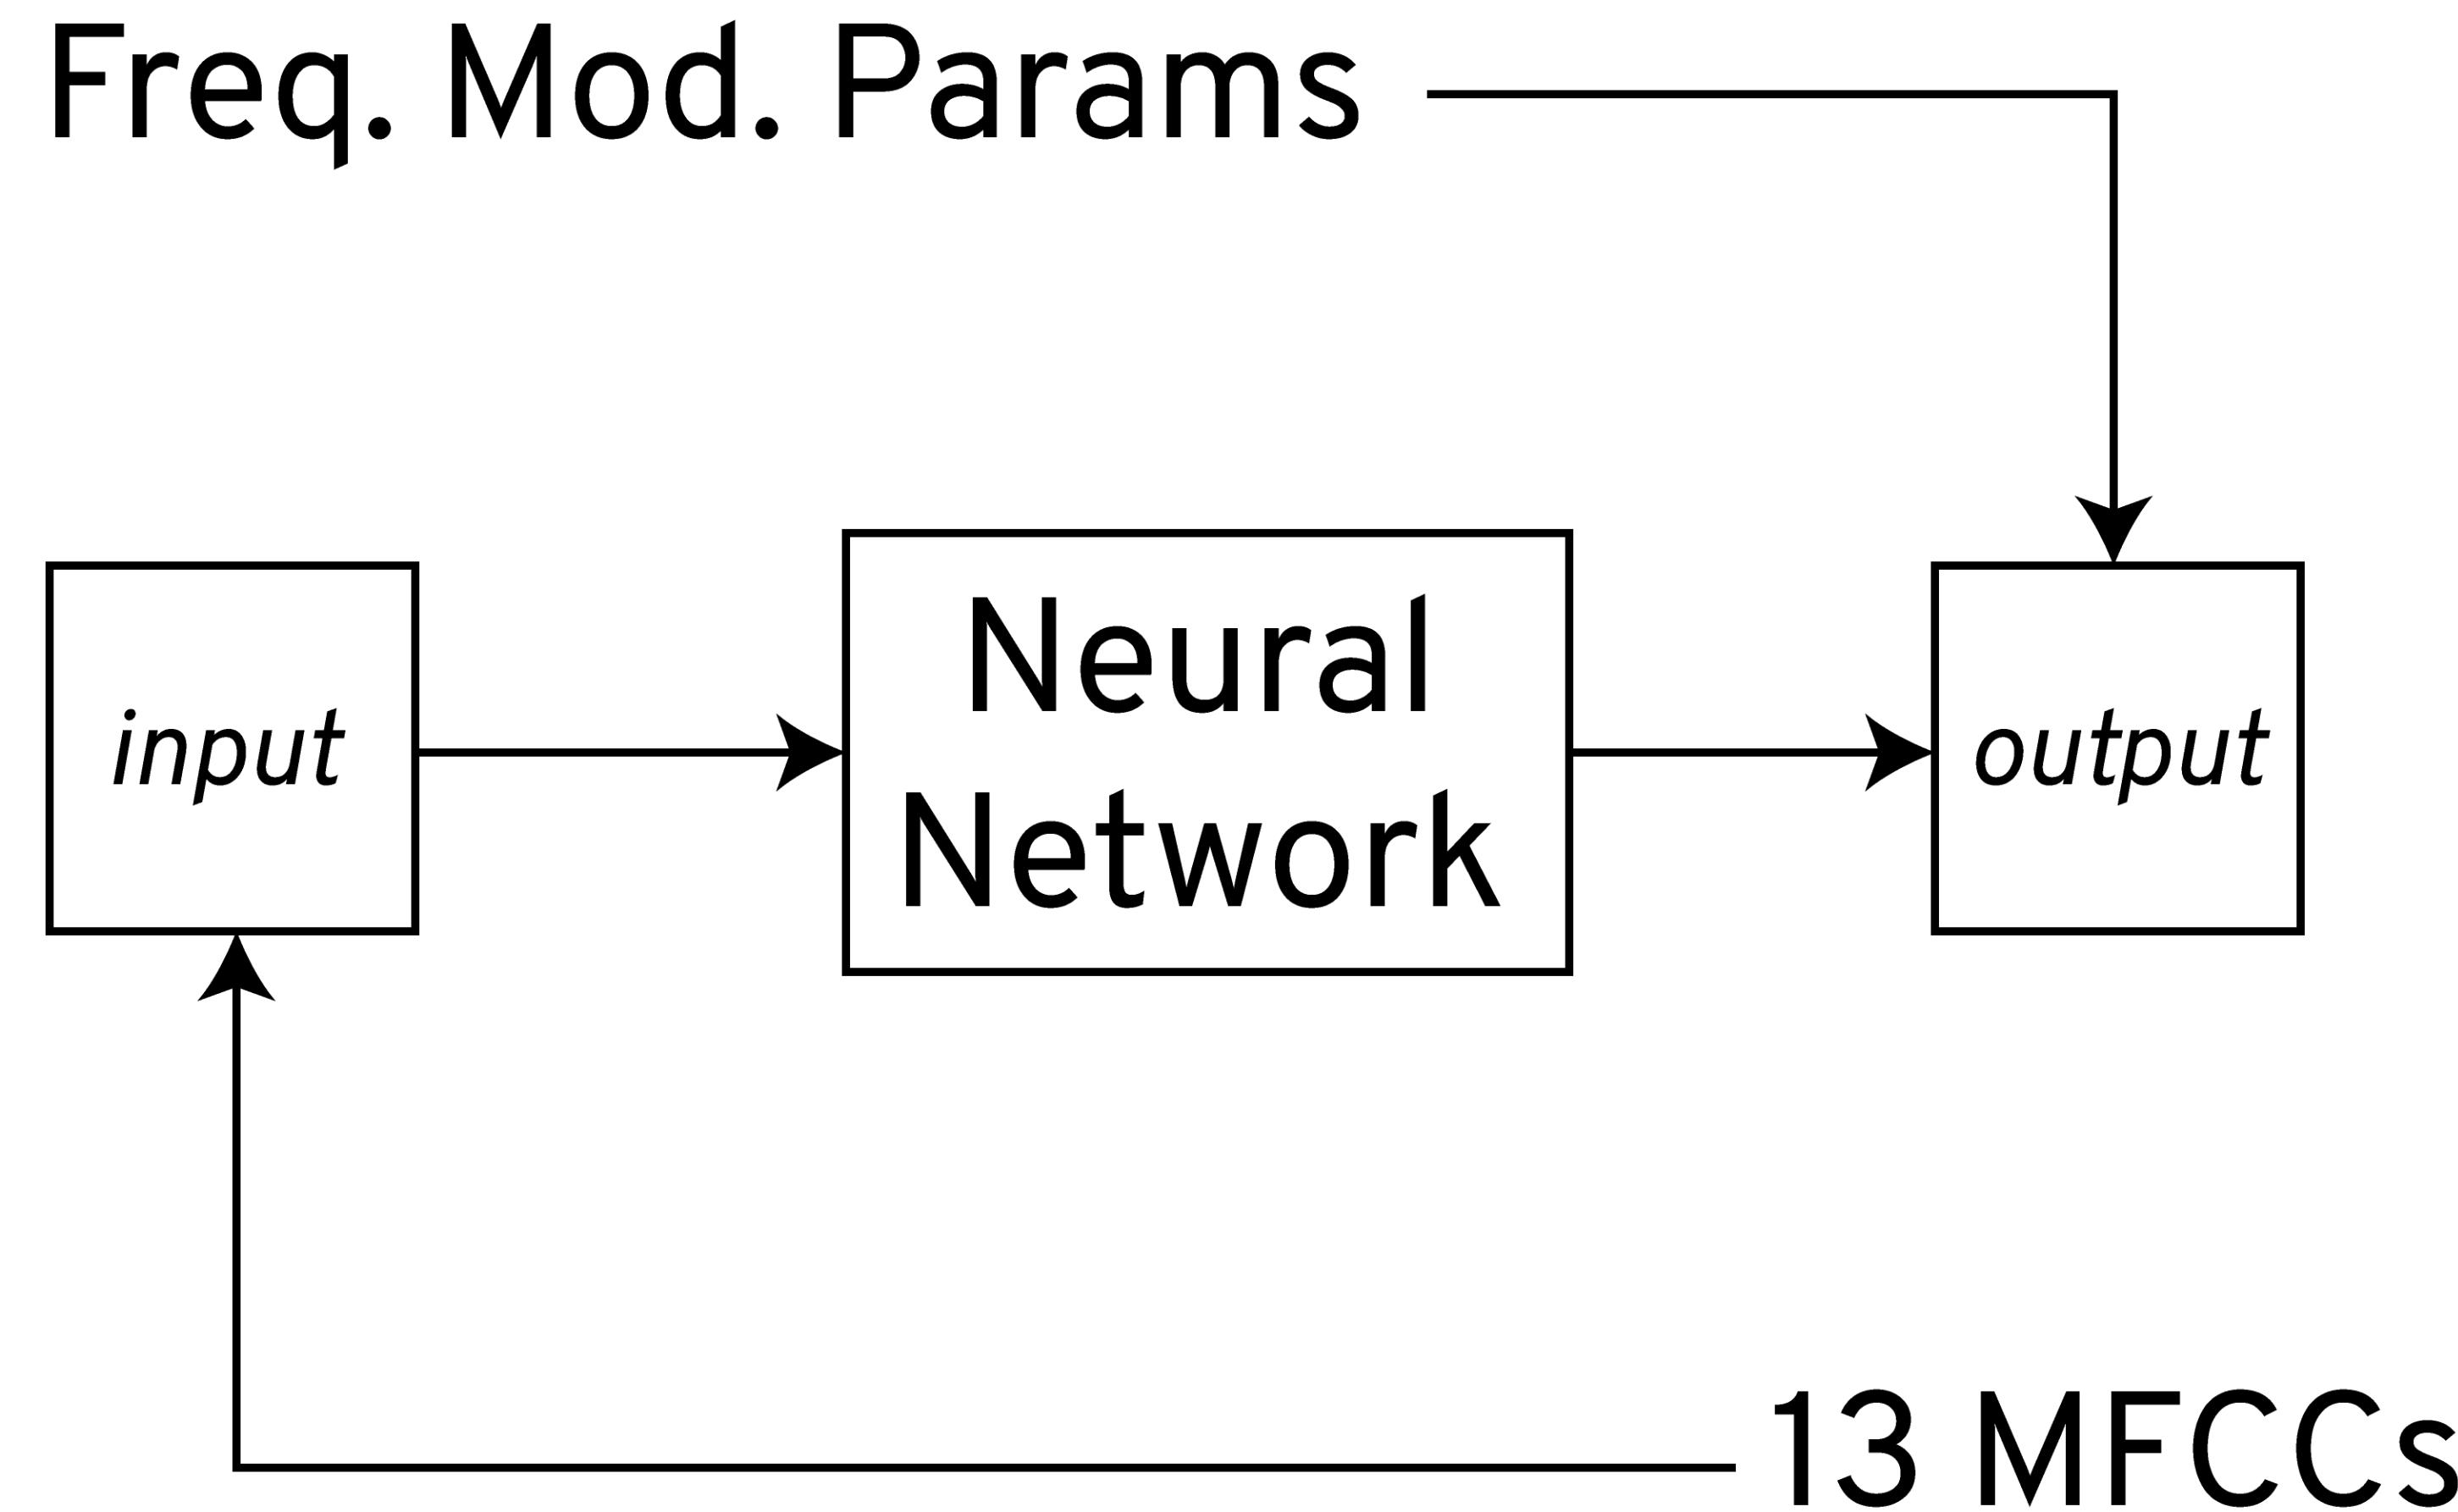
# Training

## Freq. Mod. Params

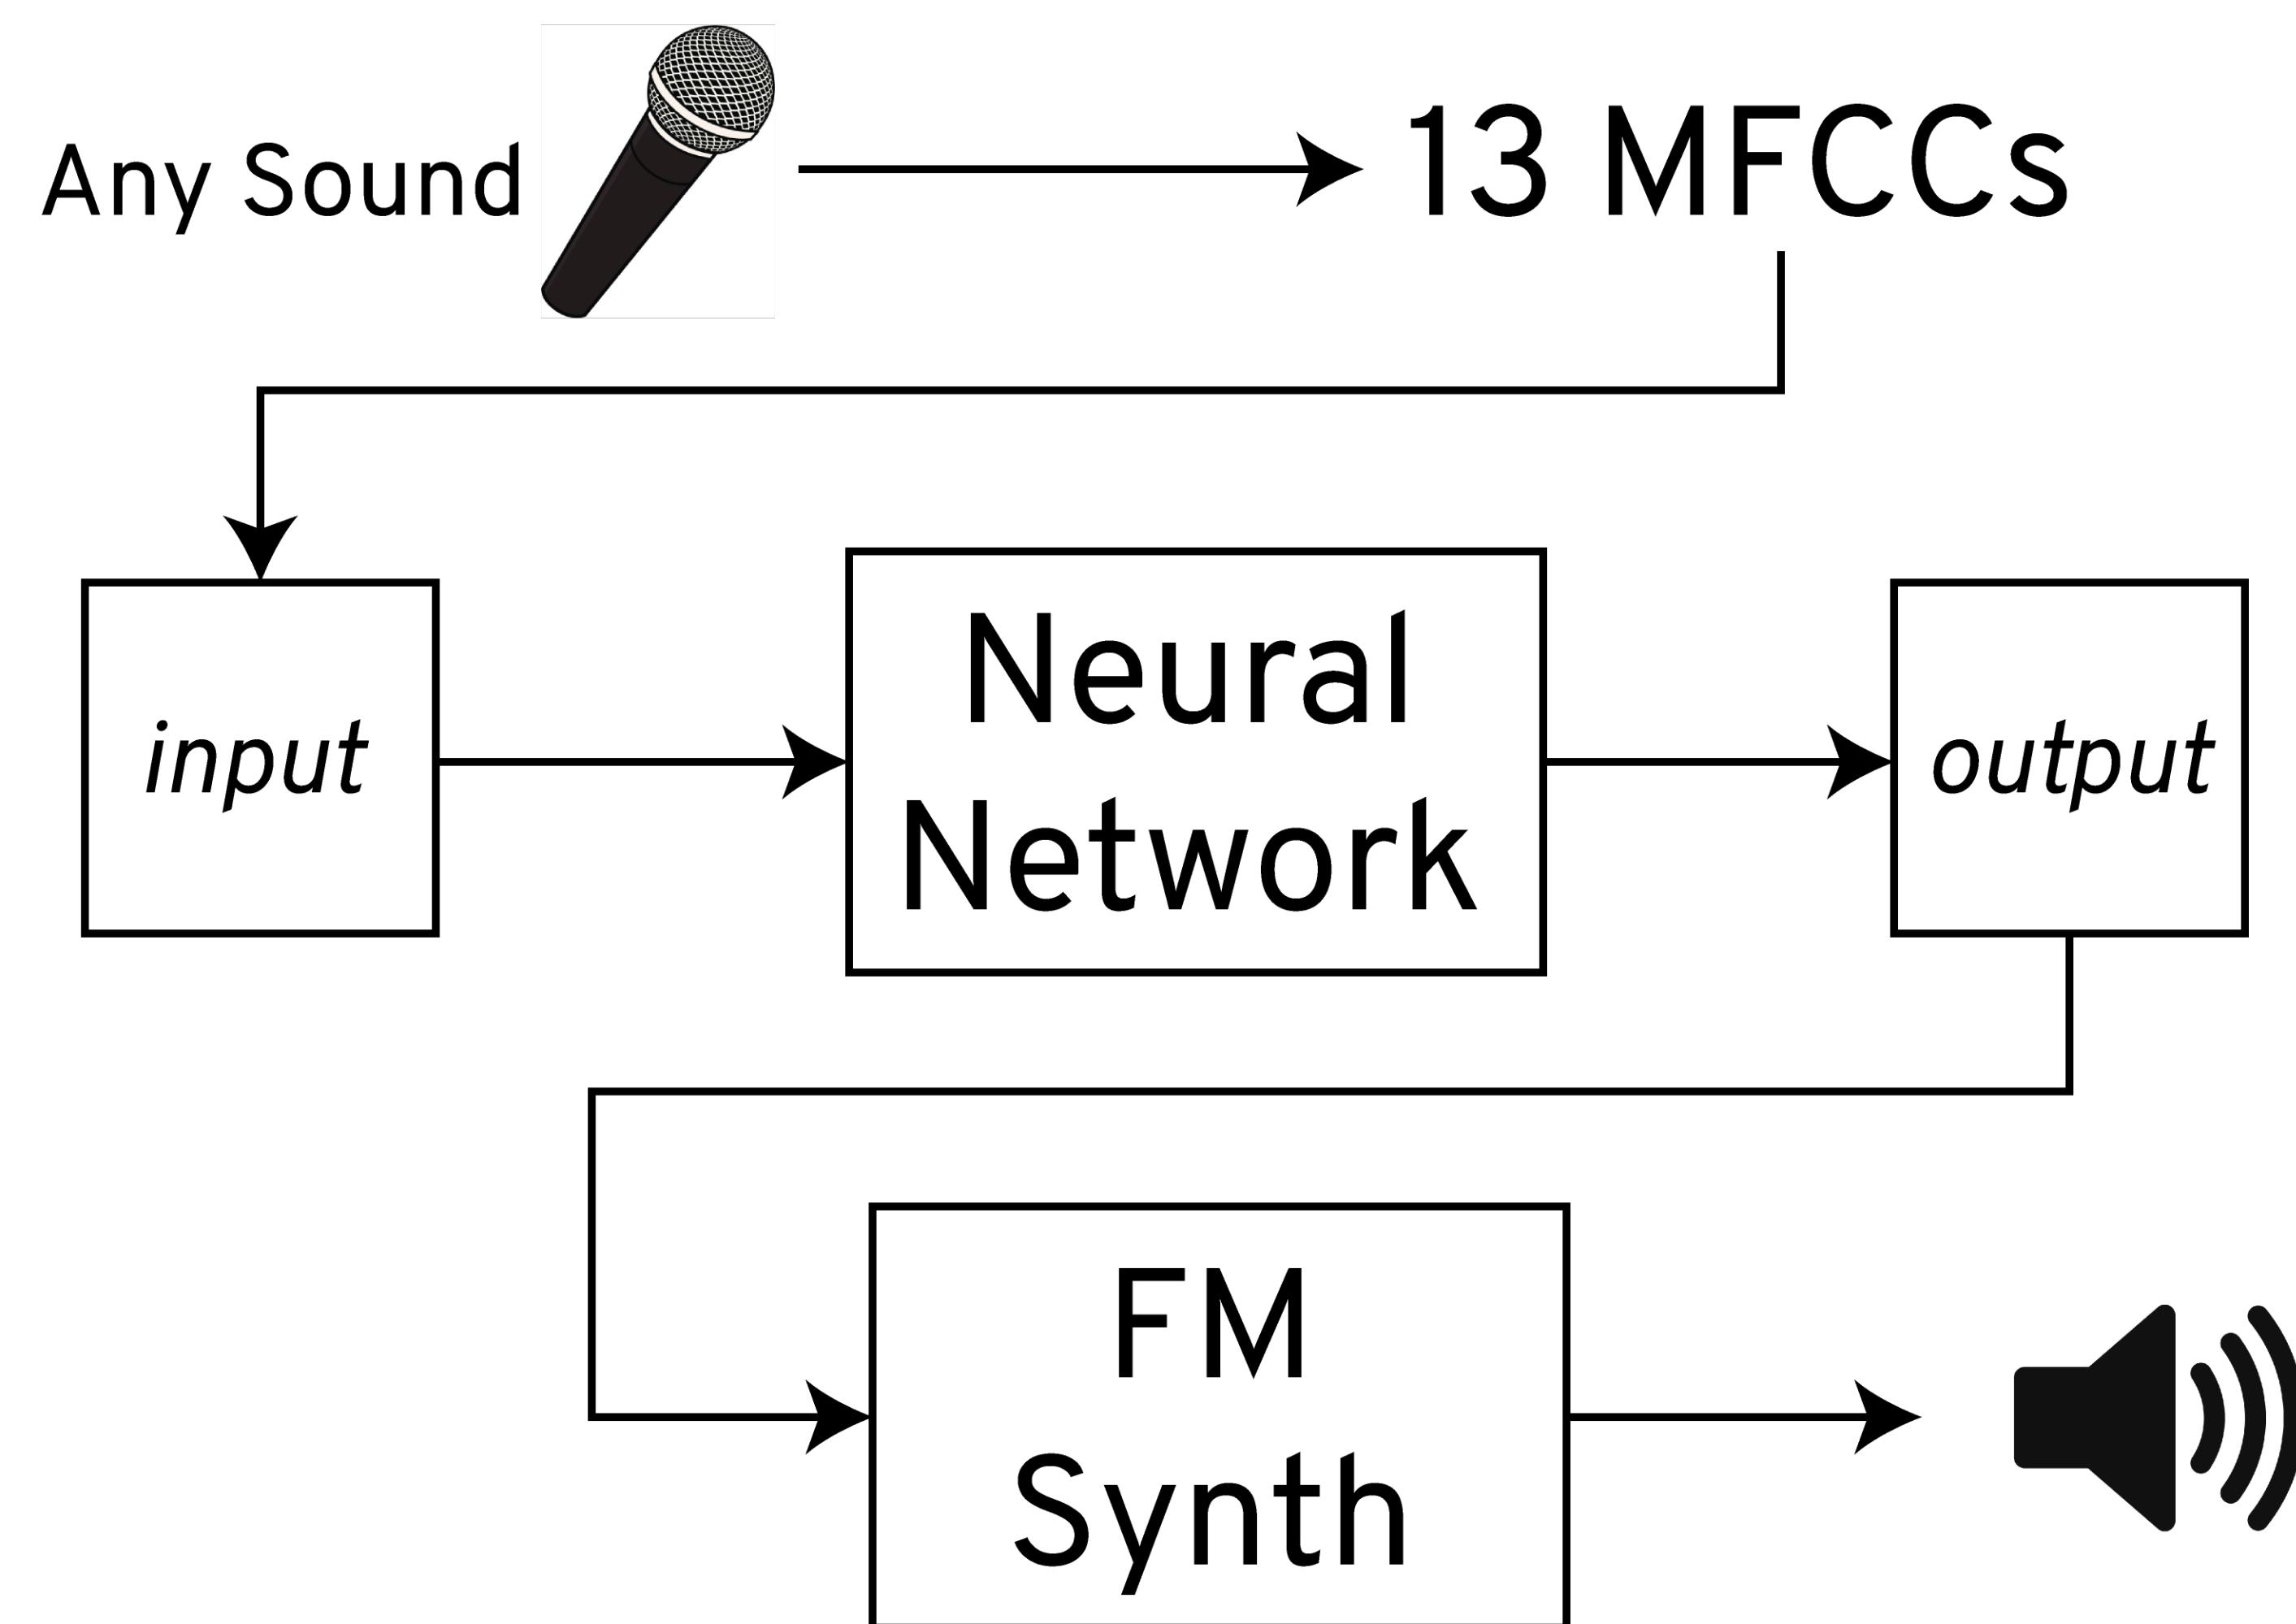


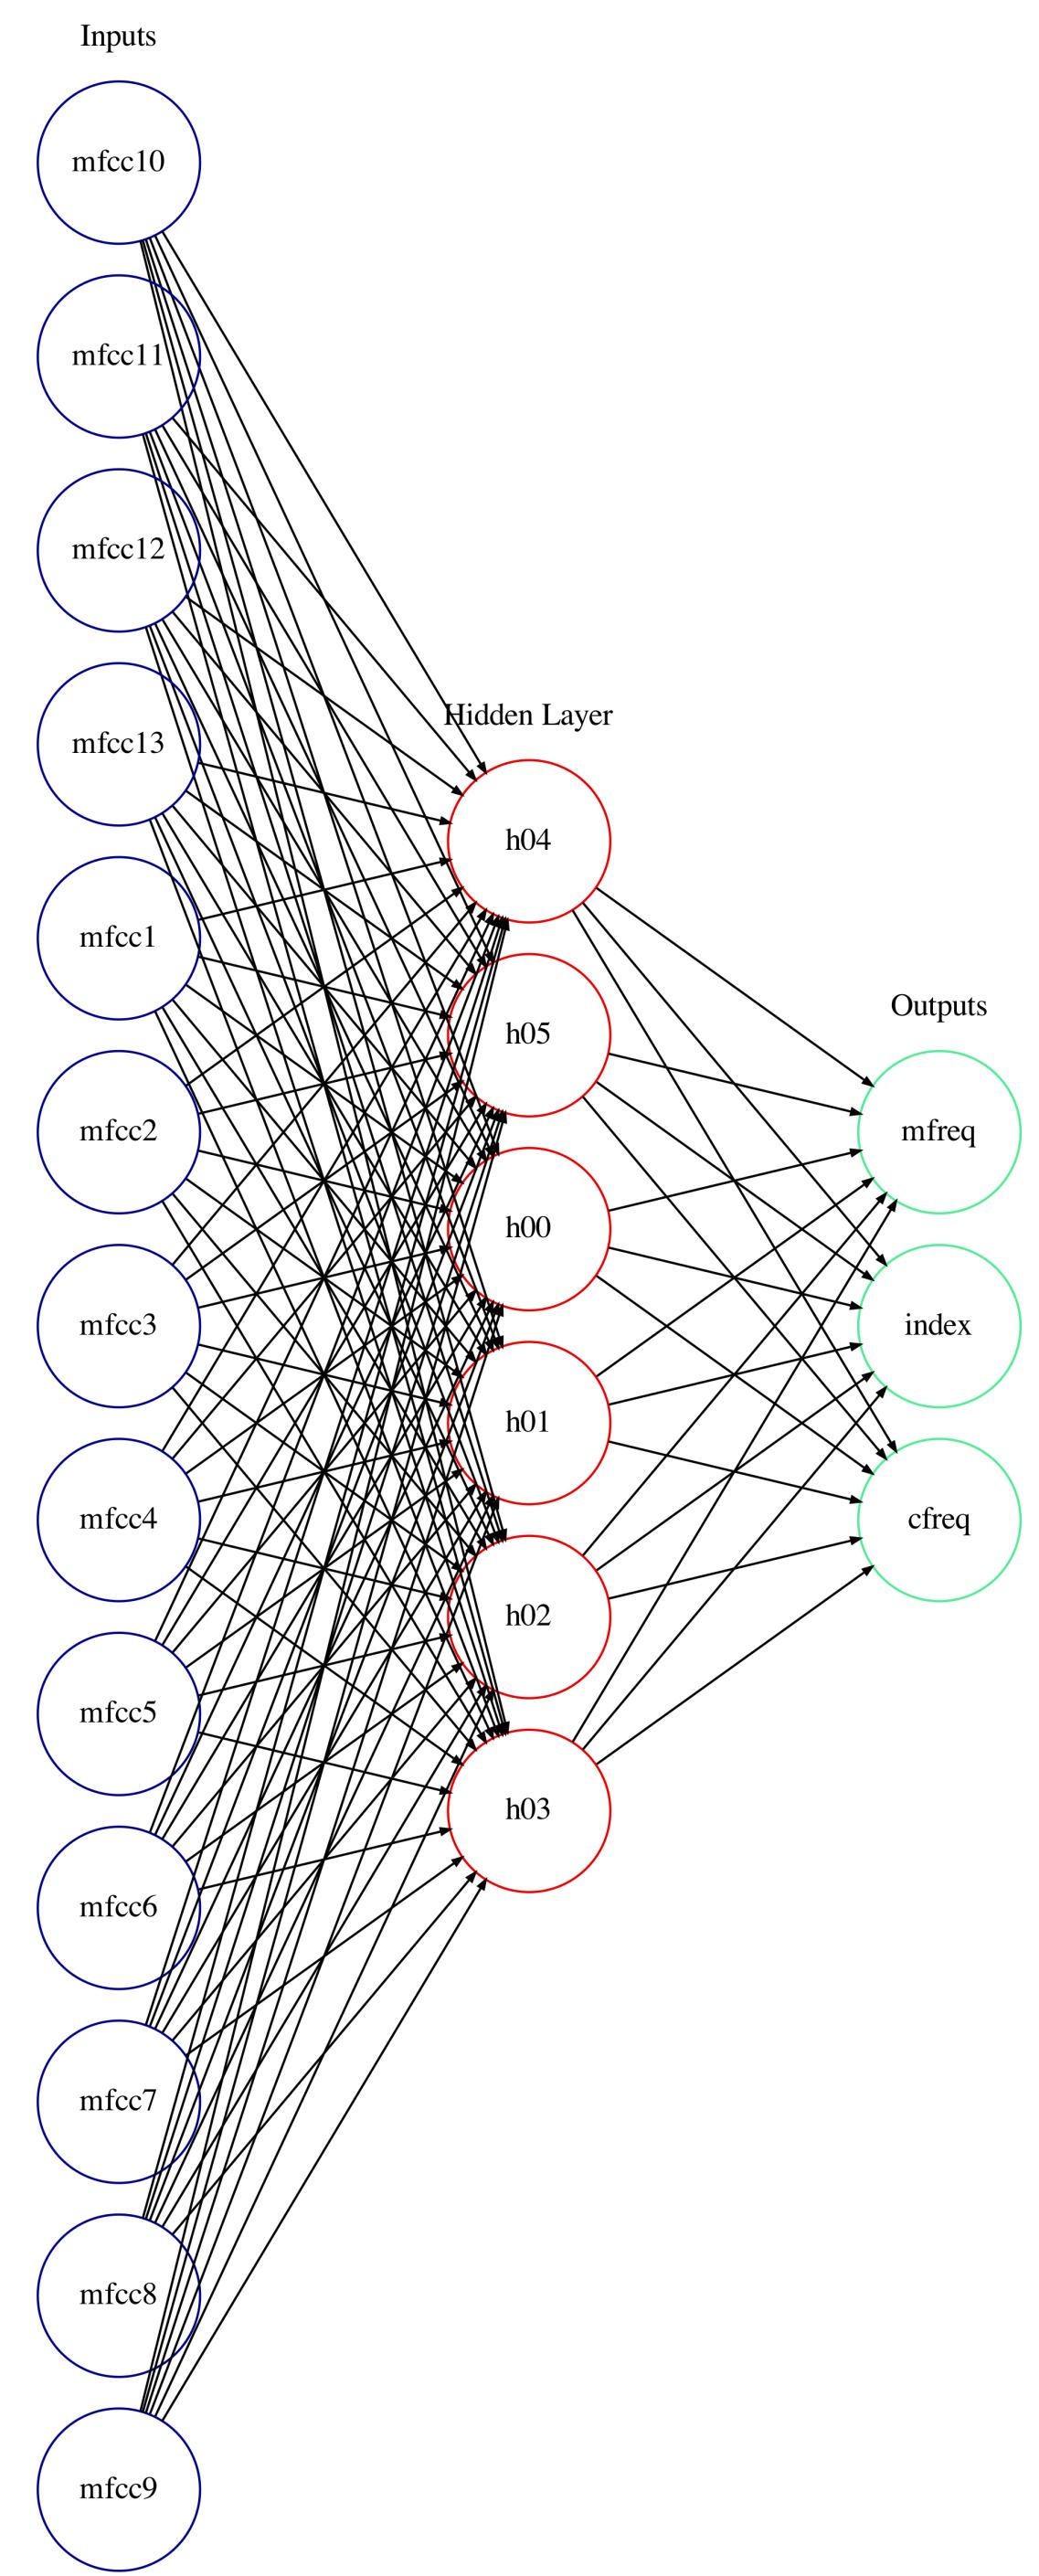
13 MFCCs

# Training



# Performing



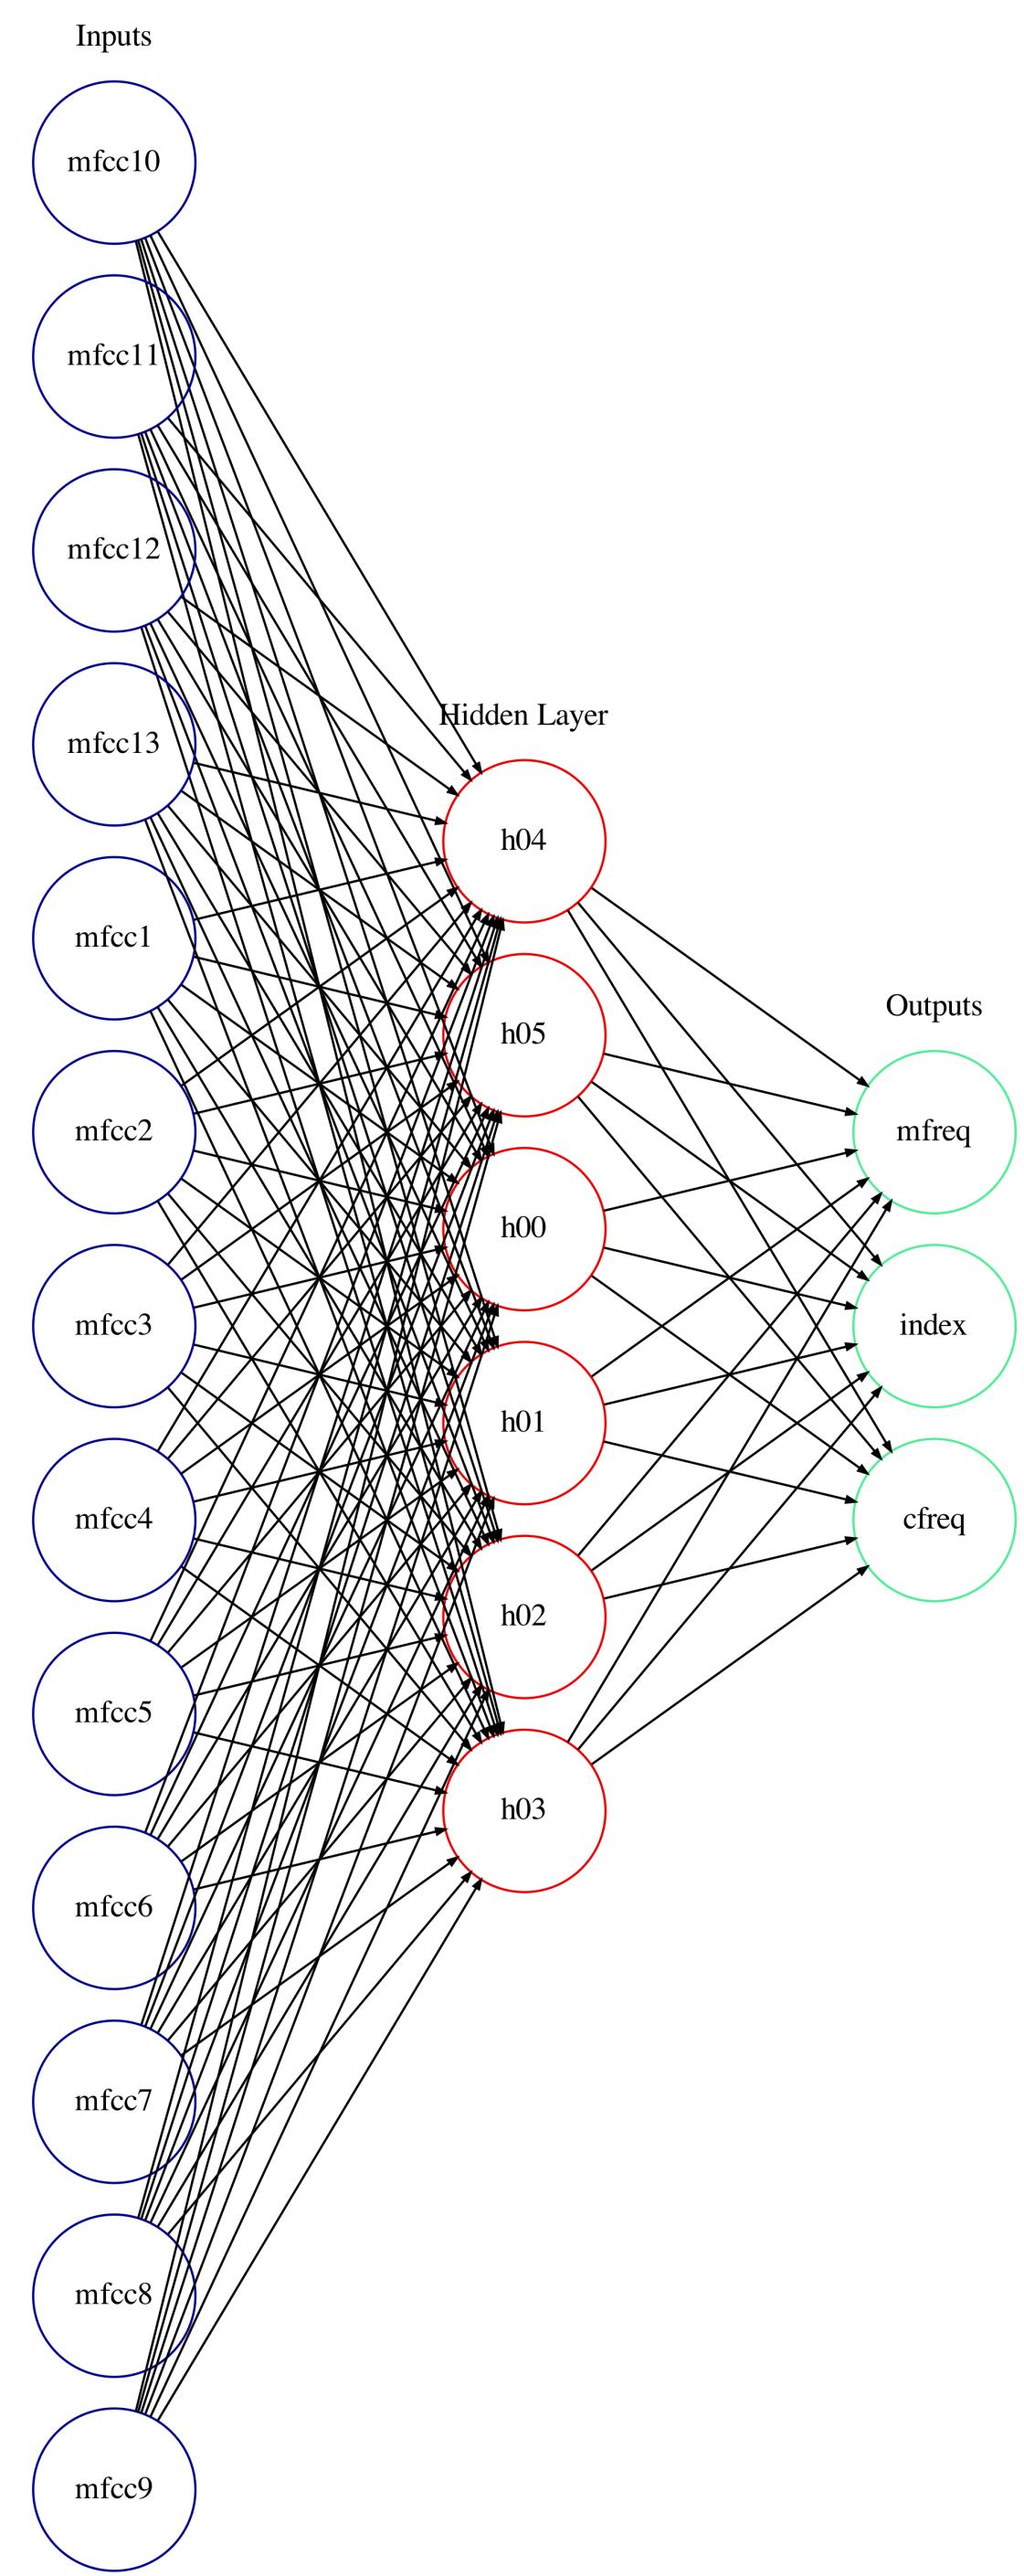




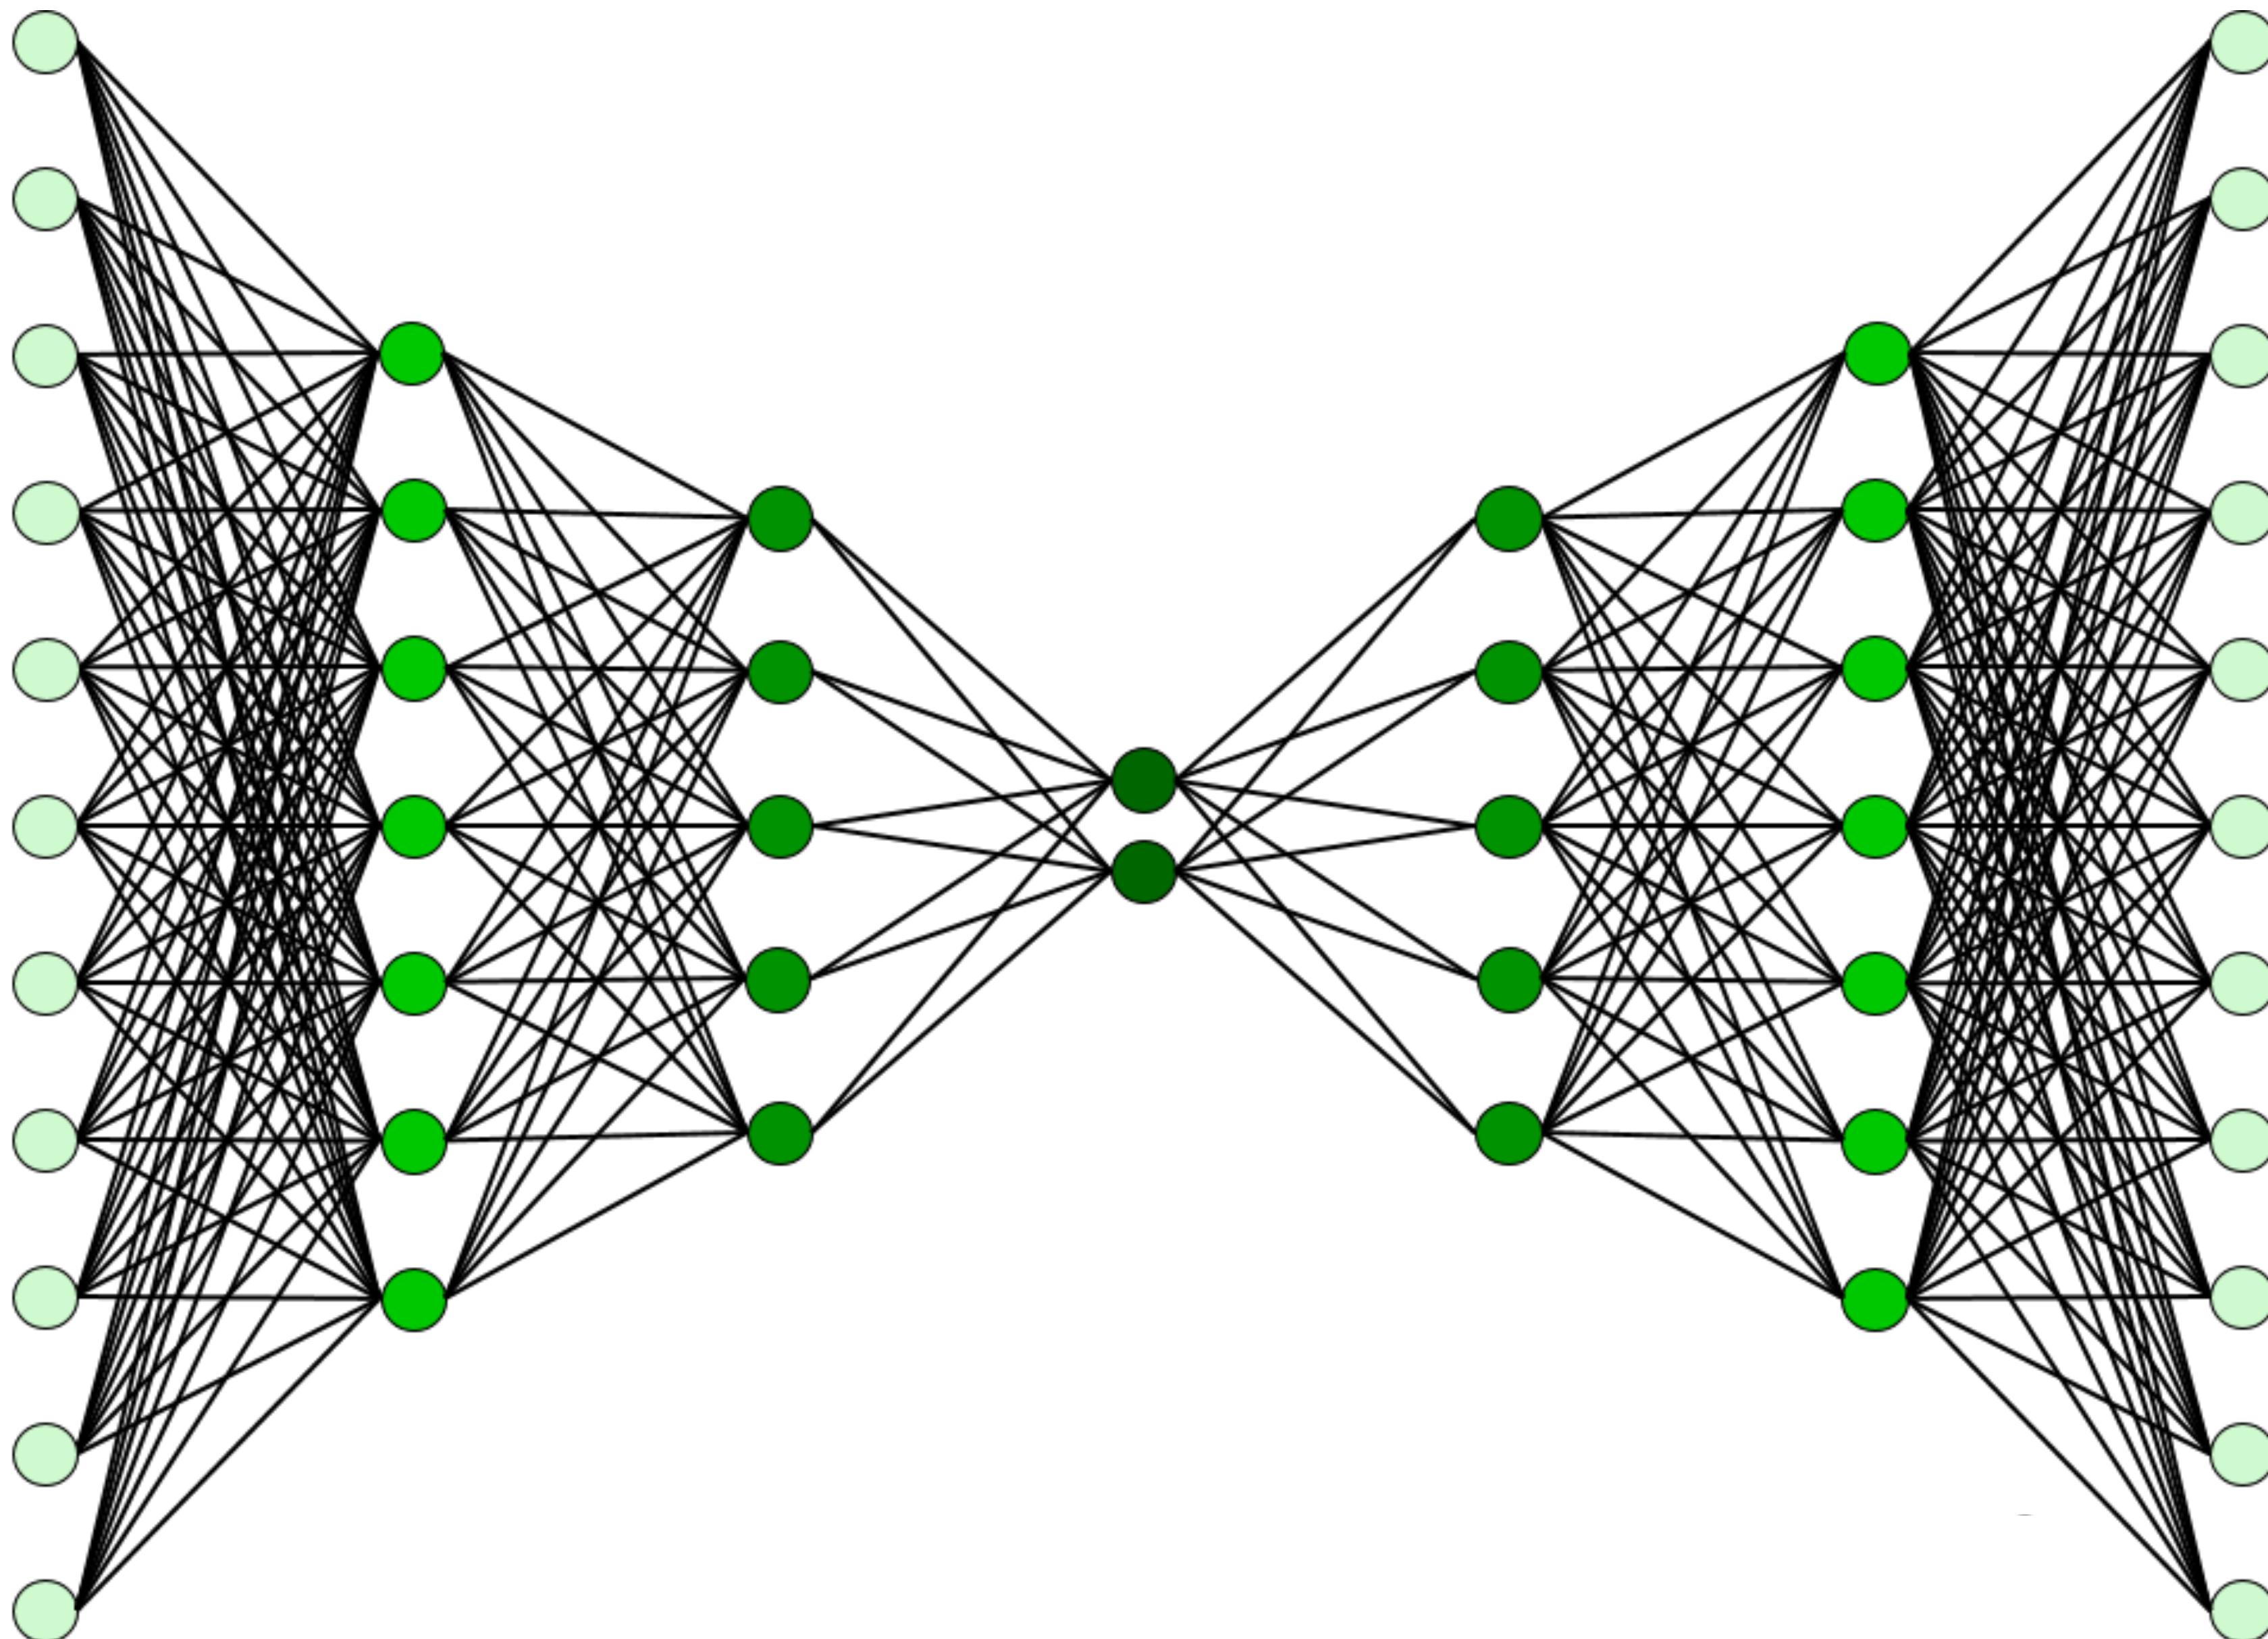
# *Wavetable Autoencoder*

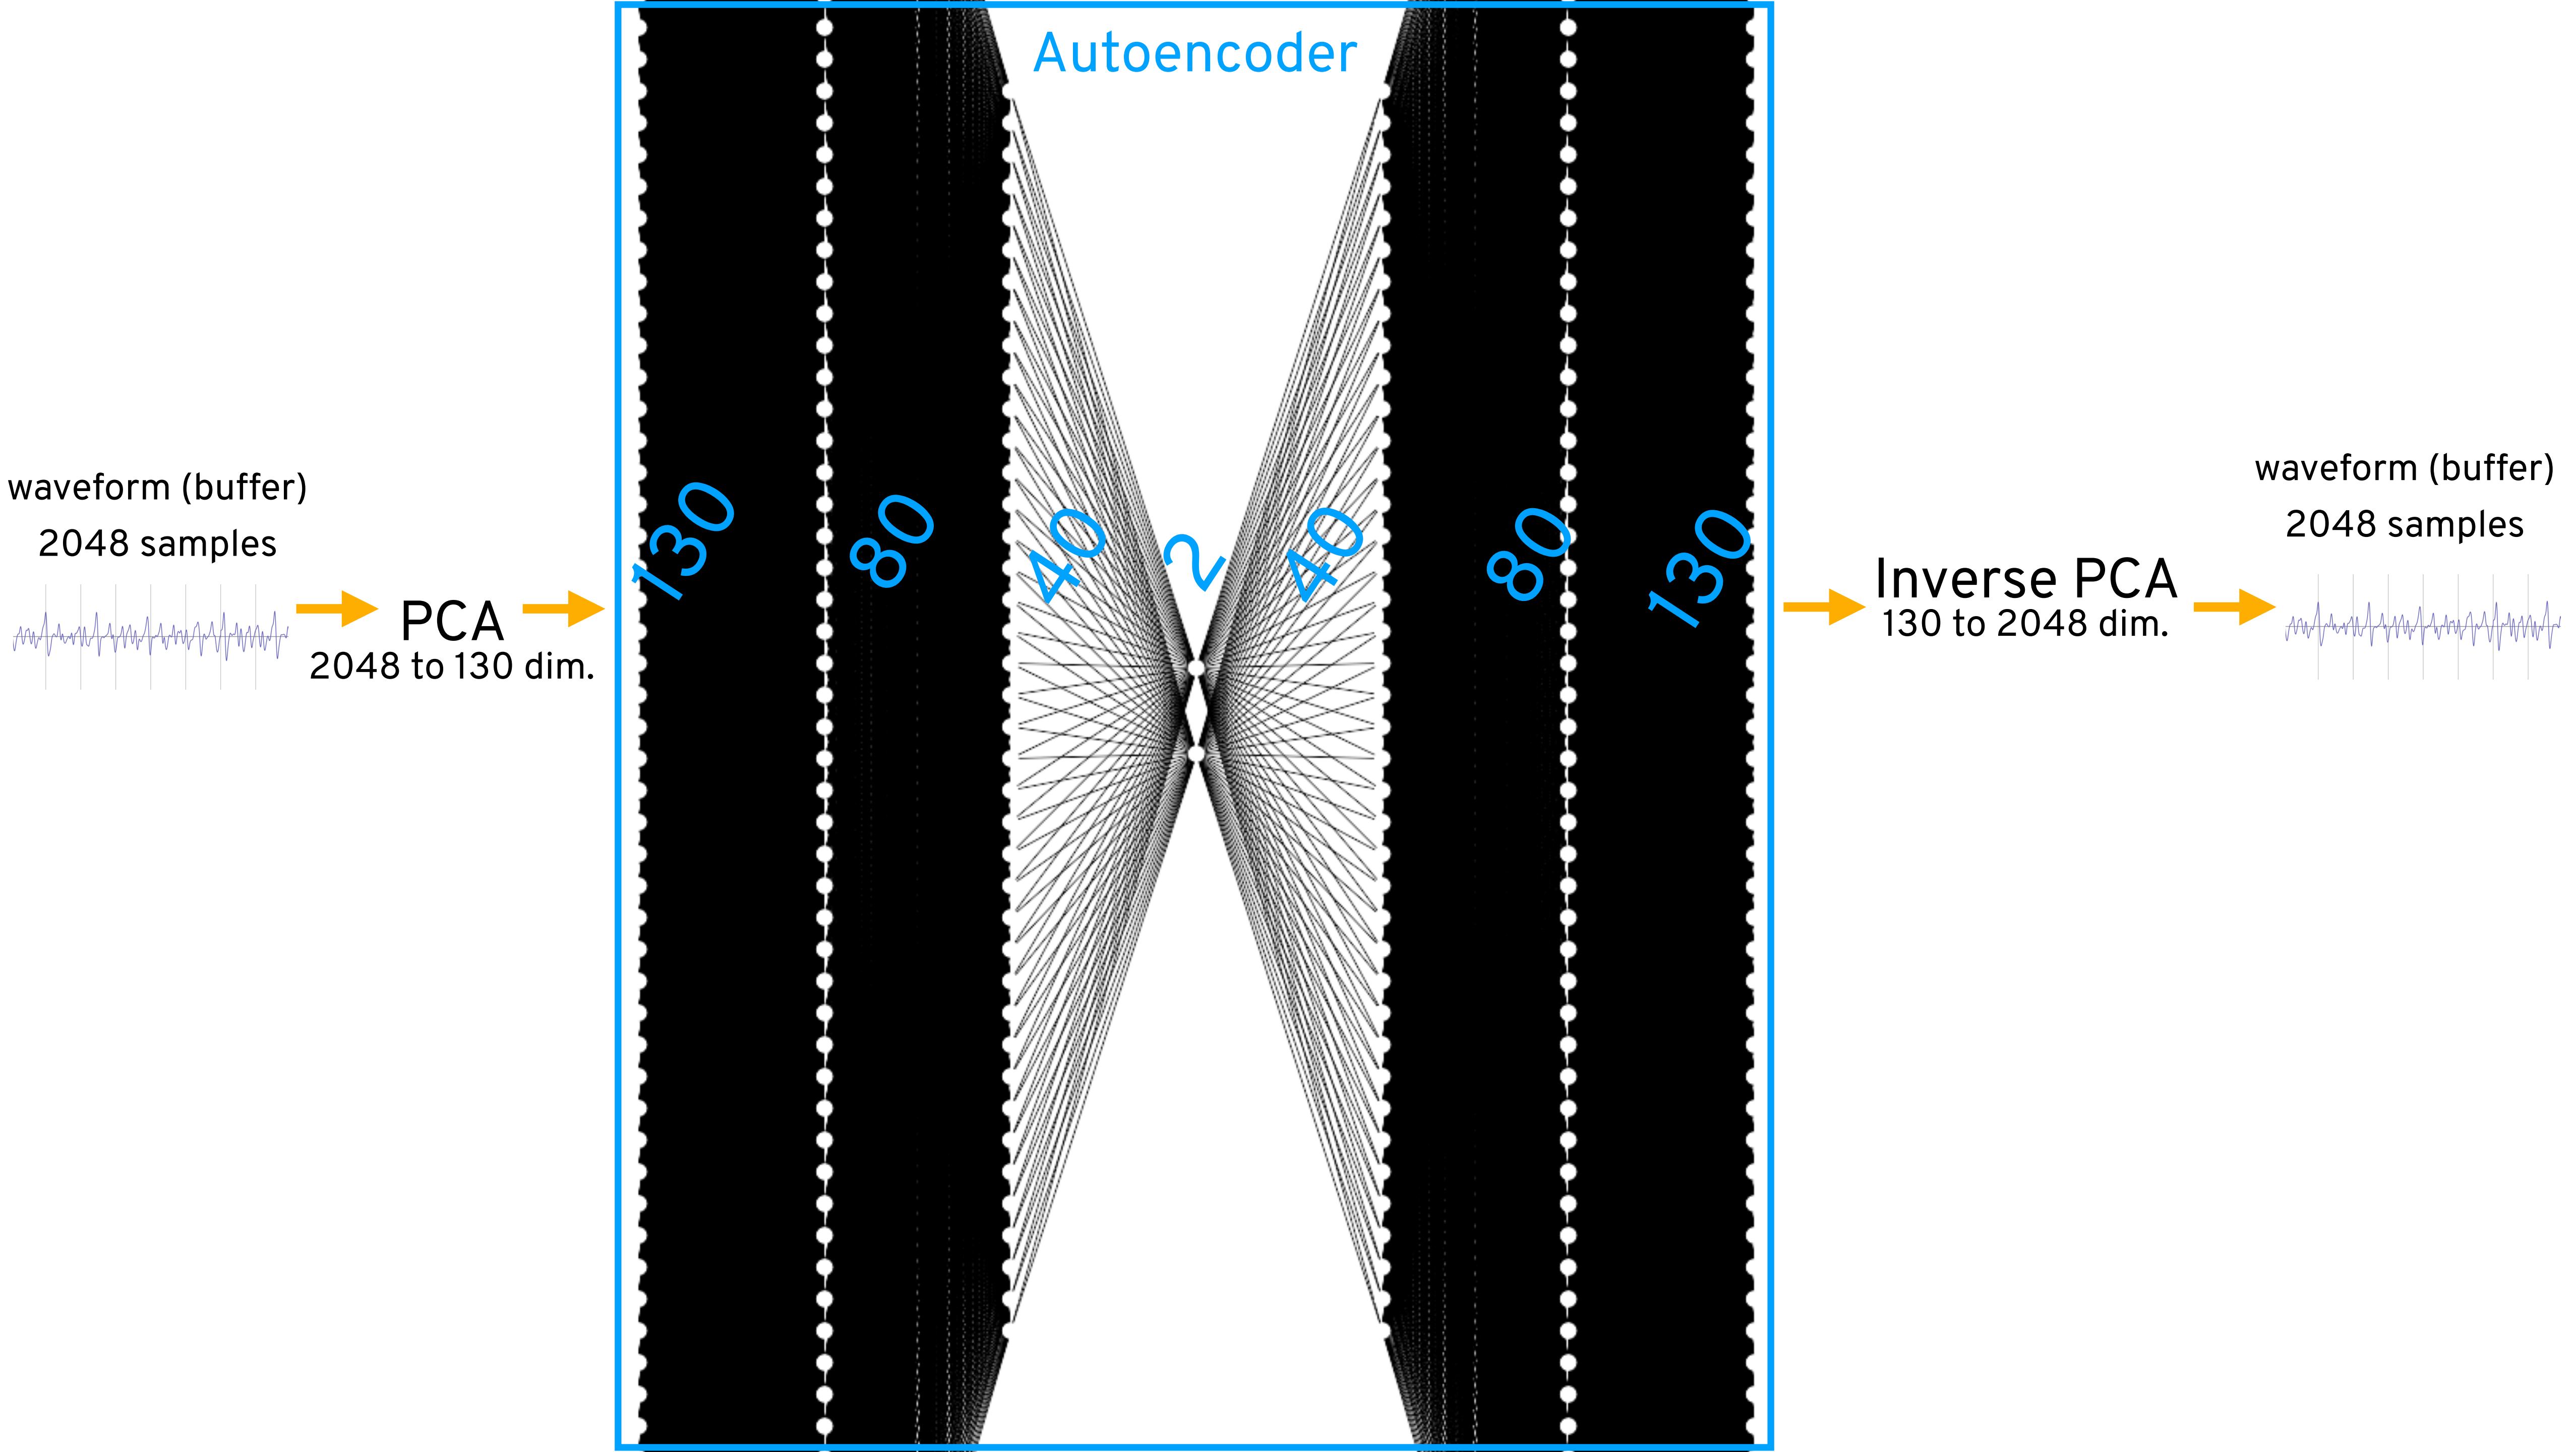


# MFCCs → FM Params



# Autoencoder





[flucoma.org](http://flucoma.org)  
[learn.flucoma.org](http://learn.flucoma.org)



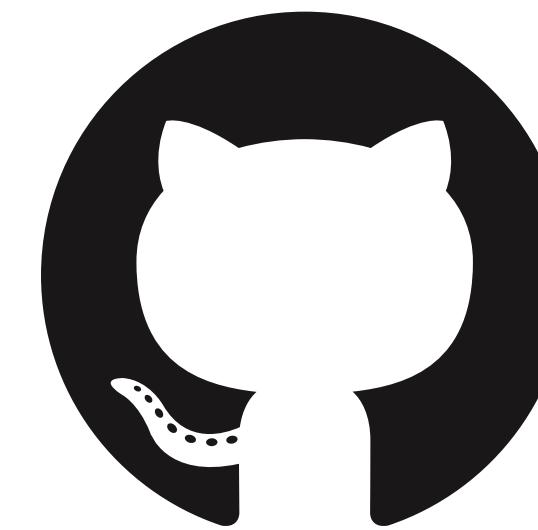
download FluCoMa package  
toolkit reference  
learn from tutorials  
hear & see works made with FluCoMa

[discourse.flucoma.org](https://discourse.flucoma.org)



share works in progress  
ask questions  
get ideas

<https://github.com/flucoma>



bug reports  
pull requests

[https://www.youtube.com/c/  
fluidcorpusmanipulation](https://www.youtube.com/c/fluidcorpusmanipulation)



tutorials  
performances

